

Monitoring & Observability with Istio

Emma - Yang XY Yang/China/IBM
bjyangyy@cn.ibm.com

About me

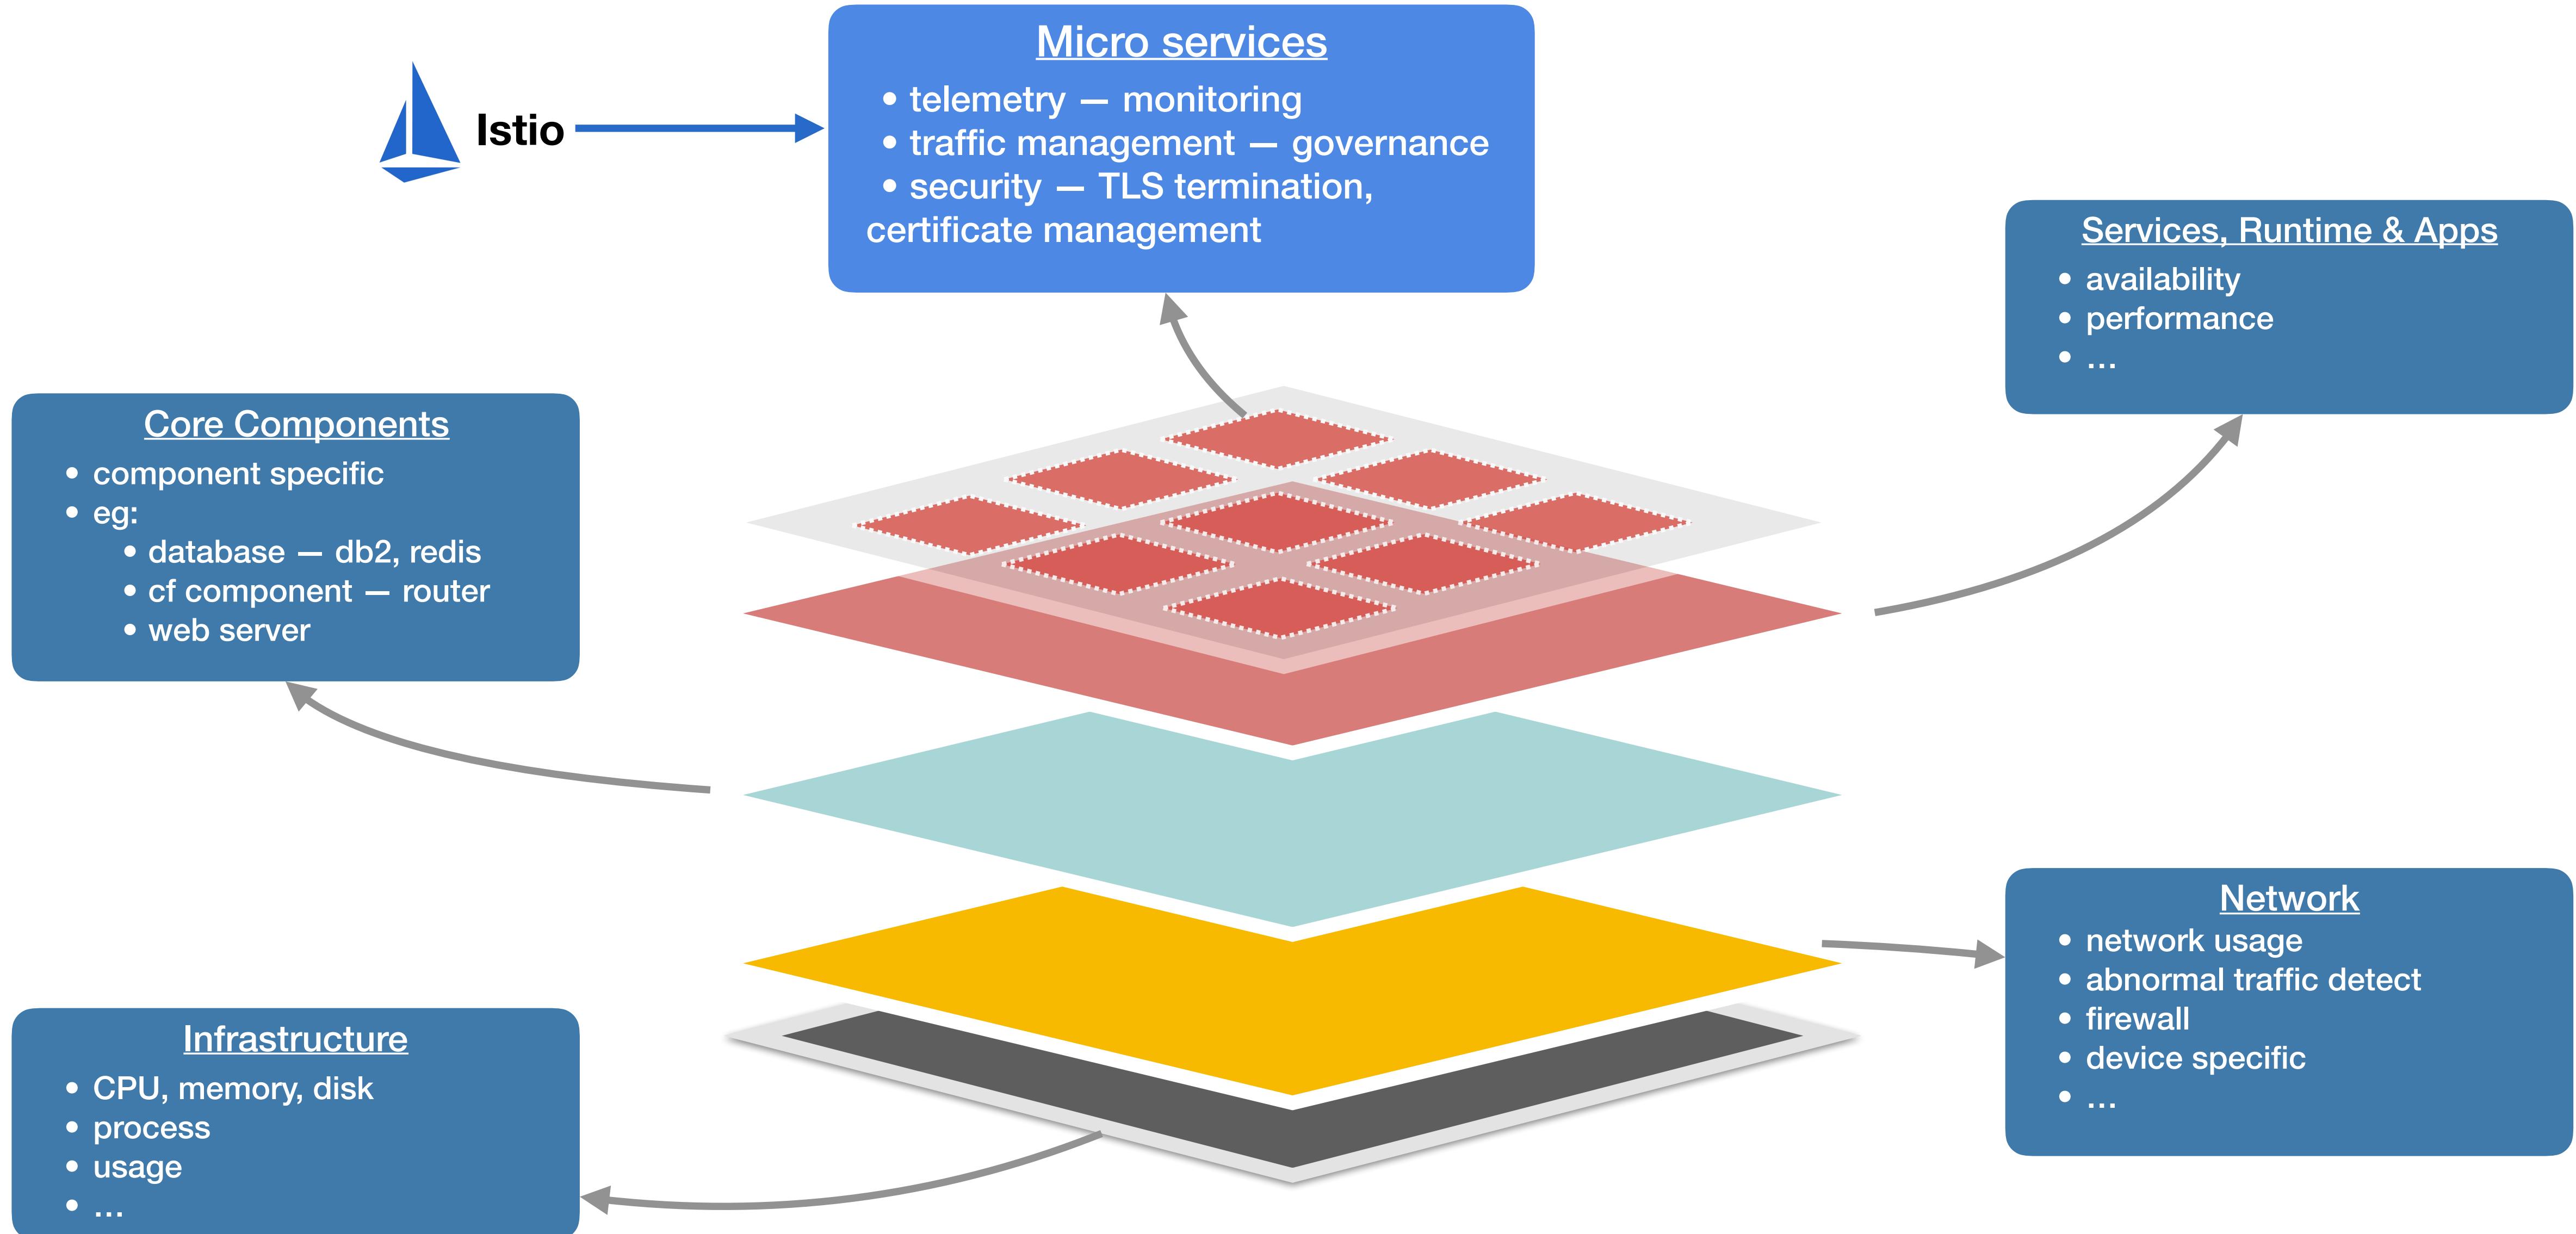


杨洋

bjyangyy@cn.ibm.com

Advisory Software Engineer
CTO OSS Monitor team, IBM

Why Istio for monitoring?

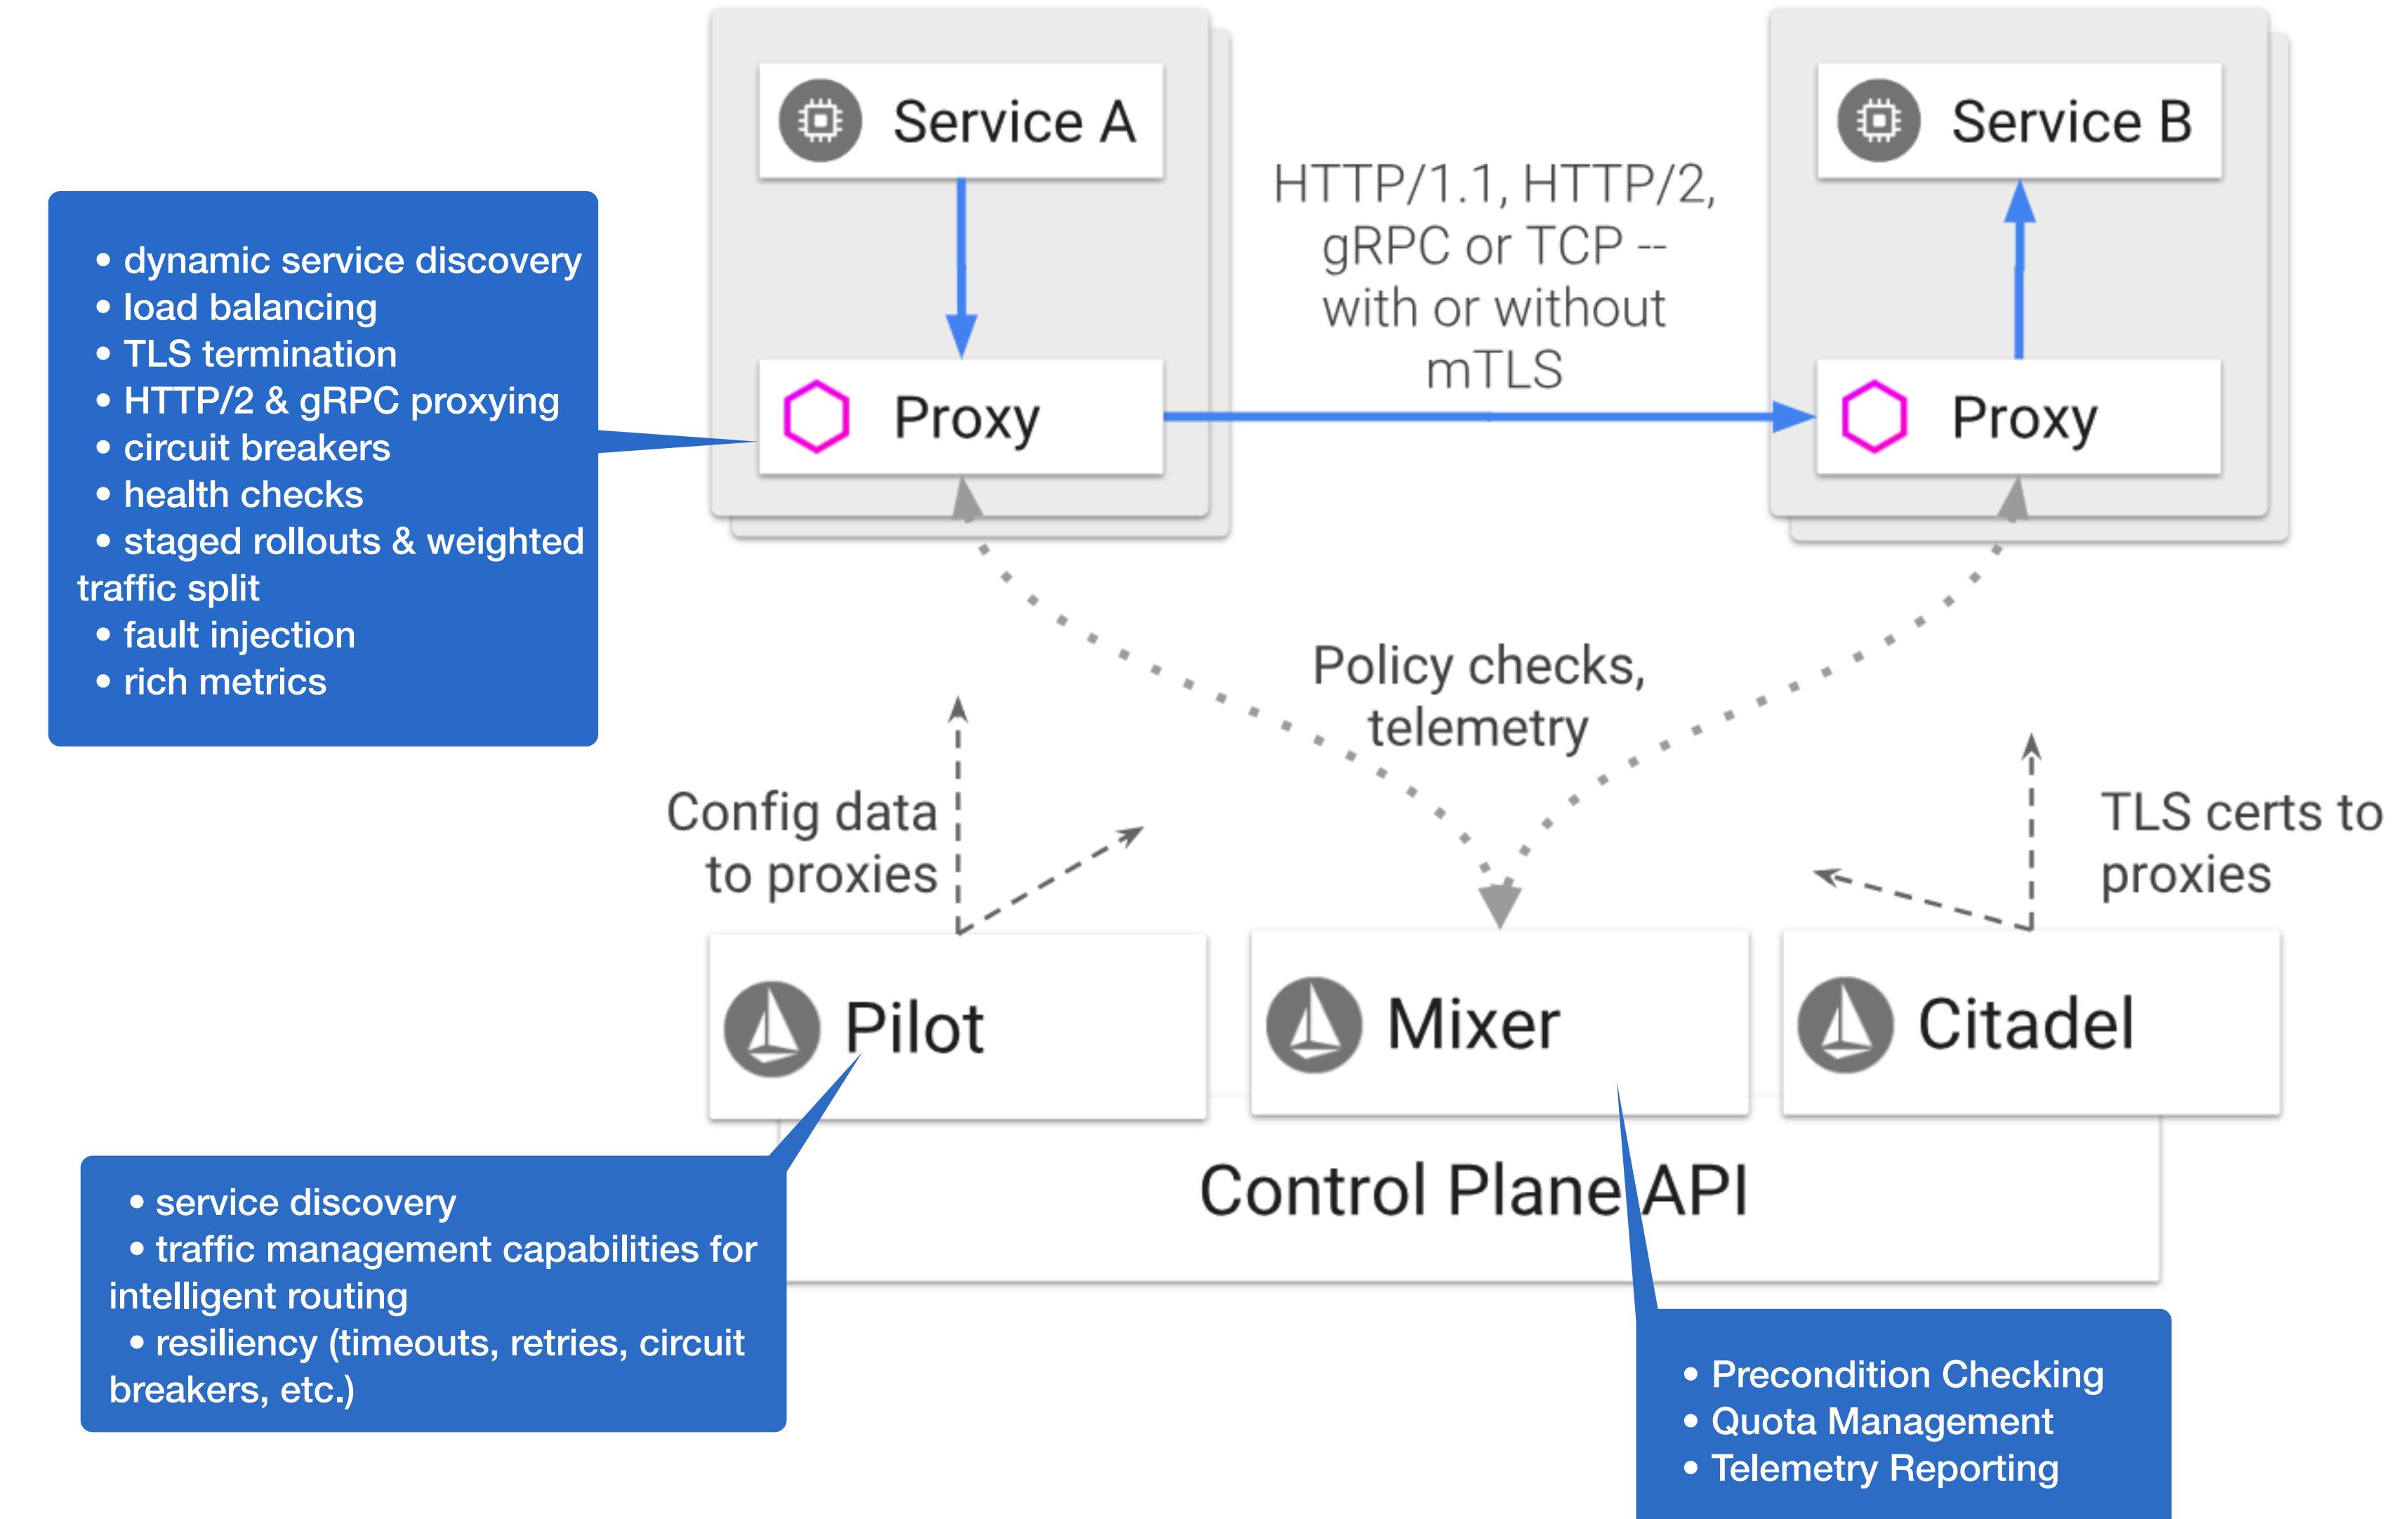


Istio

- **Istio** lets you connect, secure, control, and observe services
- Heavy investments from industry-leading companies such as **Google, IBM, and Lyft**
- Release **1.1**
- Active community

Core Features

- **Telemetry**
- **Fine-grained traffic management**
- **Rich load balancing**
- **Circuit breaking, timeouts, retries**
- **Request Tracing**
- Fault Injection
- Rate limits
- Load shedding

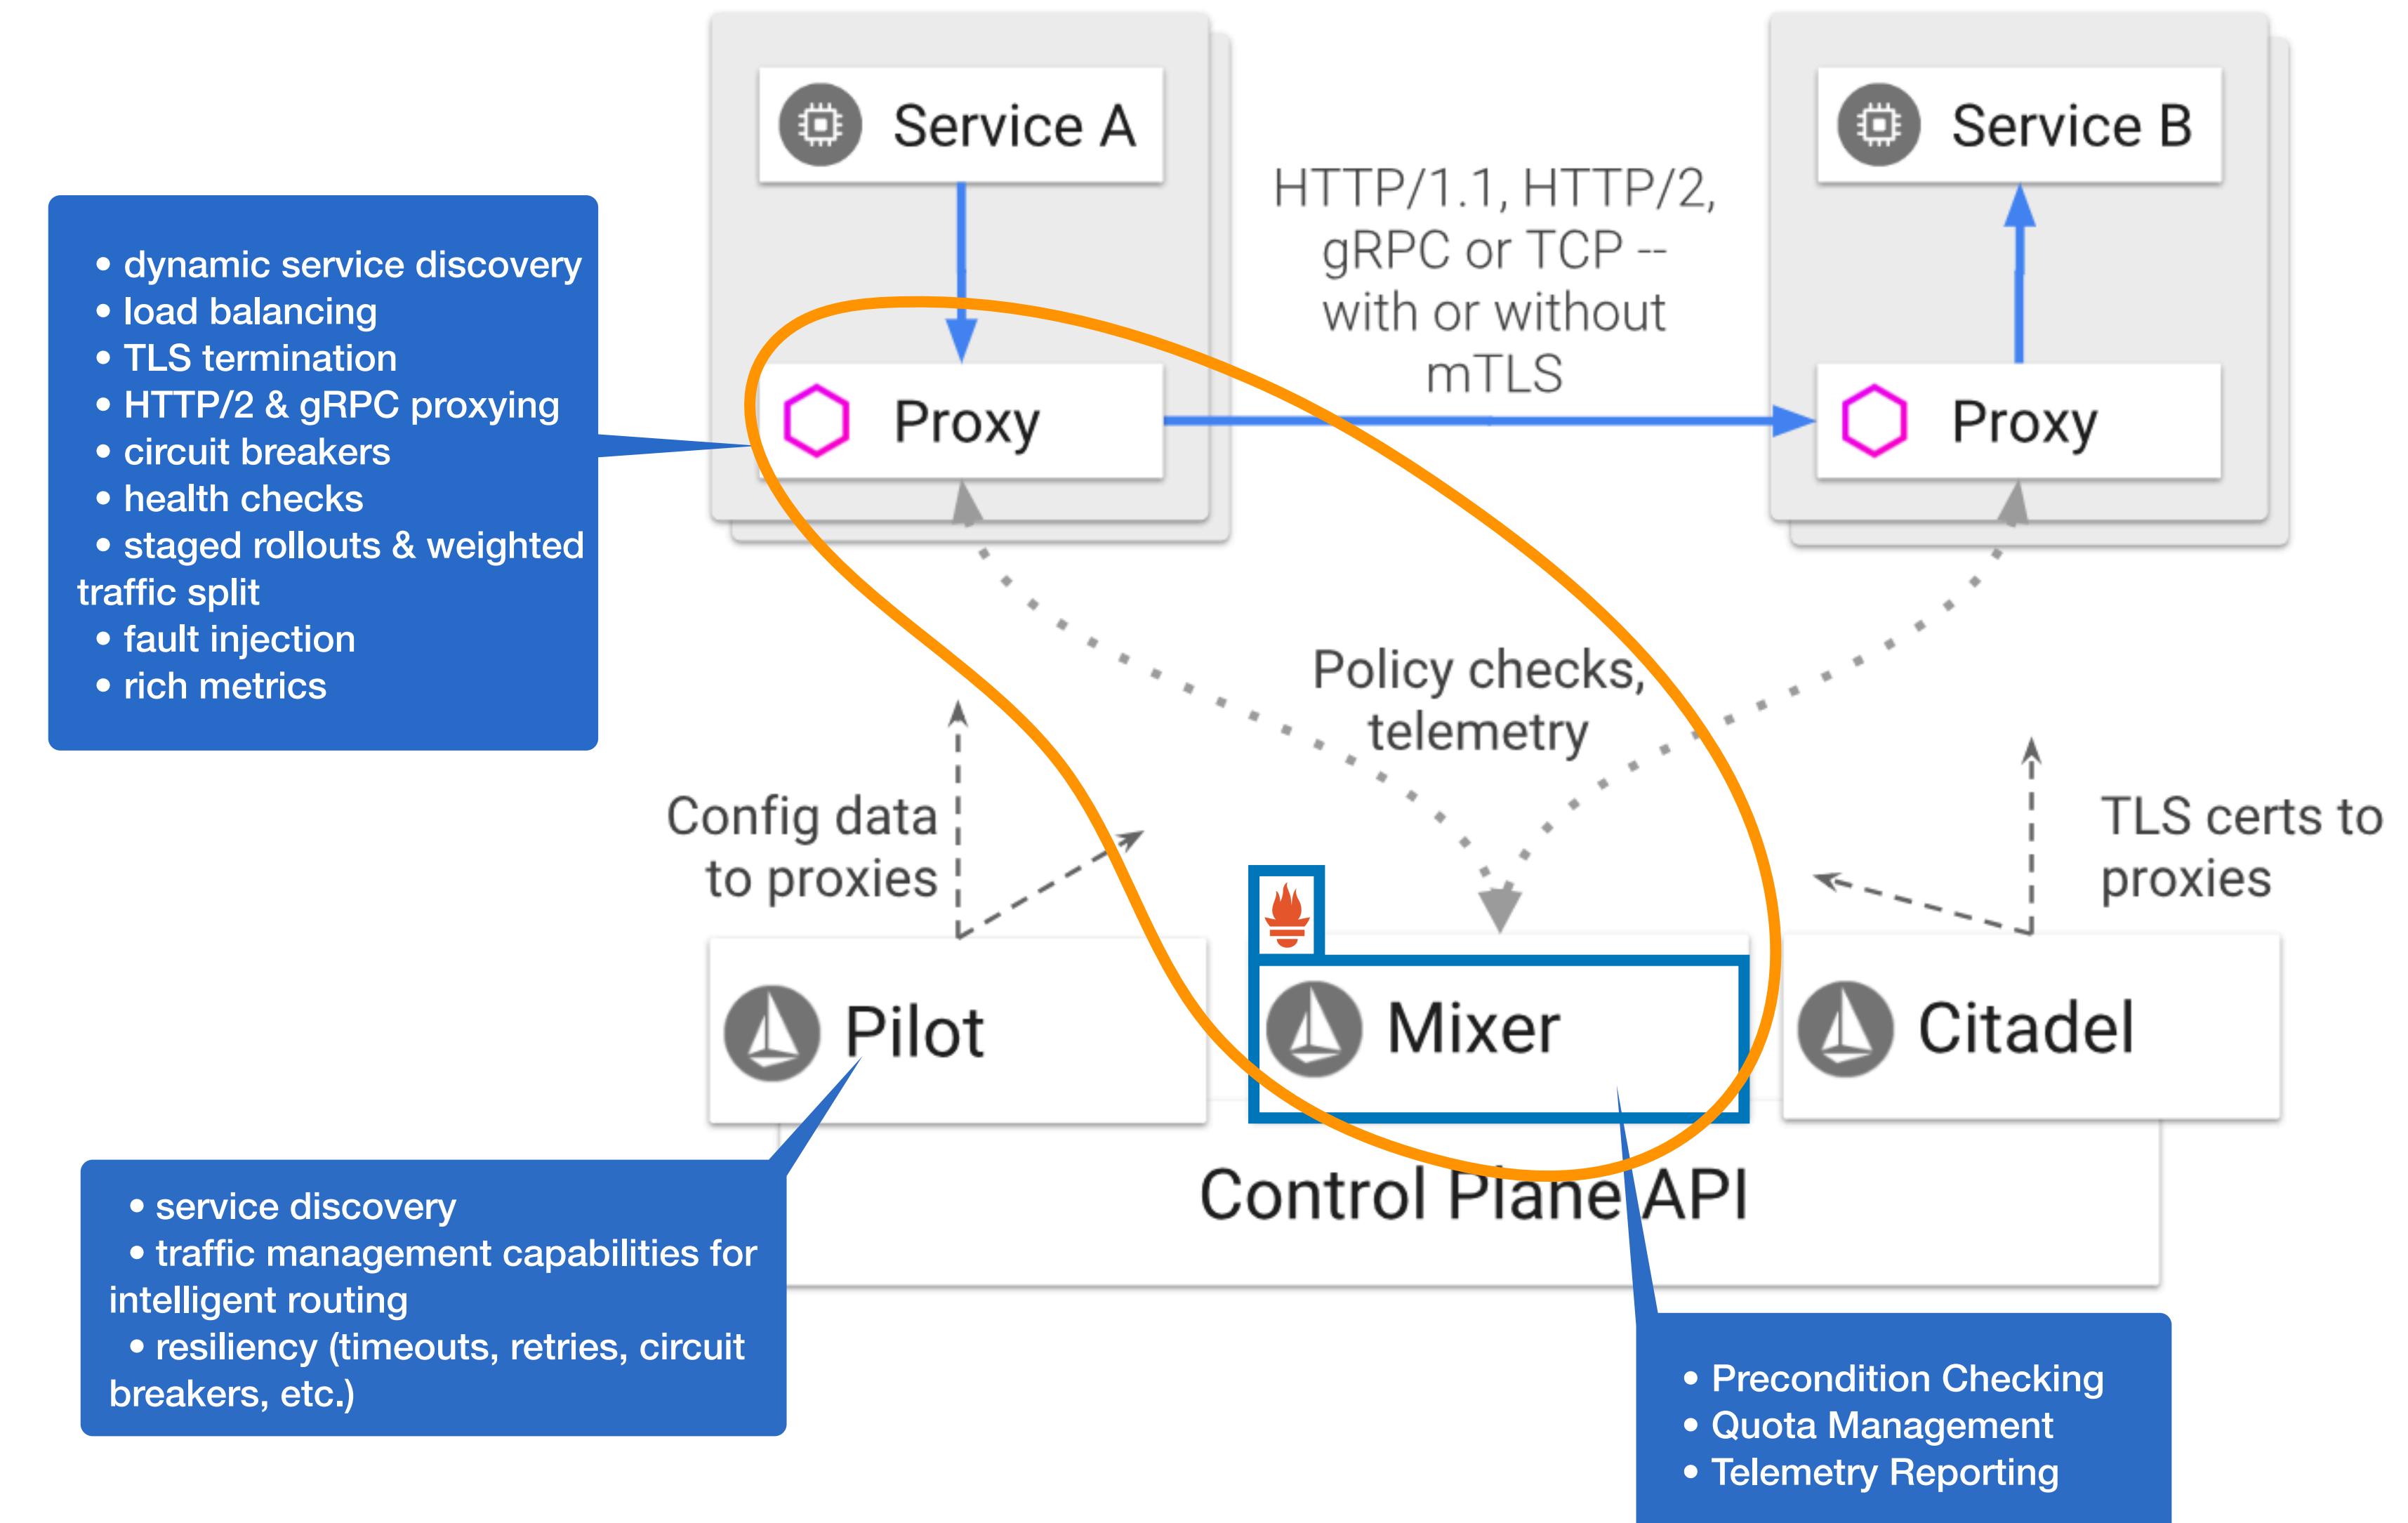


Istio

- **Istio** lets you connect, secure, control, and observe services
- Heavy investments from industry-leading companies such as **Google, IBM, and Lyft**
- Release **1.1**
- Active community

Core Features

- **Telemetry**
- **Fine-grained traffic management**
- **Rich load balancing**
- **Circuit breaking, timeouts, retries**
- **Request Tracing**
- Fault Injection
- Rate limits
- Load shedding



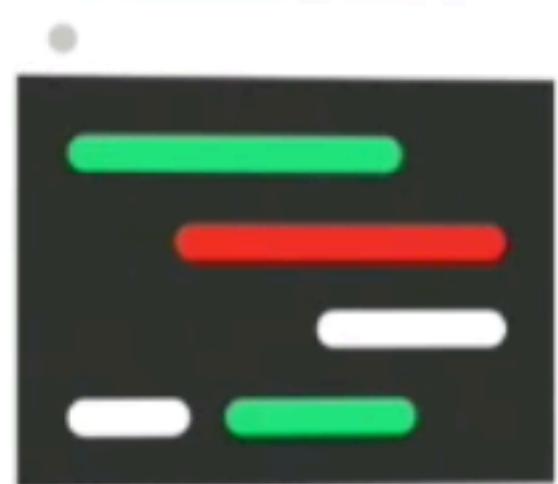
Monitoring vs. Observability

Pillars of Observability

Observability is a superset of monitoring and consists of four pillars that should help us reason about and respond to modern software systems



Alerts



Tracing



Logging



KubeCon



CloudNativeCon

China 2018

Build your own definition: <https://medium.com/@copyconstruct/monitoring-and-observability-8417d1952e1c>



Monitoring vs. Observability

Pillars of Observability

Observability is a superset of monitoring and consists of four pillars that should help us reason about and respond to modern software systems



Metrics



Tracing



Logging



Build your own definition: <https://medium.com/@copyconstruct/monitoring-and-observability-8417d1952e1c>



KubeCon



CloudNativeCon

China 2018



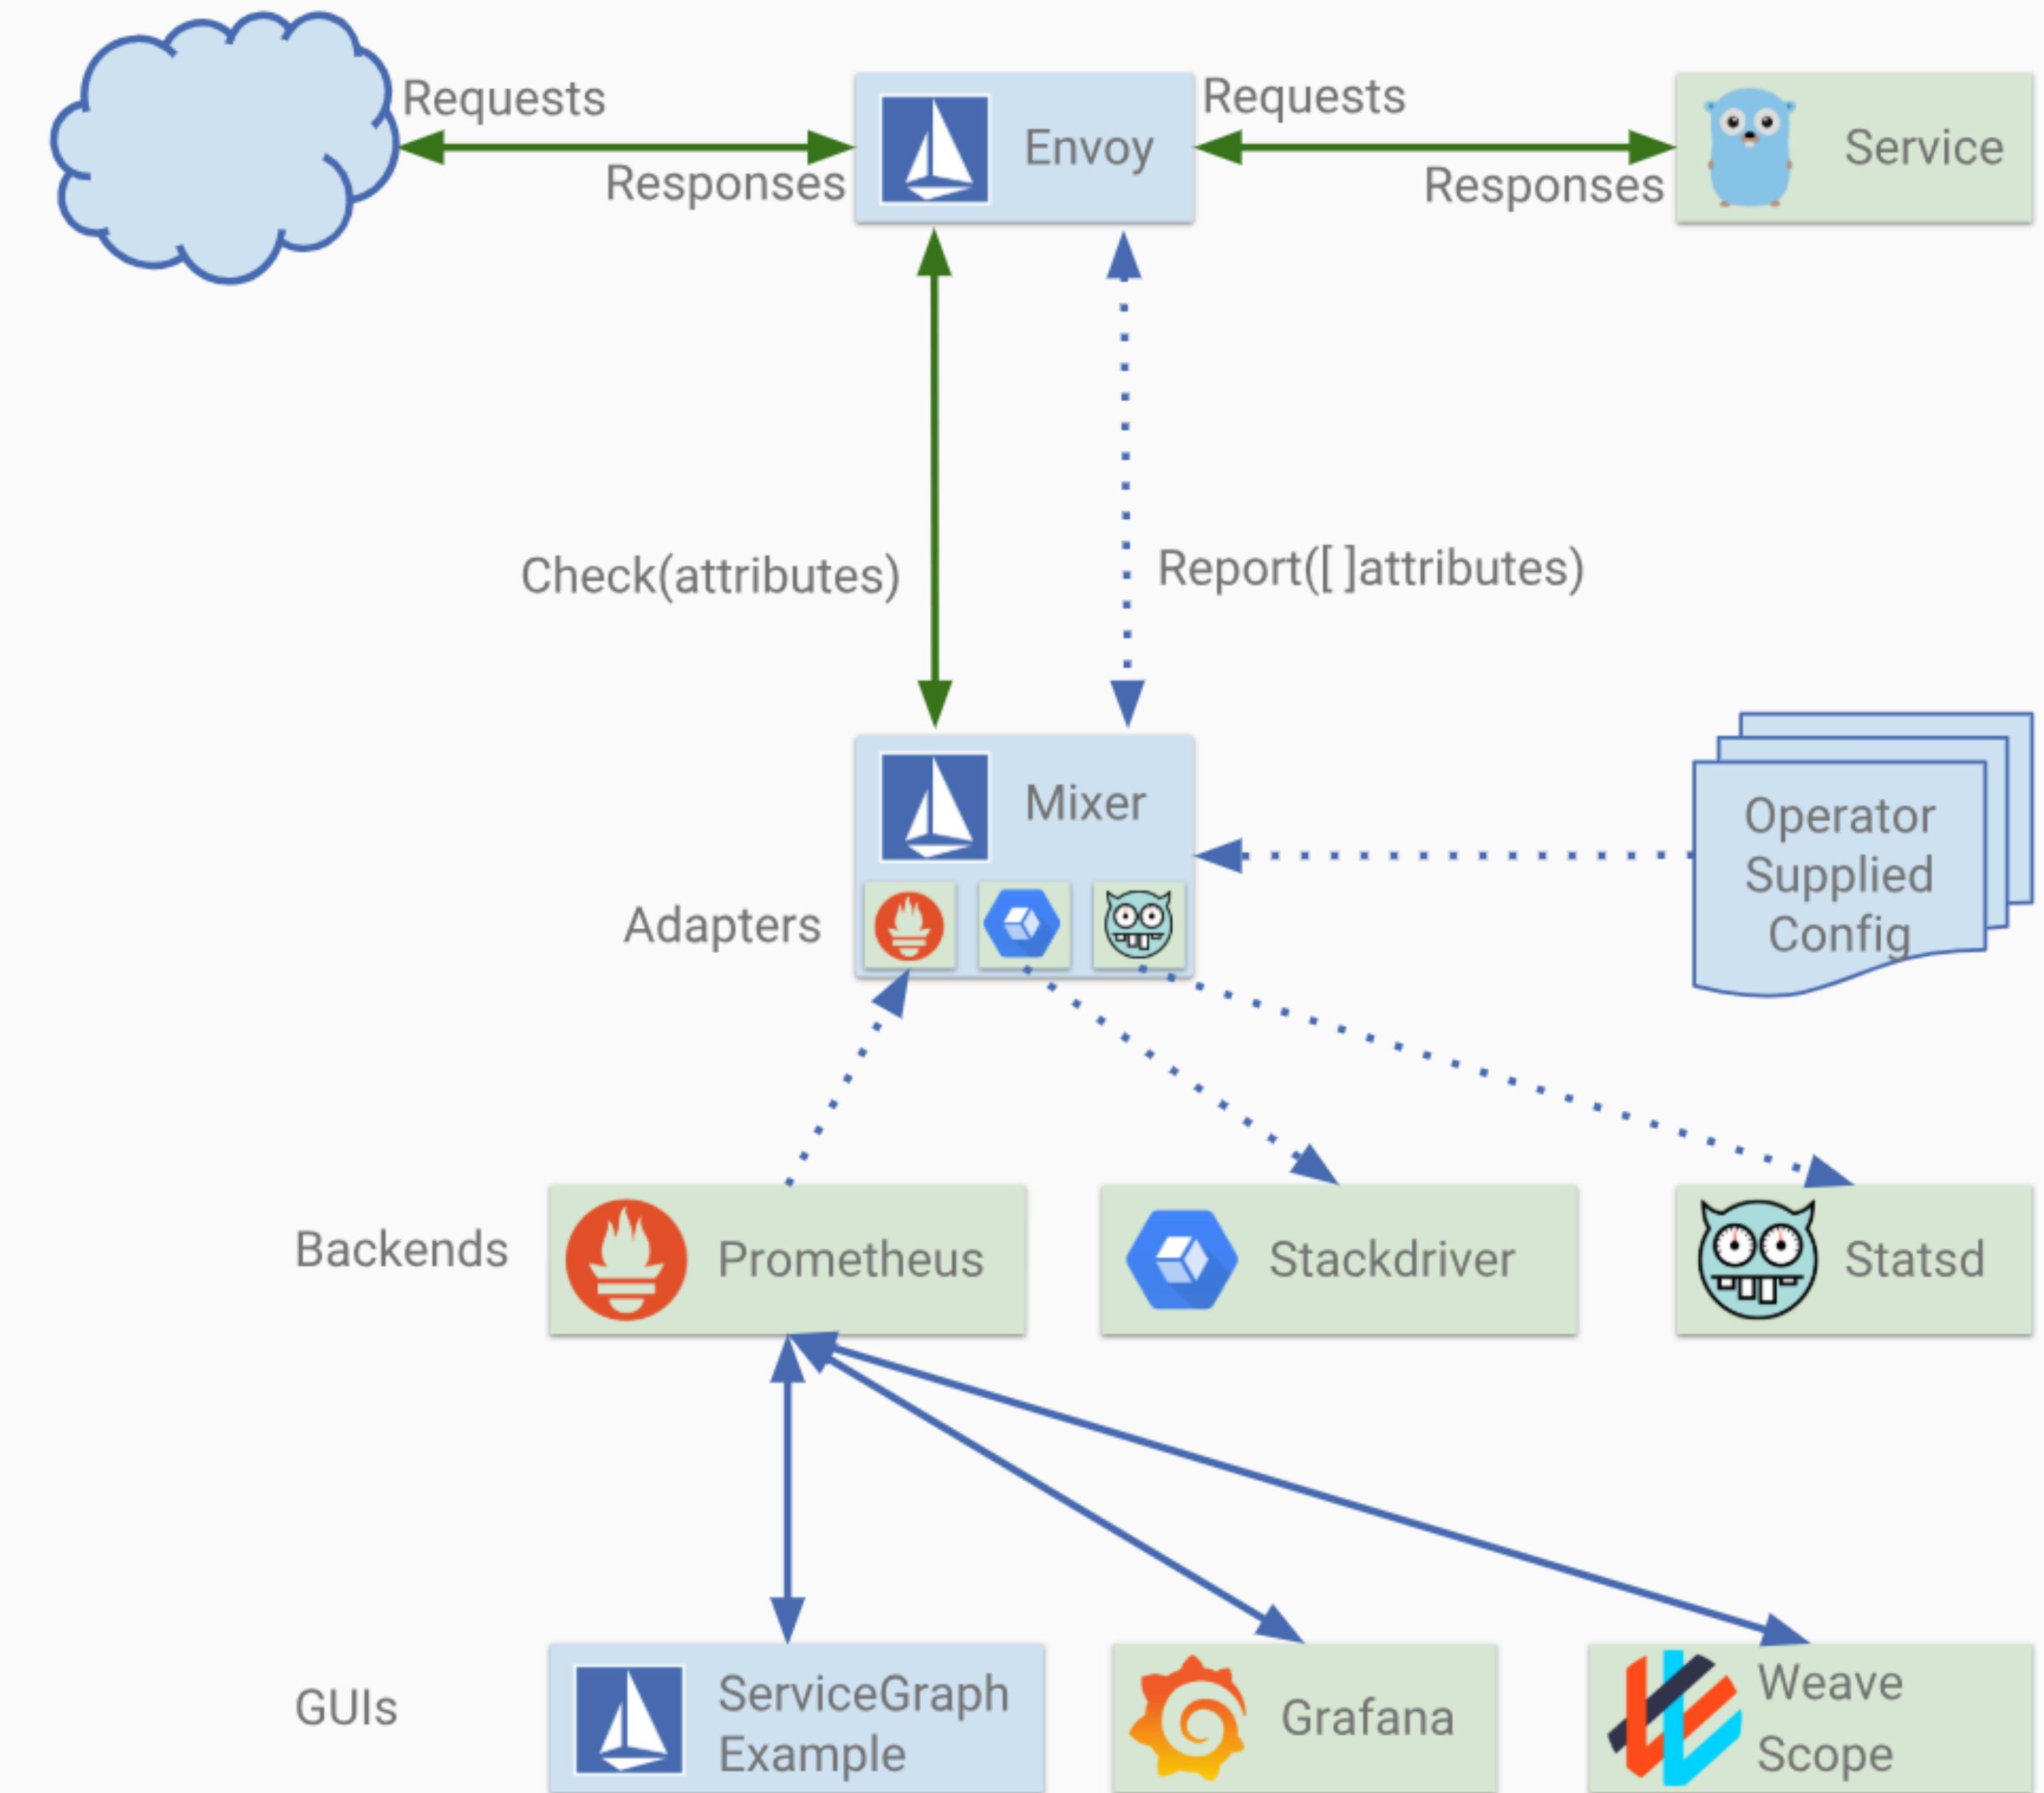
Telemetry Generation

Envoy sends information to Mixer about requests and responses in the form of *attributes* (typed key-value pairs).

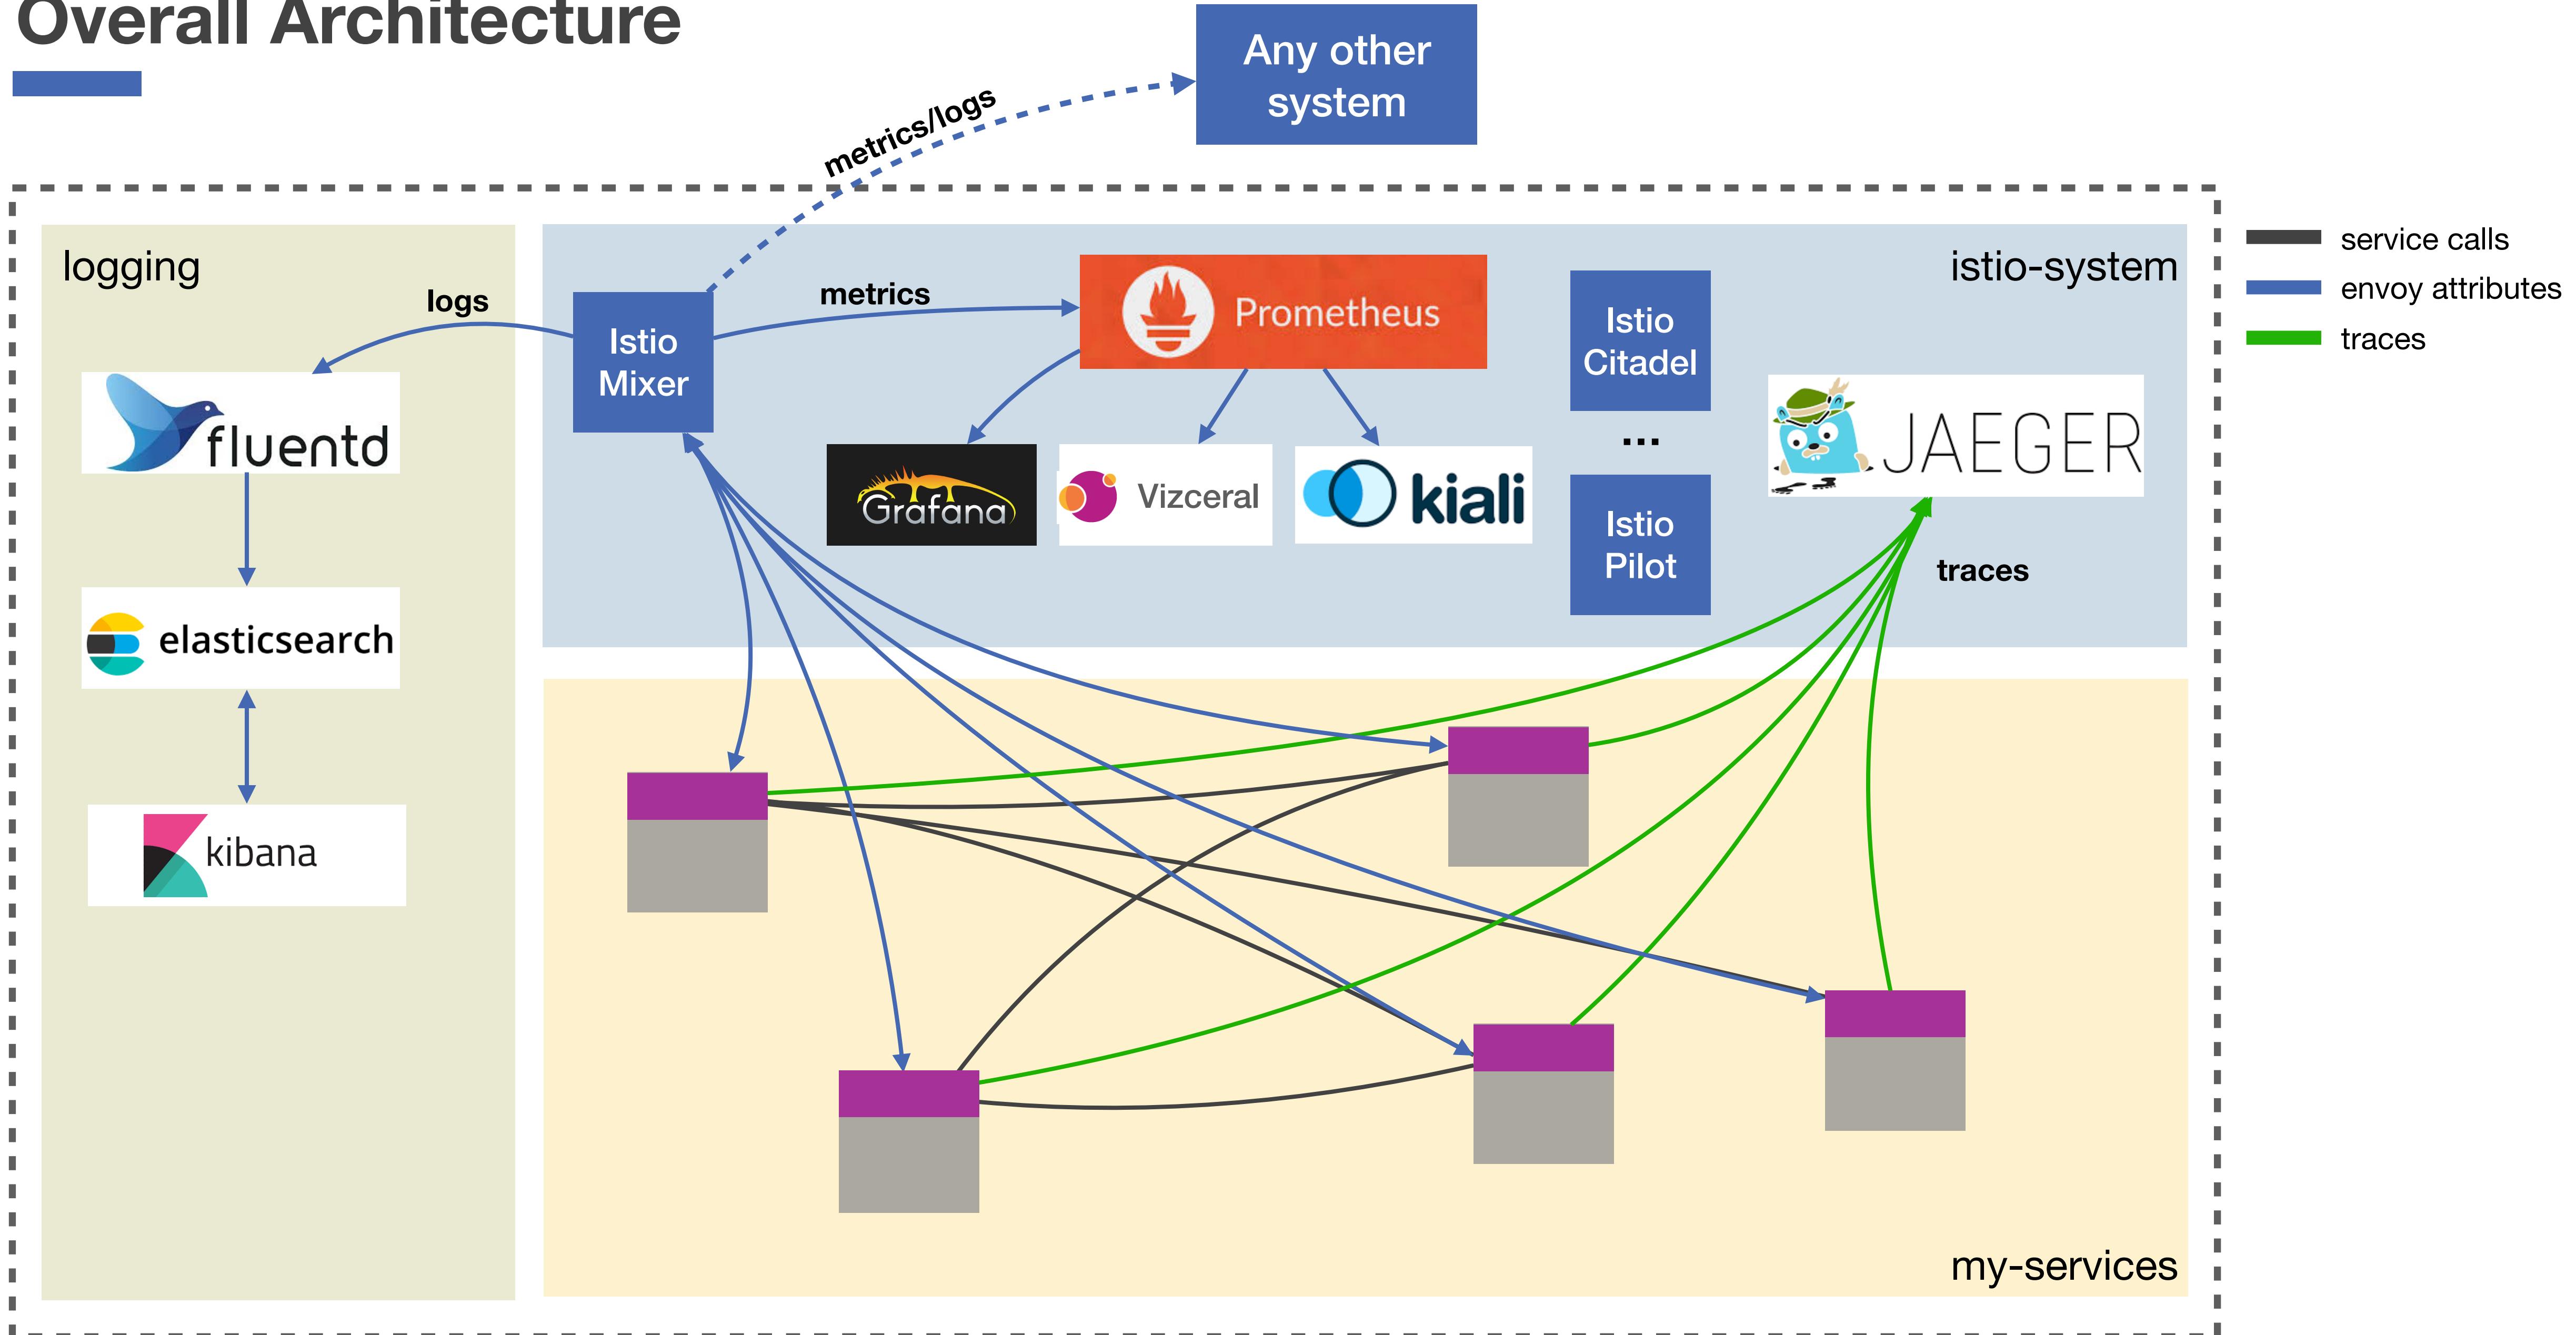
Mixer transforms attribute sets into structured values, according to operator-supplied configuration.

Mixer dispatches the derived values to a set of *adapters*, according to the operator-supplied configuration.

Adapters publish telemetry data to backend systems, making it available for further consumption and analysis.



Overall Architecture



Metrics

- By default, everything is in **Prometheus**
- **istio-mesh**, service mesh telemetry data generated by envoy
- **istio-telemetry**, metrics of telemetry itself
- Grafana for metrics visualization — Make full usage of Grafana
- Demo

```
root@qccloud:~/workspace# k get po -n istio-system -o wide
NAME                         READY   STATUS    RESTARTS   AGE     IP
grafana-85f97b7b6-pj68c       1/1     Running   0          23h    10.244.1.32
istio-citadel-65777464f-dz8pp  1/1     Running   0          23h    10.244.1.25
istio-cleanup-secrets-dxbn6  0/1     Completed  0          23h    10.244.0.38
istio-egressgateway-6497dfdfc-v4spf 1/1     Running   0          23h    10.244.0.40
istio-galley-5664cf4486-czx2z   1/1     Running   0          23h    10.244.0.39
istio-grafana-post-install-4tf9m 0/1     Completed  0          23h    10.244.1.31
istio-ingressgateway-787bf4cb5b-tmtkn 1/1     Running   0          23h    10.244.0.41
istio-pilot-54b6f64977-vbf6v    2/2     Running   0          23h    10.244.0.44
istio-policy-659559675c-trx8c   2/2     Running   0          23h    10.244.0.43
istio-sidecar-injector-7ddfb55469-99jkj 1/1     Running   0          23h    10.244.1.35
istio-statsd-prom-bridge-67bbcc746c-fsvzq 1/1     Running   0          23h    10.244.1.34
istio-telemetry-7d4f794c8d-q7vct    2/2     Running   0          23h    10.244.2.24
istio-tracing-6445d6dbbf-wtqj6   1/1     Running   0          23h    10.244.2.26
kiali-b765dfc99-swbv7          1/1     Running   0          23h    10.244.0.42
prometheus-6967c997cf-r4nl2    1/1     Running   0          23h    10.244.1.33
```

Targets

All Unhealthy

envoy (1/1 up) [show more](#)

galley (1/1 up) [show more](#)

istio-mesh (1/1 up) [show less](#)

Endpoint	State
http://10.244.2.24:42422/metrics	UP

istio-policy (1/1 up) [show more](#)

istio-telemetry (1/1 up) [show less](#)

Endpoint	State
http://10.244.2.24:9093/metrics	UP

kubernetes-apiservers (1/1 up) [show more](#)

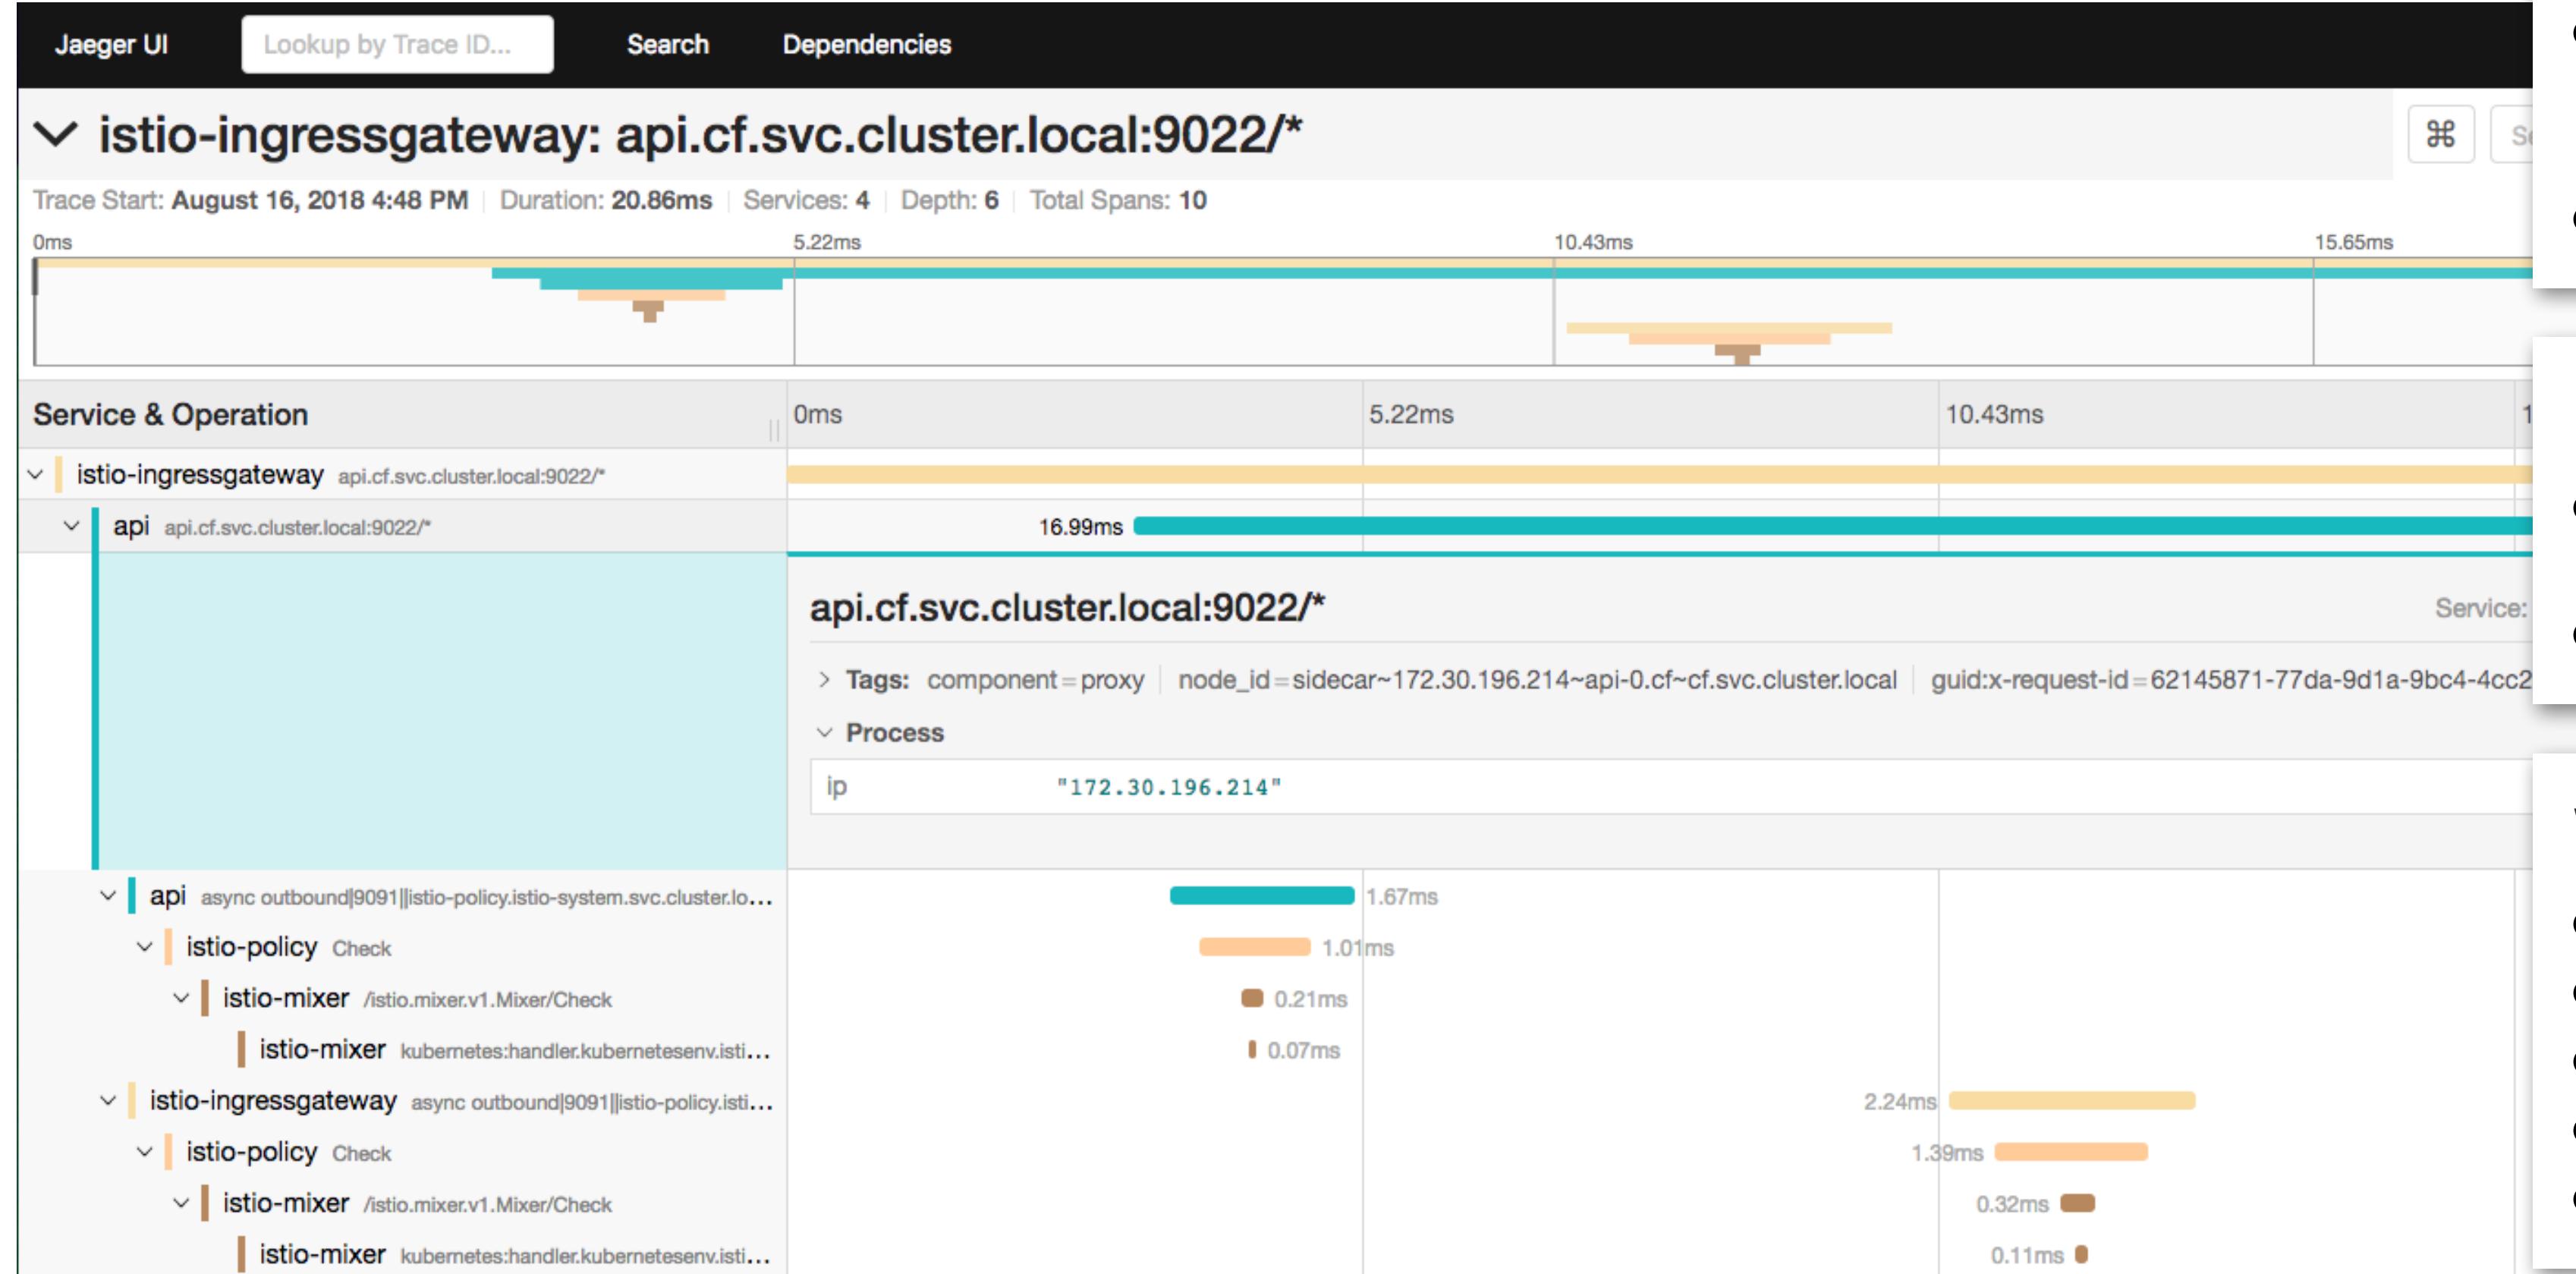
kubernetes-cadvisor (3/3 up) [show more](#)

kubernetes-nodes (3/3 up) [show more](#)

kubernetes-service-endpoints (3/3 up) [show more](#)

pilot (1/1 up) [show more](#)

Tracing – Jaeger



Problems that Jaeger addresses

distributed transaction monitoring

performance and latency optimization

root cause analysis

service dependency analysis

distributed context propagation

What is Jaeger

- CNCF project, native support for OpenTracing
- Find service upstream/downstream dependencies by checking opentracing standard headers
- Supports Istio by default

How we want to use it – scenarios

- Find out where consumes most of the time in request chain – Bottle neck identification
- Trouble shooting for specific requests

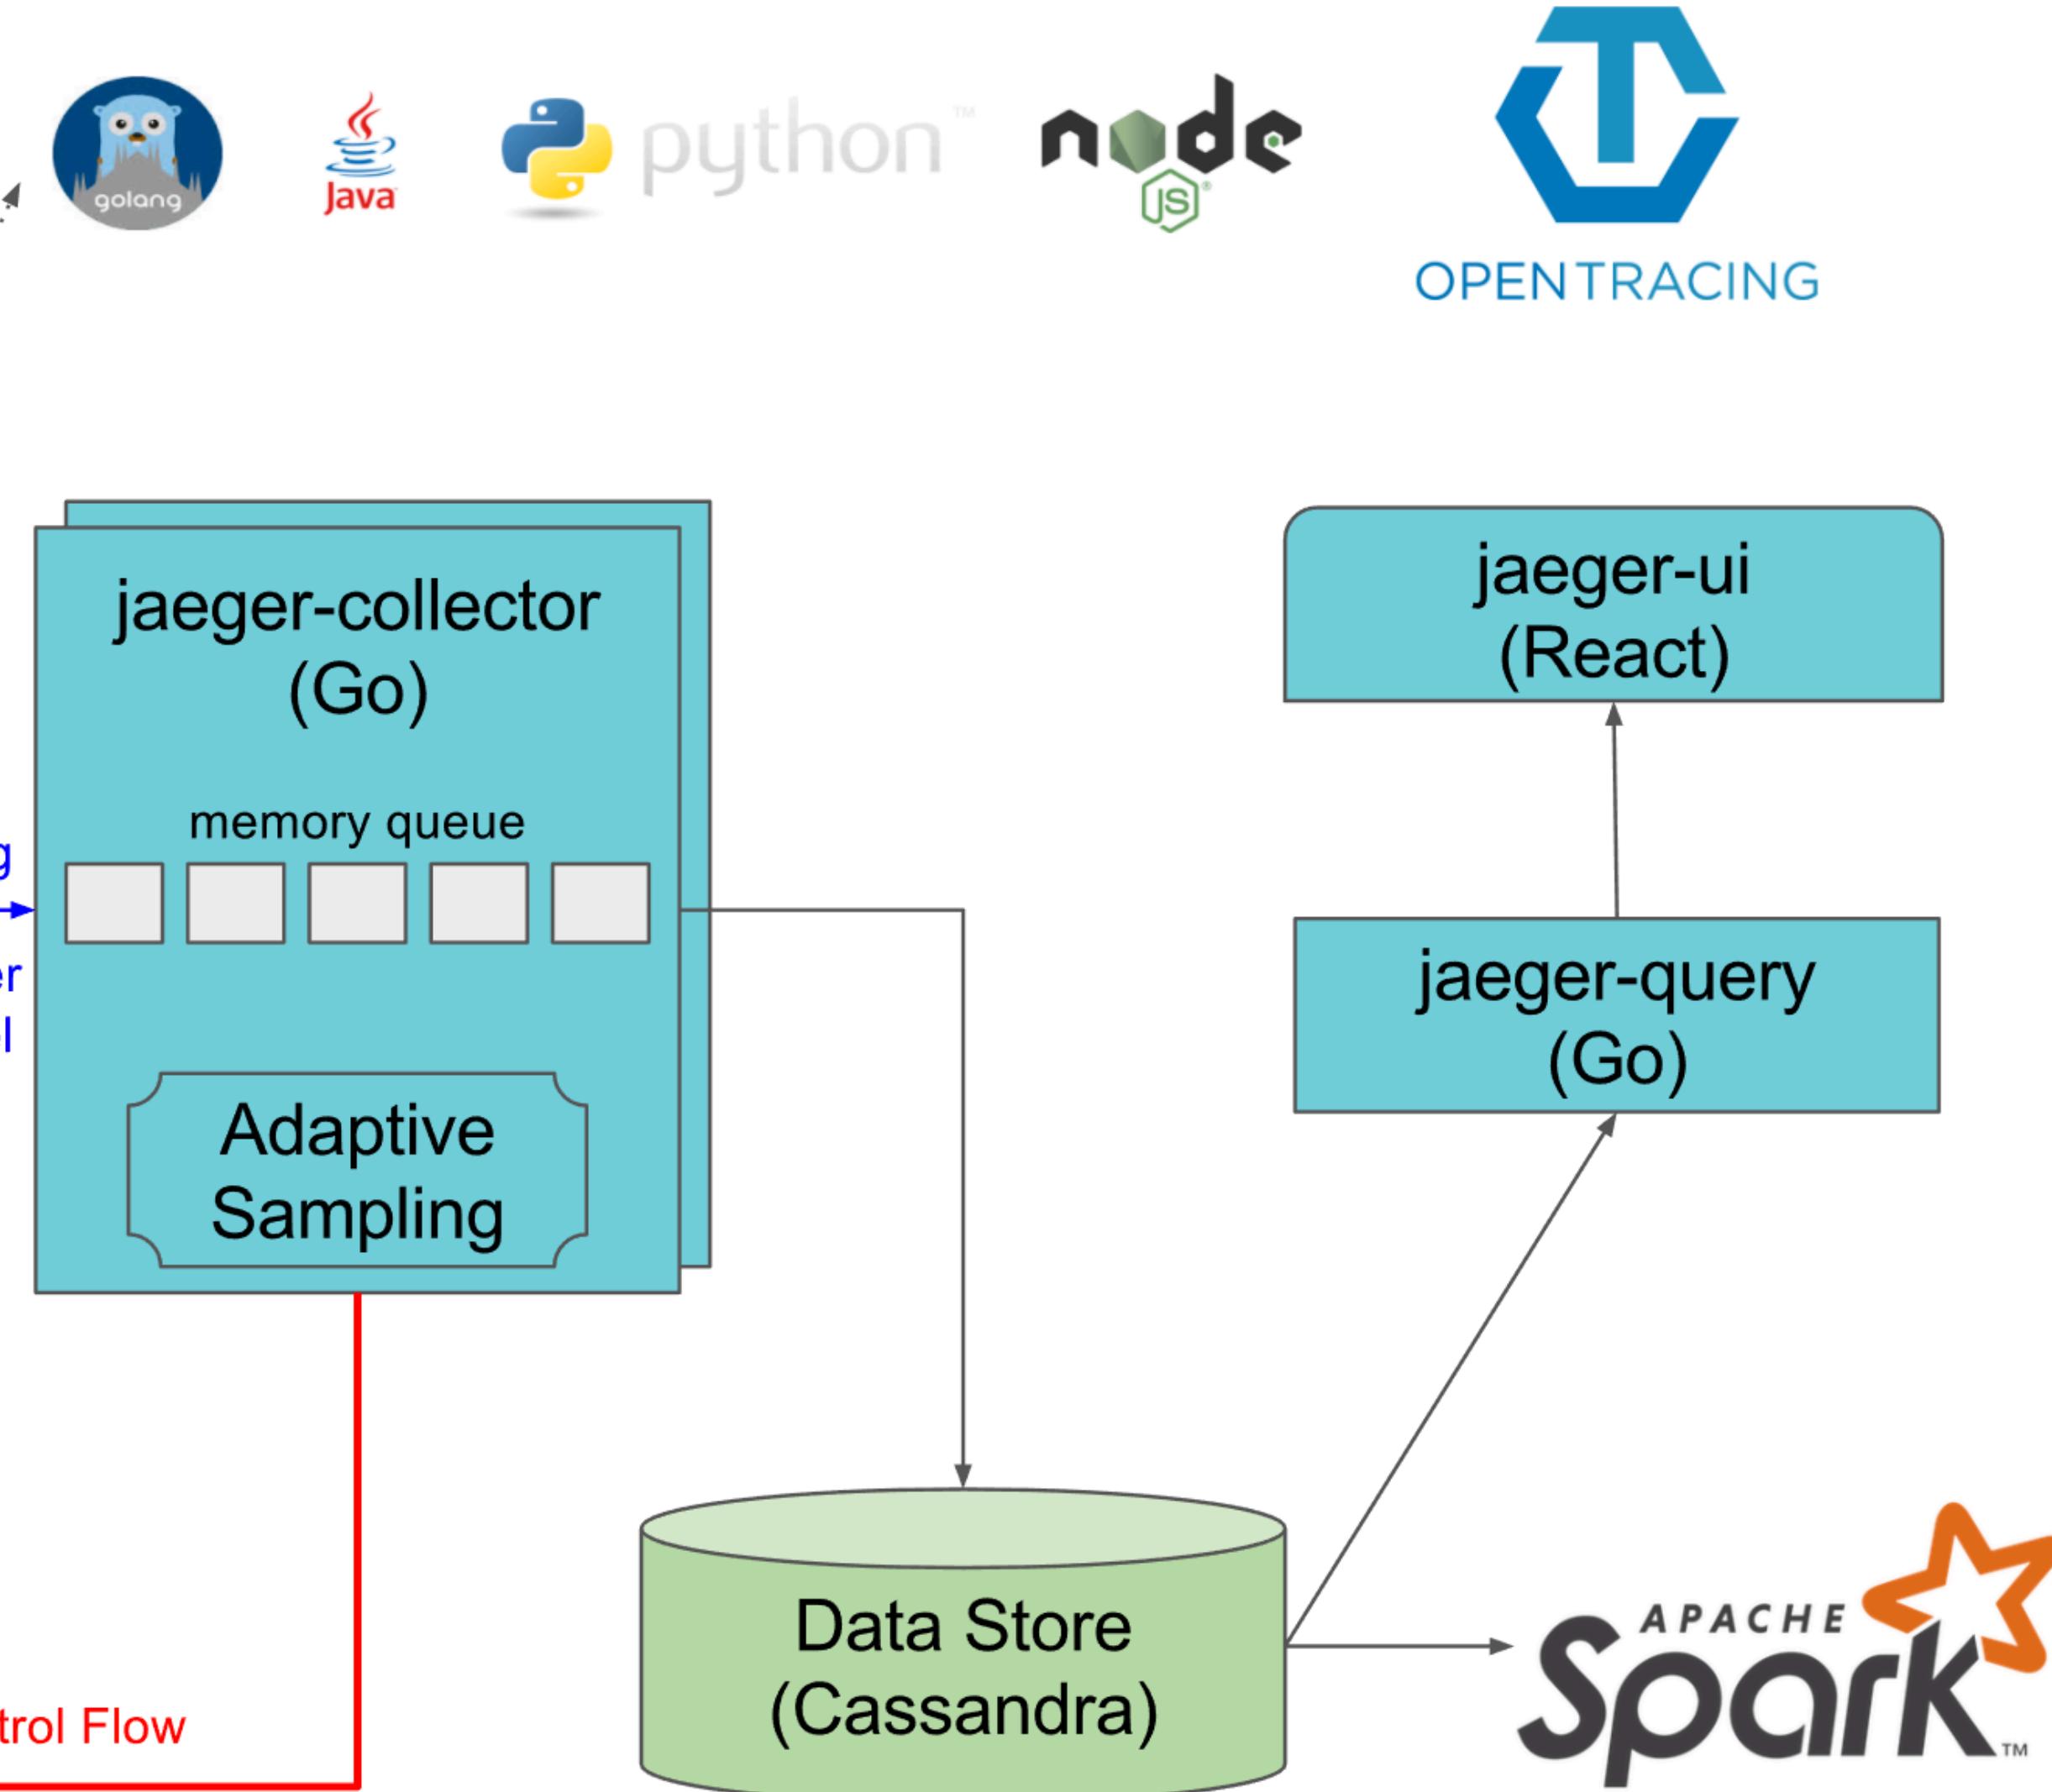
What is SPAN in OpenTracing

- An operation name
- A start timestamp and finish timestamp
- A set of key:value span Tags
- A set of key:value span Logs
- A SpanContext

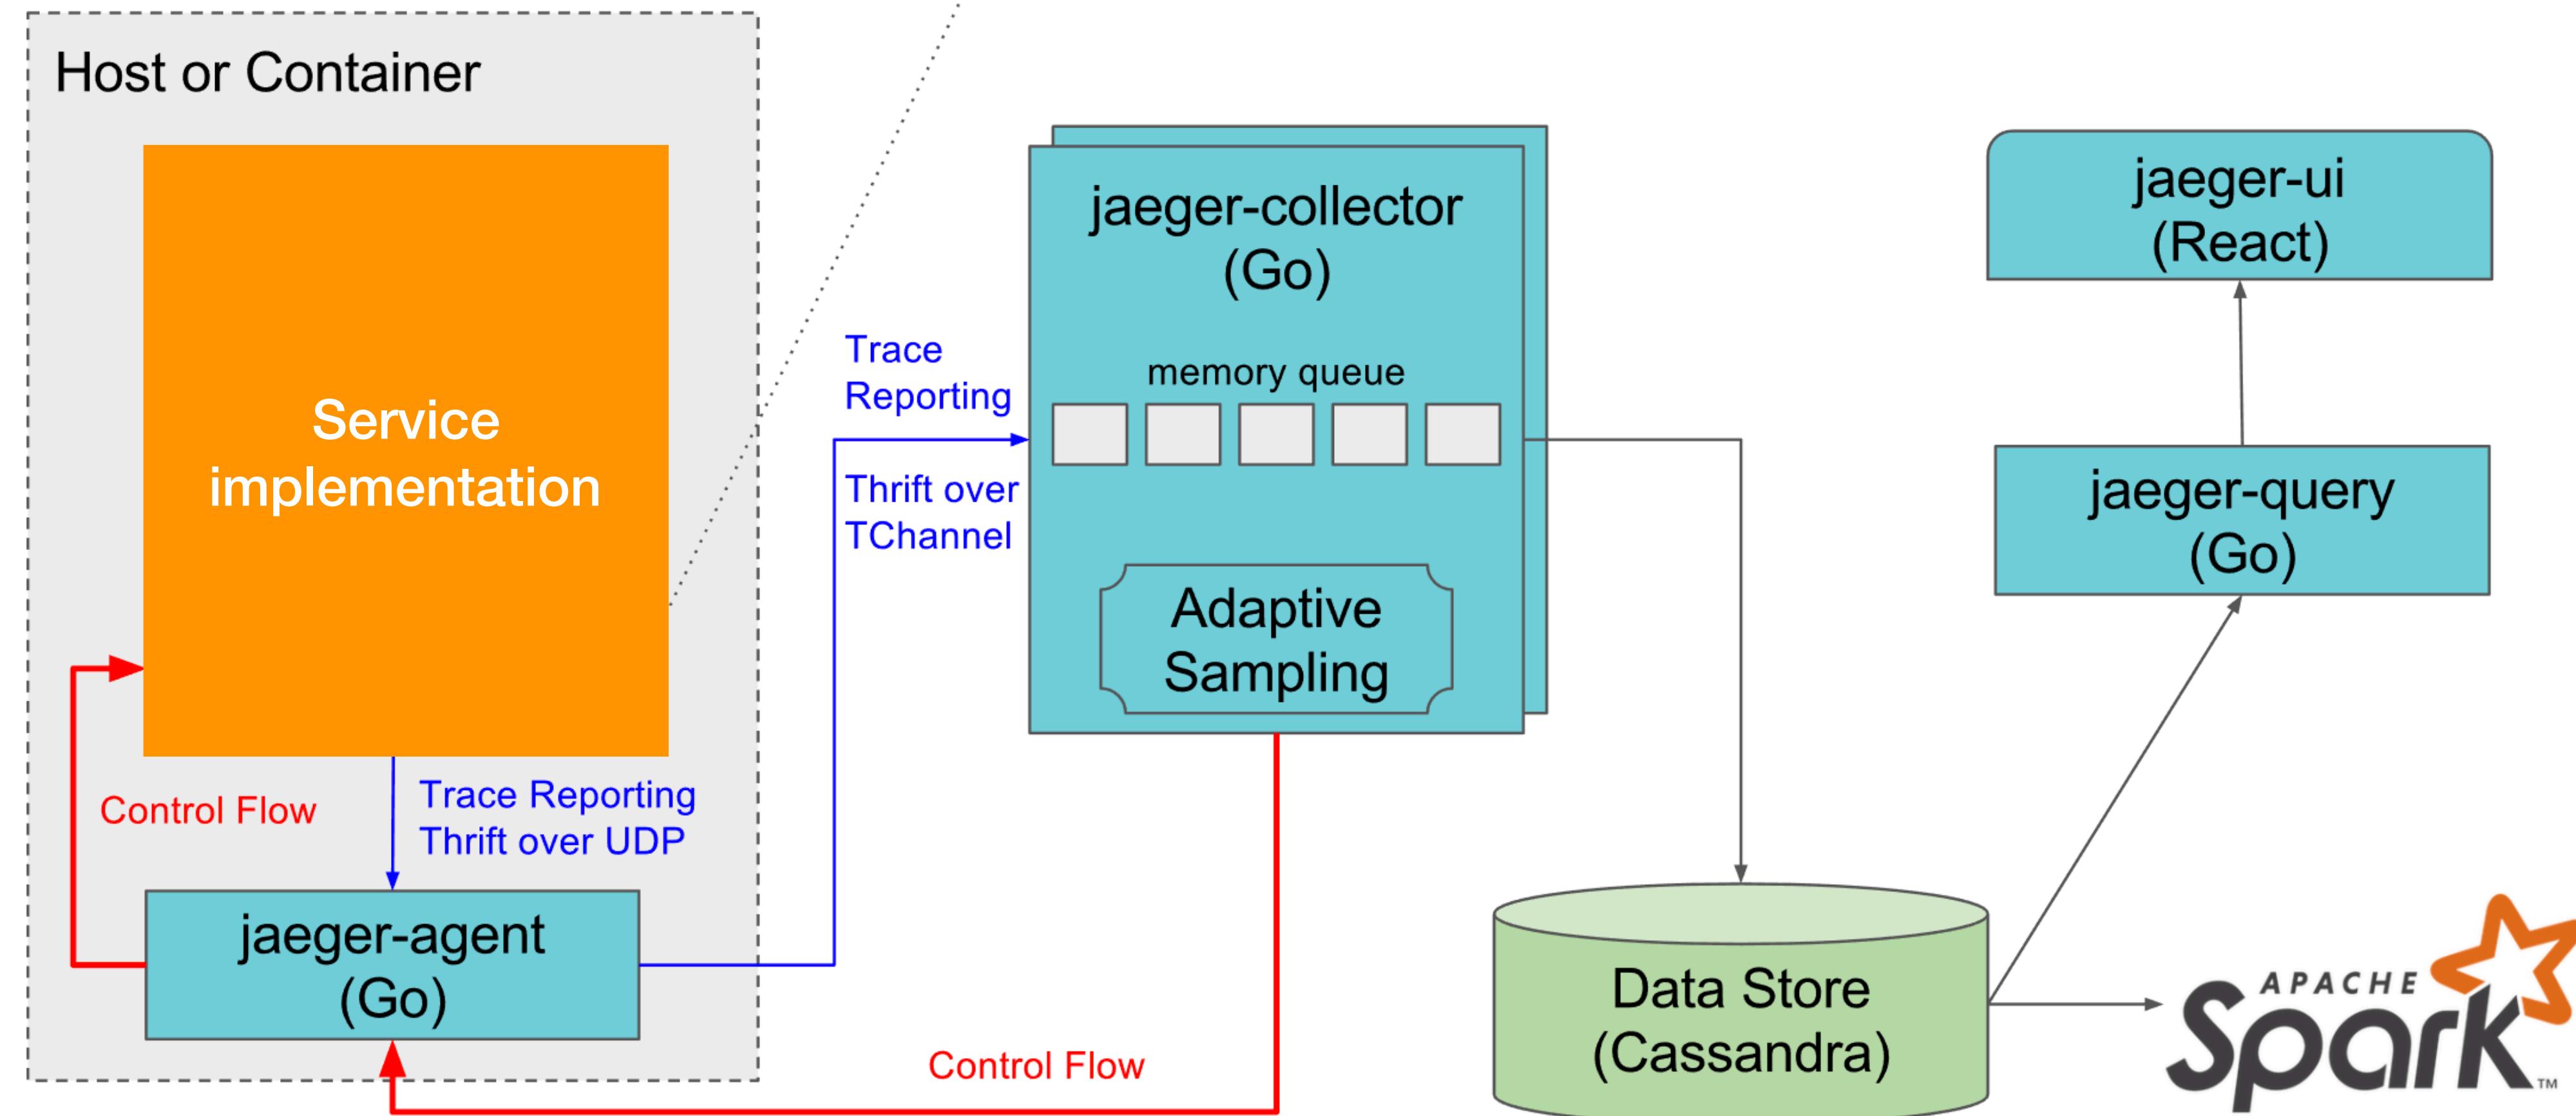
What is SPAN in Jaeger

A **span** represents a logical unit of work in Jaeger that has an operation name, the start time of the operation, and the duration. Spans may be nested and ordered to model causal relationships.

Tracing – Jaeger



Tracing – Jaeger



Tracing

- Tracing is a **super powerful tool** to help with trouble shooting and performance analysis in real world operation
- Inspired by Dapper and Zipkin, initiated by Uber
- Implemented by following OpenTracing <https://opentracing.io/docs/overview/>
- In Istio, spans will be generated by each envoy sidecar and send to Jaeger by default — separated spans
- To link spans together, code change is required — “context  propagation”

bookinfo — details.rb

```
def get_forward_headers(request)
  headers = {}
  incoming_headers = [ 'x-request-id',
    'x-b3-traceid',
    'x-b3-spanid',
    'x-b3-parentspanid',
    'x-b3-sampled',
    'x-b3-flags',
    'x-ot-span-context' ]
  request.each do |header, value|
    if incoming_headers.include? header then
      headers[header] = value
    end
  end
  return headers
end

server.start
```

Tracing

- Tracing is a **super powerful tool** to help with trouble shooting and performance analysis in real world operation
- Inspired by Dapper and Zipkin, initiated by Uber
- Implemented by following OpenTracing <https://opentracing.io/docs/overview/>
- In Istio, spans will be generated by each envoy sidecar and send to Jaeger by default — separated spans
- To link spans together, code change is required — “context 

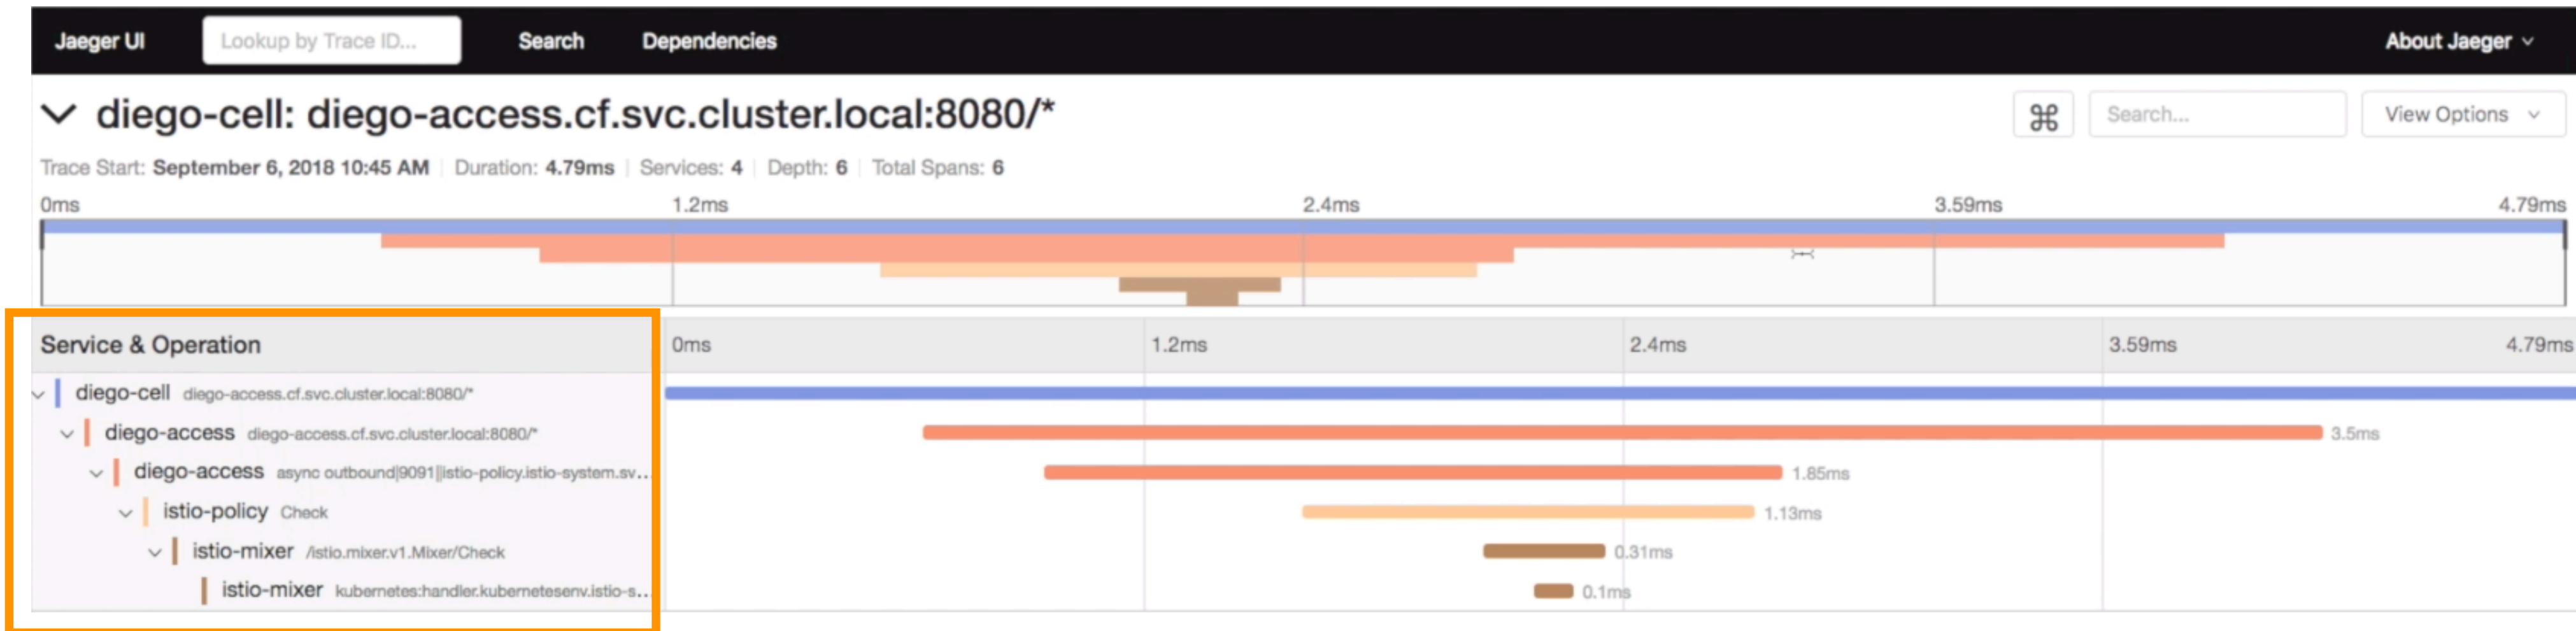
```
graph LR; A[envoy A] --> B[envoy B]; B --> C[envoy C]; C --> Out[ ]; A -- "A->B" --> B; B -- "B->C" --> C; C -- "B->C, C->*" --> Out;
```

bookinfo – details.rb

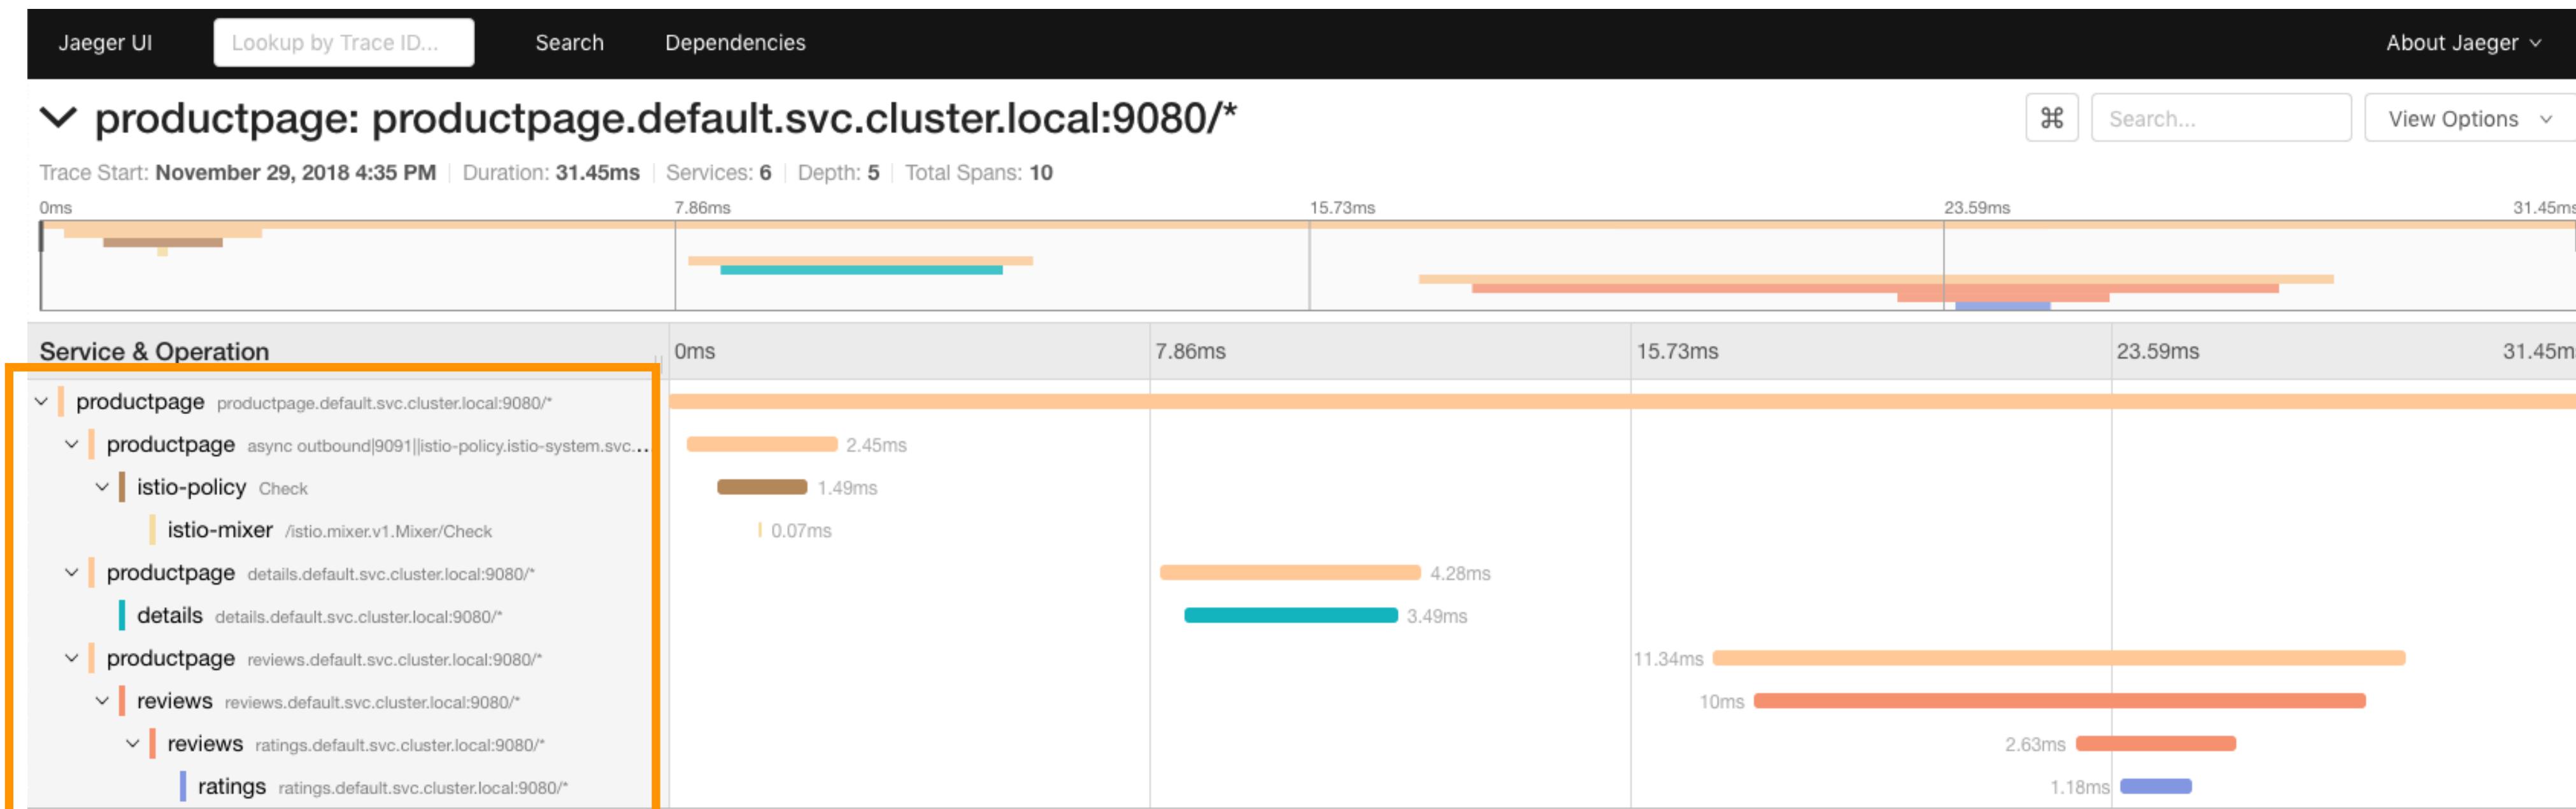
```
def get_forward_headers(request)
  headers = {}
  incoming_headers = [ 'x-request-id',
    'x-b3-traceid',
    'x-b3-spanid',
    'x-b3-parentspanid',
    'x-b3-sampled',
    'x-b3-flags',
    'x-ot-span-context' ]
  request.each do |header, value|
    if incoming_headers.include? header then
      headers[header] = value
    end
  end
  return headers
end

server.start
```

Tracing



Separated



Linked

Tracing - Production Readiness

-  Current steps described in [official document "Distributed Tracing"](#) is not a production ready solution – “ALL_IN_ONE” mode, and **SPAN_STORAGE_TYPEIS is in memory**
- Change backend storage to [Cassandra](#) or [Elasticsearch](#) or [Kafka](#) for data persistent
- HPA for jaeger-collector

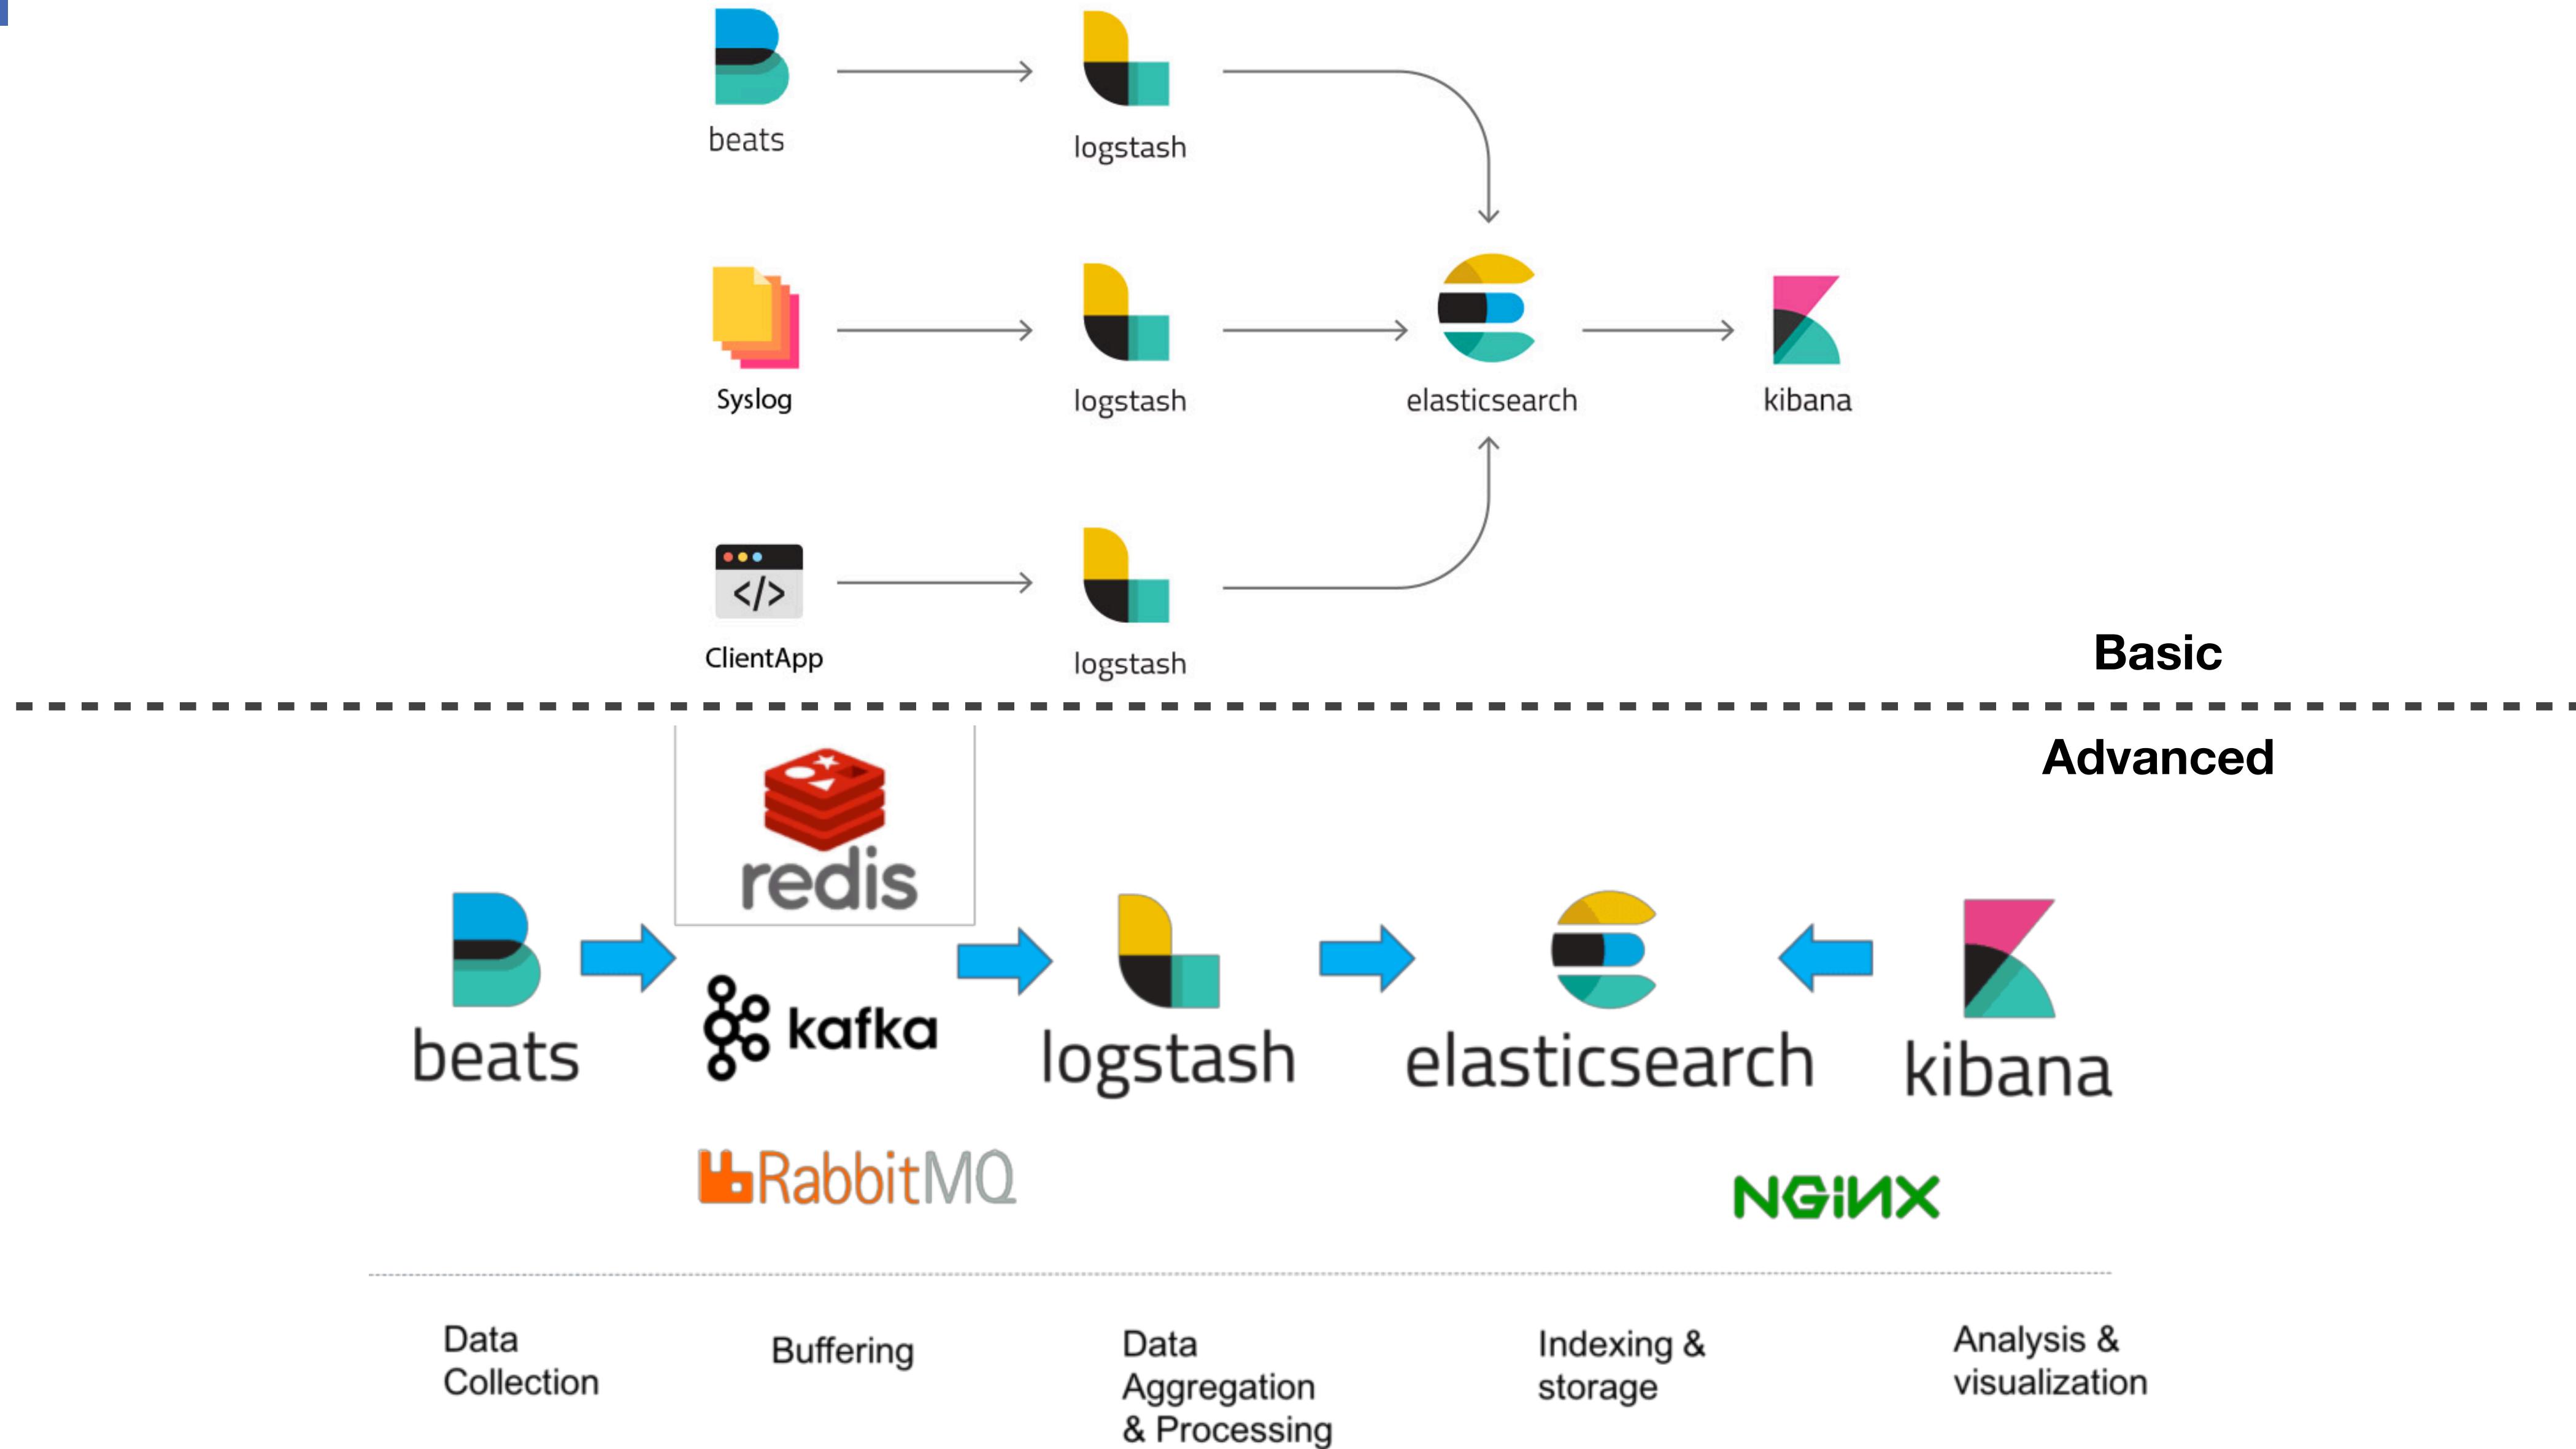
```
cmd [master••] k get po -n istio-system
NAME                                         READY   STATUS
grafana-546d9997bb-fq5fg                   1/1     Running
istio-citadel-6955bc9cb7-5r47x              1/1     Running
istio-egressgateway-7dc5cbbc56-w7tvb       1/1     Running
istio-galley-545b6b8f5b-mrnch               1/1     Running
istio-ingressgateway-7958d776b5-kvjqvf      1/1     Running
istio-pilot-56bfdbffff-xfd2w                2/2     Running
istio-policy-5c689f446f-g948f                2/2     Running
istio-sidecar-injector-99b476b7b-zdlrf       1/1     Running
istio-telemetry-55d68b5dfb-d65v9             2/2     Running
istio-tracing-6445d6dbbf-8zzrn             1/1     Running
kiali-ddf8fb8b-qvmgt                         1/1     Running
prometheus-65d6f6b6c-59jq7                  1/1     Running
servicegraph-57c8cbc56f-trvmn                1/1     Running

►cmd [master••] k get svc -n istio-system
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP
grafana       NodePort    10.111.216.82 <none>
istio-citadel ClusterIP  10.102.88.145  <none>
istio-egressgateway ClusterIP 10.109.197.155 <none>
istio-galley   ClusterIP  10.96.50.109  <none>
istio-ingressgateway LoadBalancer 10.108.209.24 <none>
istio-pilot    ClusterIP  10.99.181.25  <none>
istio-policy   ClusterIP  10.109.129.152 <none>
istio-sidecar-injector ClusterIP 10.102.110.132 <none>
istio-telemetry NodePort    10.96.87.57   <none>
jaeger-agent  ClusterIP  None          <none>
jaeger-collector ClusterIP  10.102.108.54 <none>
jaeger-query  ClusterIP  10.98.62.81   <none>
kiali         NodePort    10.98.86.230 <none>
prometheus    NodePort    10.97.65.25   <none>
servicegraph   ClusterIP  10.105.130.215 <none>
tracing        ClusterIP  10.105.44.158 <none>
zipkin         ClusterIP  10.109.202.116 <none>

►cmd [master••]
```

```
spec:
  containers:
    - name: jaeger-agent
      image: jaegertracing/jaeger-agent
      imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 5778
          protocol: TCP
        - containerPort: 5775
          protocol: UDP
        - containerPort: 6831
          protocol: UDP
        - containerPort: 6832
          protocol: UDP
    - name: jaeger-query
      image: jaegertracing/jaeger-query
      imagePullPolicy: IfNotPresent
      args:
        - "--es.server-urls=http://k8s-support-cluster2.fyre.ibm.com:9200"
      ports:
        - containerPort: 16686
          protocol: TCP
        - containerPort: 16687
          protocol: TCP
      env:
        - name: SPAN_STORAGE_TYPE
          value: "elasticsearch"
    - name: jaeger-collector
      image: jaegertracing/jaeger-collector
      imagePullPolicy: IfNotPresent
      args:
        - "--es.server-urls=http://k8s-support-cluster2.fyre.ibm.com:9200"
        - "--es.num-shards=1"
      ports:
        - containerPort: 14269
          protocol: TCP
        - containerPort: 14268
          protocol: TCP
        - containerPort: 14267
          protocol: TCP
        - containerPort: 9411
          protocol: TCP
      env:
        - name: COLLECTOR_ZIPKIN_HTTP_PORT
          value: "9411"
        - name: SPAN_STORAGE_TYPE
          value: "elasticsearch"
```

Logging – in the old days...

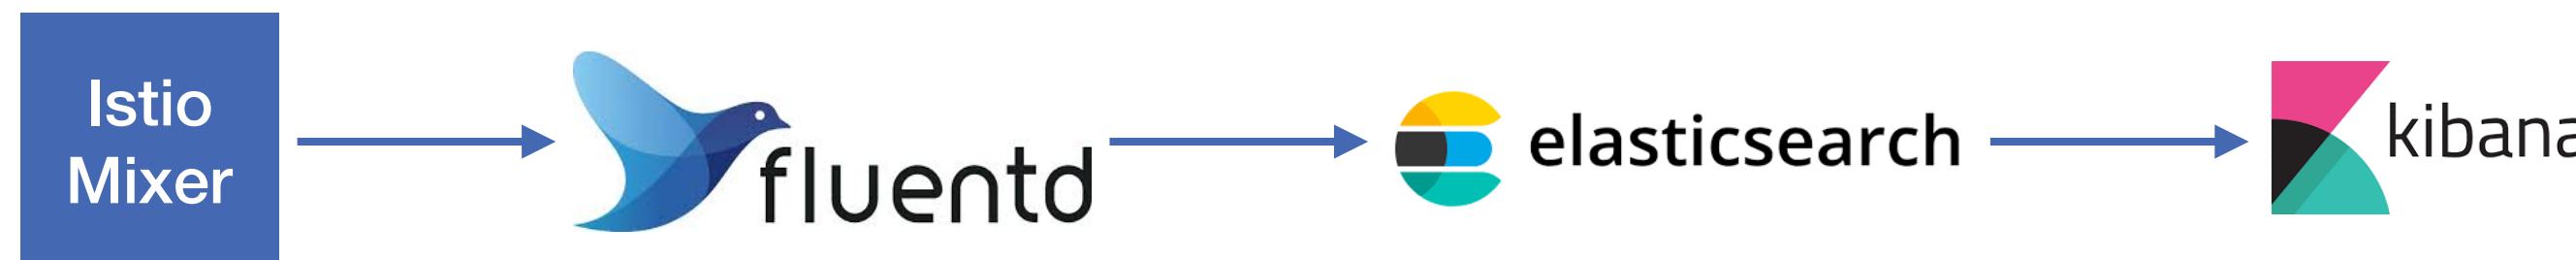


Logging



What is fluentd

- CNCF project, work with Elasticsearch and Kibana well, and also Kubernetes
- Light weight, very little foot print
- Rich plugins and community support



fluentd config from `logging-stack.yaml`

```
# Fluentd ConfigMap, contains config files.
kind: ConfigMap
apiVersion: v1
data:
  forward.input.conf: |-
    # Takes the messages sent over TCP
    <source>
      type forward
    </source>
  output.conf: |-
    <match **>
      type elasticsearch
      log_level info
      include_tag_key true
      host elasticsearch
      port 9200
      logstash_format true
      # Set the chunk limits.
      buffer_chunk_limit 2M
      buffer_queue_limit 8
      flush_interval 5s
      # Never wait longer than 5 minutes between
      retries.
      max_retry_wait 30
      # Disable the limit on the number of retries
      (retry forever).
      disable_retry_limit
      # Use multiple threads for processing.
      num_threads 2
    </match>
  metadata:
    name: fluentd-es-config
    namespace: logging
```

Logging

-  Current steps described in [official document "Logging with Fluentd"](#) is not a production ready solution
- “logging” in Istio, seems to be another way of collecting request specific data
- What about application logs?
- Investigation in progress: how to extend istio logging to include application logs as well
- Demo

```
# Configuration for logentry instances
apiVersion: "config.istio.io/v1alpha2"
kind: logentry
metadata:
  name: newlog
  namespace: istio-system
spec:
  severity: "info"
  timestamp: request.time
  variables:
    source: source.labels["app"] | source.workload.name | "unknown"
    user: source.user | "unknown"
    destination: destination.labels["app"] | destination.workload.name | "unknown"
    responseCode: response.code | 0
    responseSize: response.size | 0
    latency: response.duration | "0ms"
  monitored_resource_type: "UNSPECIFIED"
```

```
# Configuration for a fluentd handler
apiVersion: "config.istio.io/v1alpha2"
kind: fluentd
metadata:
  name: handler
  namespace: istio-system
spec:
  address: "fluentd-es.logging:24224"
```

```
# Rule to send logentry instances to the
# fluentd handler
apiVersion: "config.istio.io/v1alpha2"
kind: rule
metadata:
  name: newlogtofluentd
  namespace: istio-system
spec:
  match: "true" # match for all requests
  actions:
    - handler: handler.fluentd
      instances:
        - newlog.logentry
```

Kiali

- “Kiali is originated from the greek word κιάλι meaning monocular or spyglass.”
- Initiated by Redhat 
- <https://www.kiali.io>
- Officially included in Istio deployment since version v1.1: [Visualizing Your Mesh](#)
- Very active community and development sprints
- Demo

Features

A growing set of features

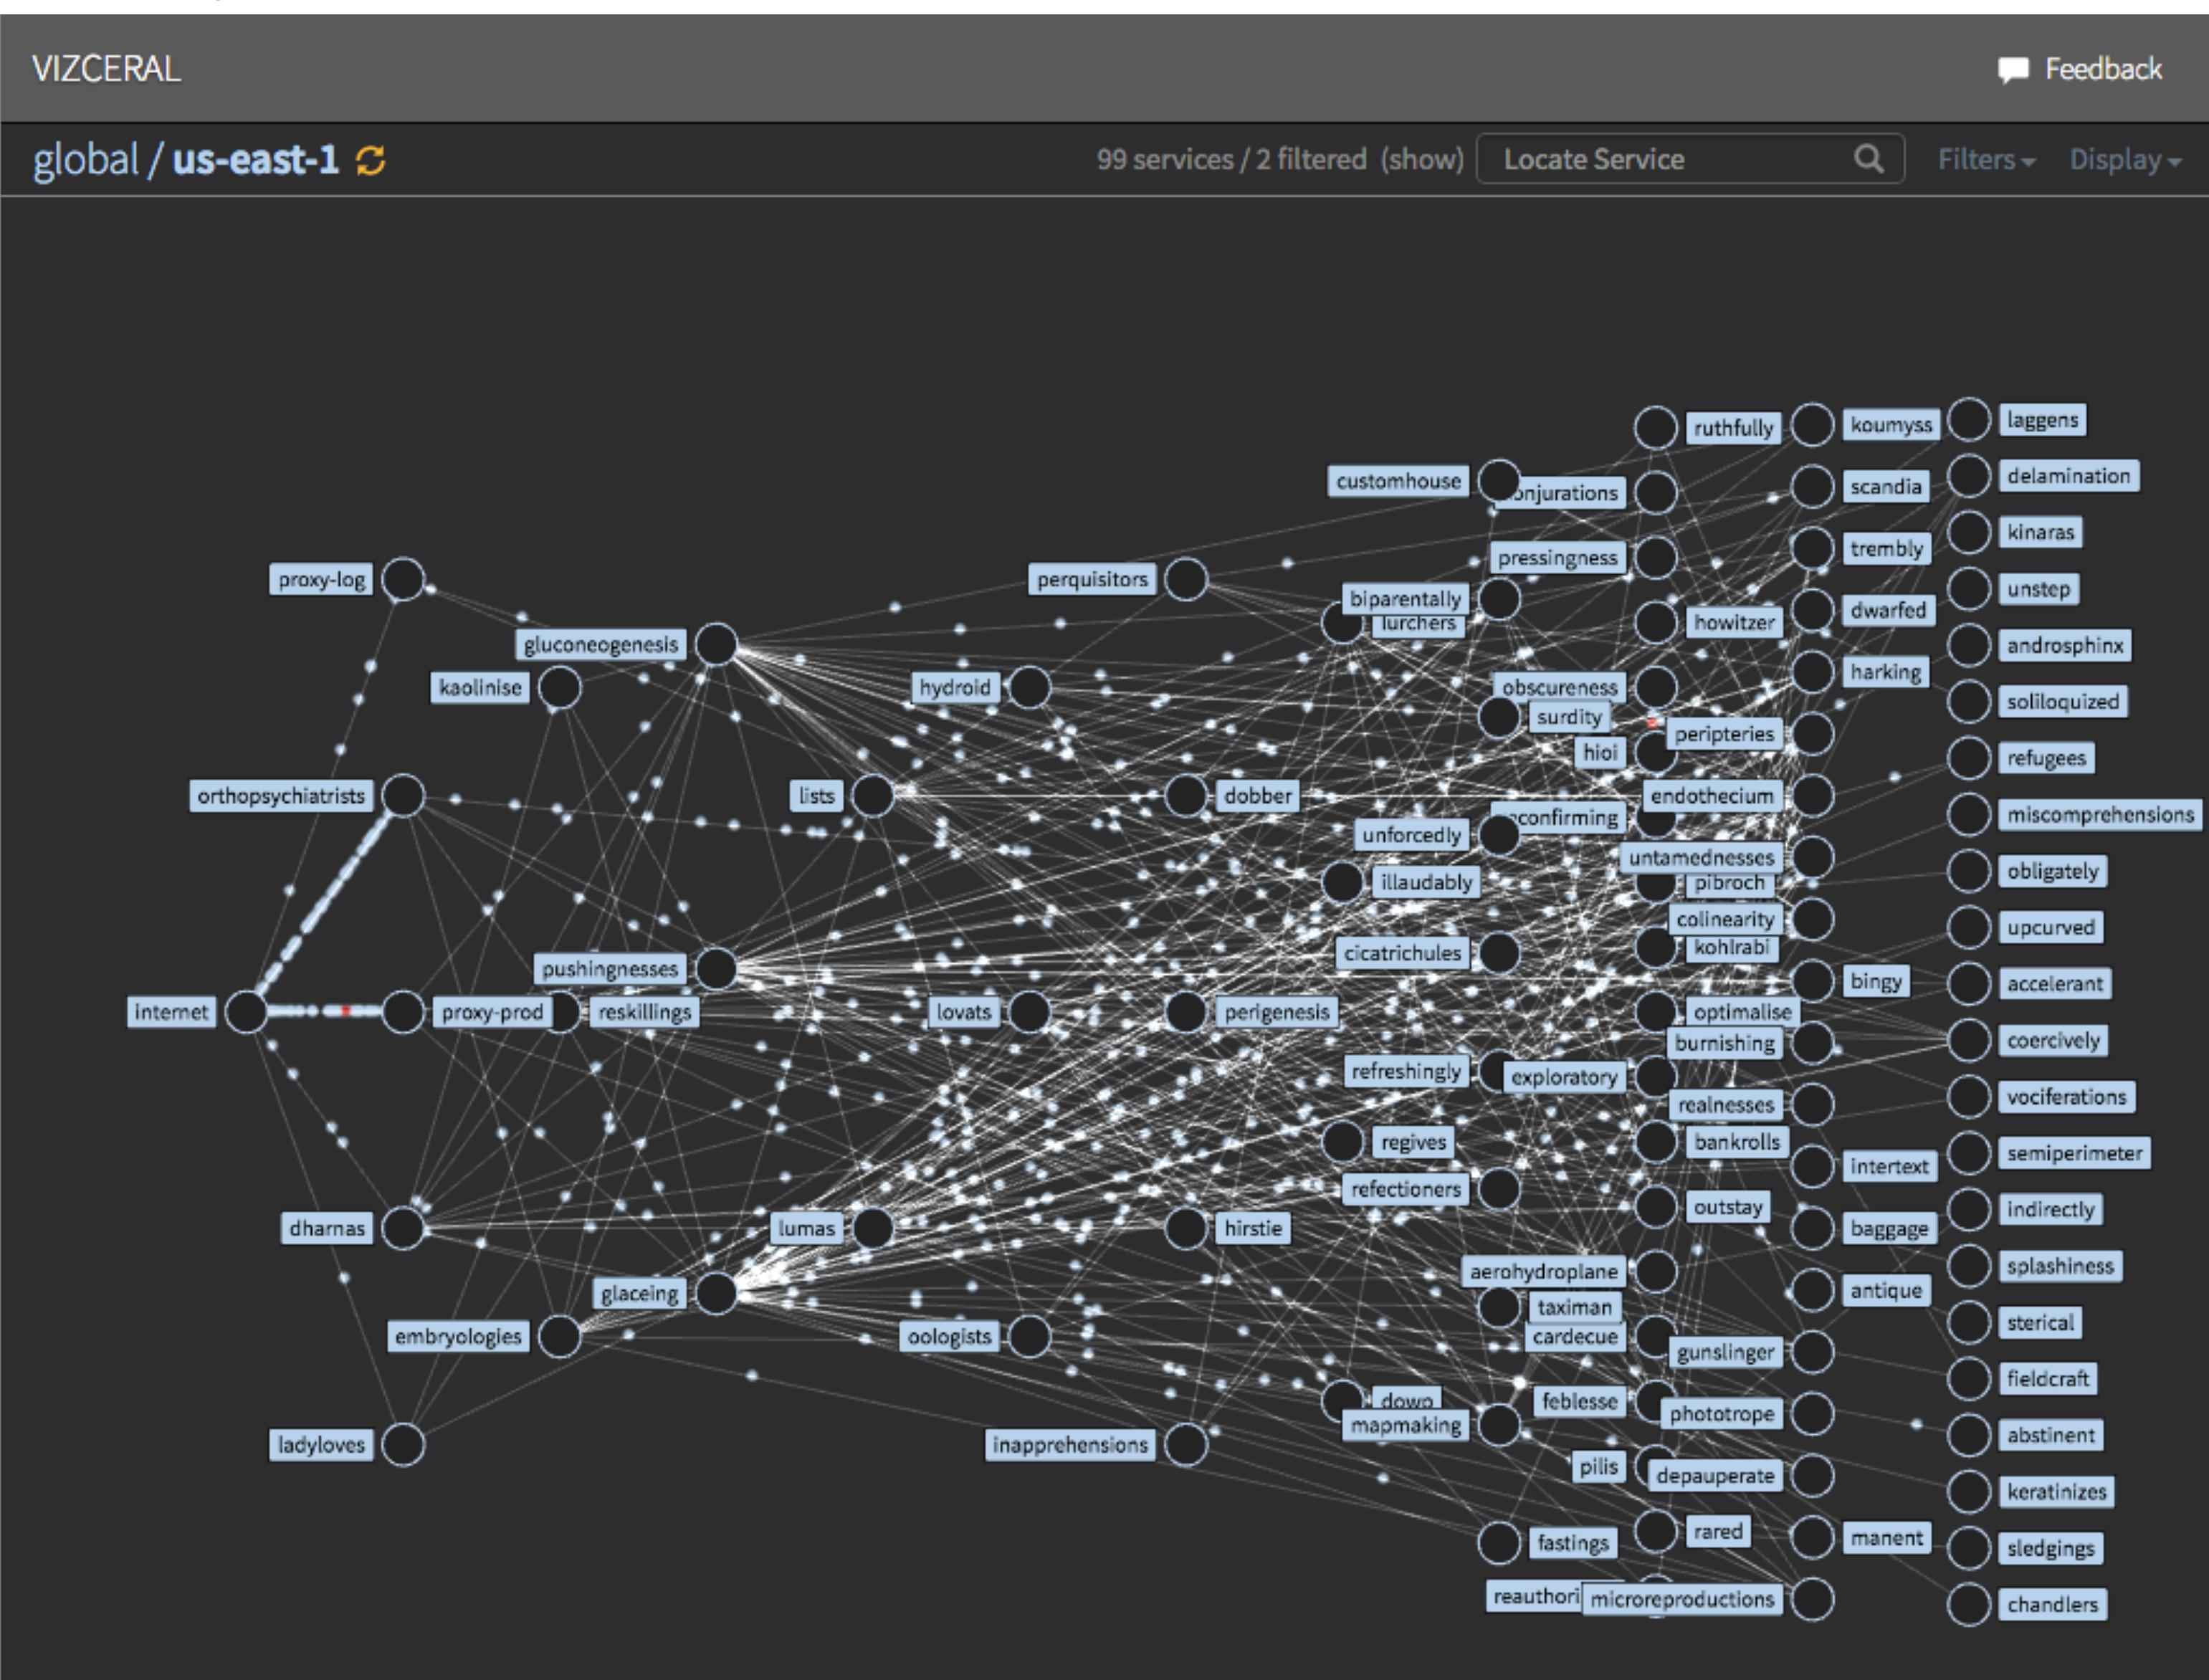
- | | |
|---|--|
|  Service graph representation |  Configuration validation |
|  Distributed tracing |  Health computation/display |
|  Metrics collection and graphs |  Service discovery |



Vizceral

- [Vizceral](#) is only a visualization font-end written in javascript, and it can support any type of datasource, as long as the data is formatted correctly.
- To have Prometheus working with Vizceral, there's a project called [promviz](#), which reads data from Prometheus, and convert the data in to Vizceral format to consume.
-  Vizceral is **not actively maintained** by Netflix anymore, there's a fork project called [promviz-front](#), and it can work with promviz perfectly, and also have several important new features, like replay, connection chart, etc.

Sample by Netflix



Resources

-  **NewRelic adapter repo:** <https://github.com/IBM/newrelic-istio-adapter>
- Istio official website: <https://istio.io>
- IBM Istio 101: <https://github.com/ibm/istio101>

Technical blogs we posted:

- Enable Istio for your service: <https://www.ibm.com/blogs/bluemix/2018/11/enable-istio-for-your-service/>
- How to deploy a gRPC mode Istio Mixer adapter into k8s: <https://www.ibm.com/blogs/bluemix/2018/09/how-to-deploy-a-grpc-mode-istio-mixer-adapter-into-kubernetes/>
- How to Enable Kiali for Istio on IBM Cloud Kubernetes Service: <https://www.ibm.com/blogs/bluemix/2018/08/how-to-enable-kiali-for-istio-on-ibm-cloud-kubernetes-service/>

Thanks!