

# Serverless IoT Applications

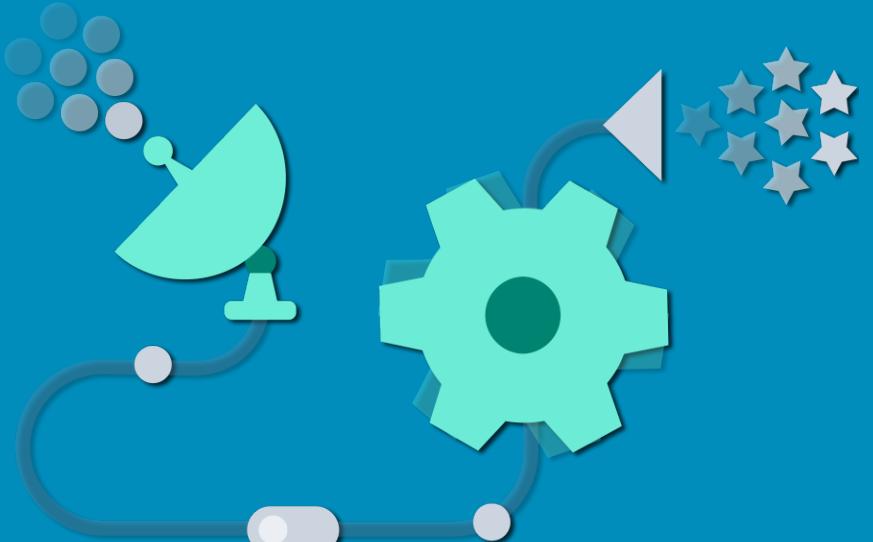
*with Apache OpenWhisk  
on IBM Bluemix*

<http://openwhisk.org/>

 #openwhisk

 <https://openwhisk-team.slack.com/>

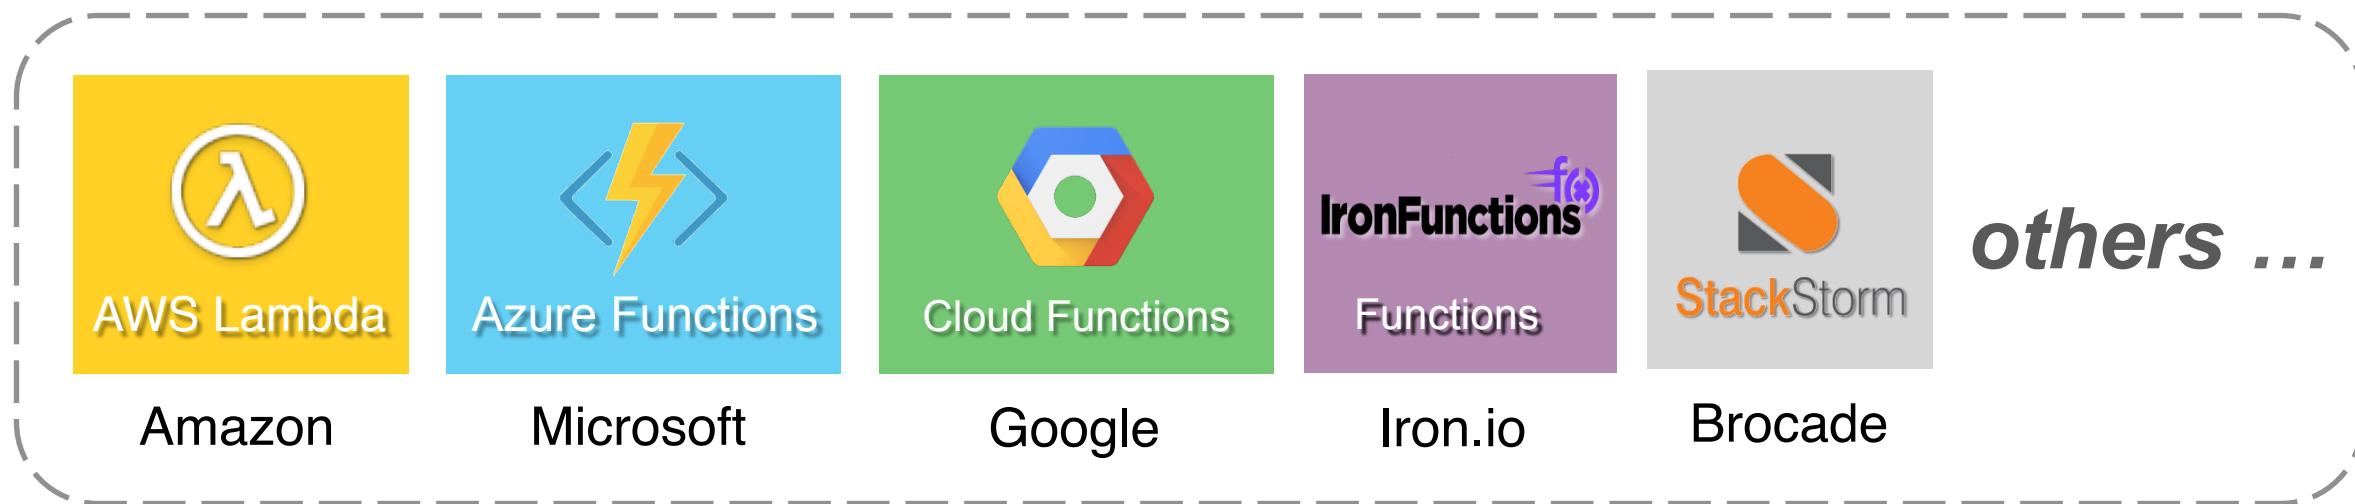
 <https://console.bluemix.net/openwhisk/>



- 1. Overview of the Serverless Paradigm and Apache OpenWhisk**
- 2. The Internet of Things (IoT) and IoT Applications**
- 3. Serverless Application Patterns for the IoT**

# 1. Overview of the Serverless Paradigm and Apache OpenWhisk

# Many Cloud providers offer some form of Serverless framework

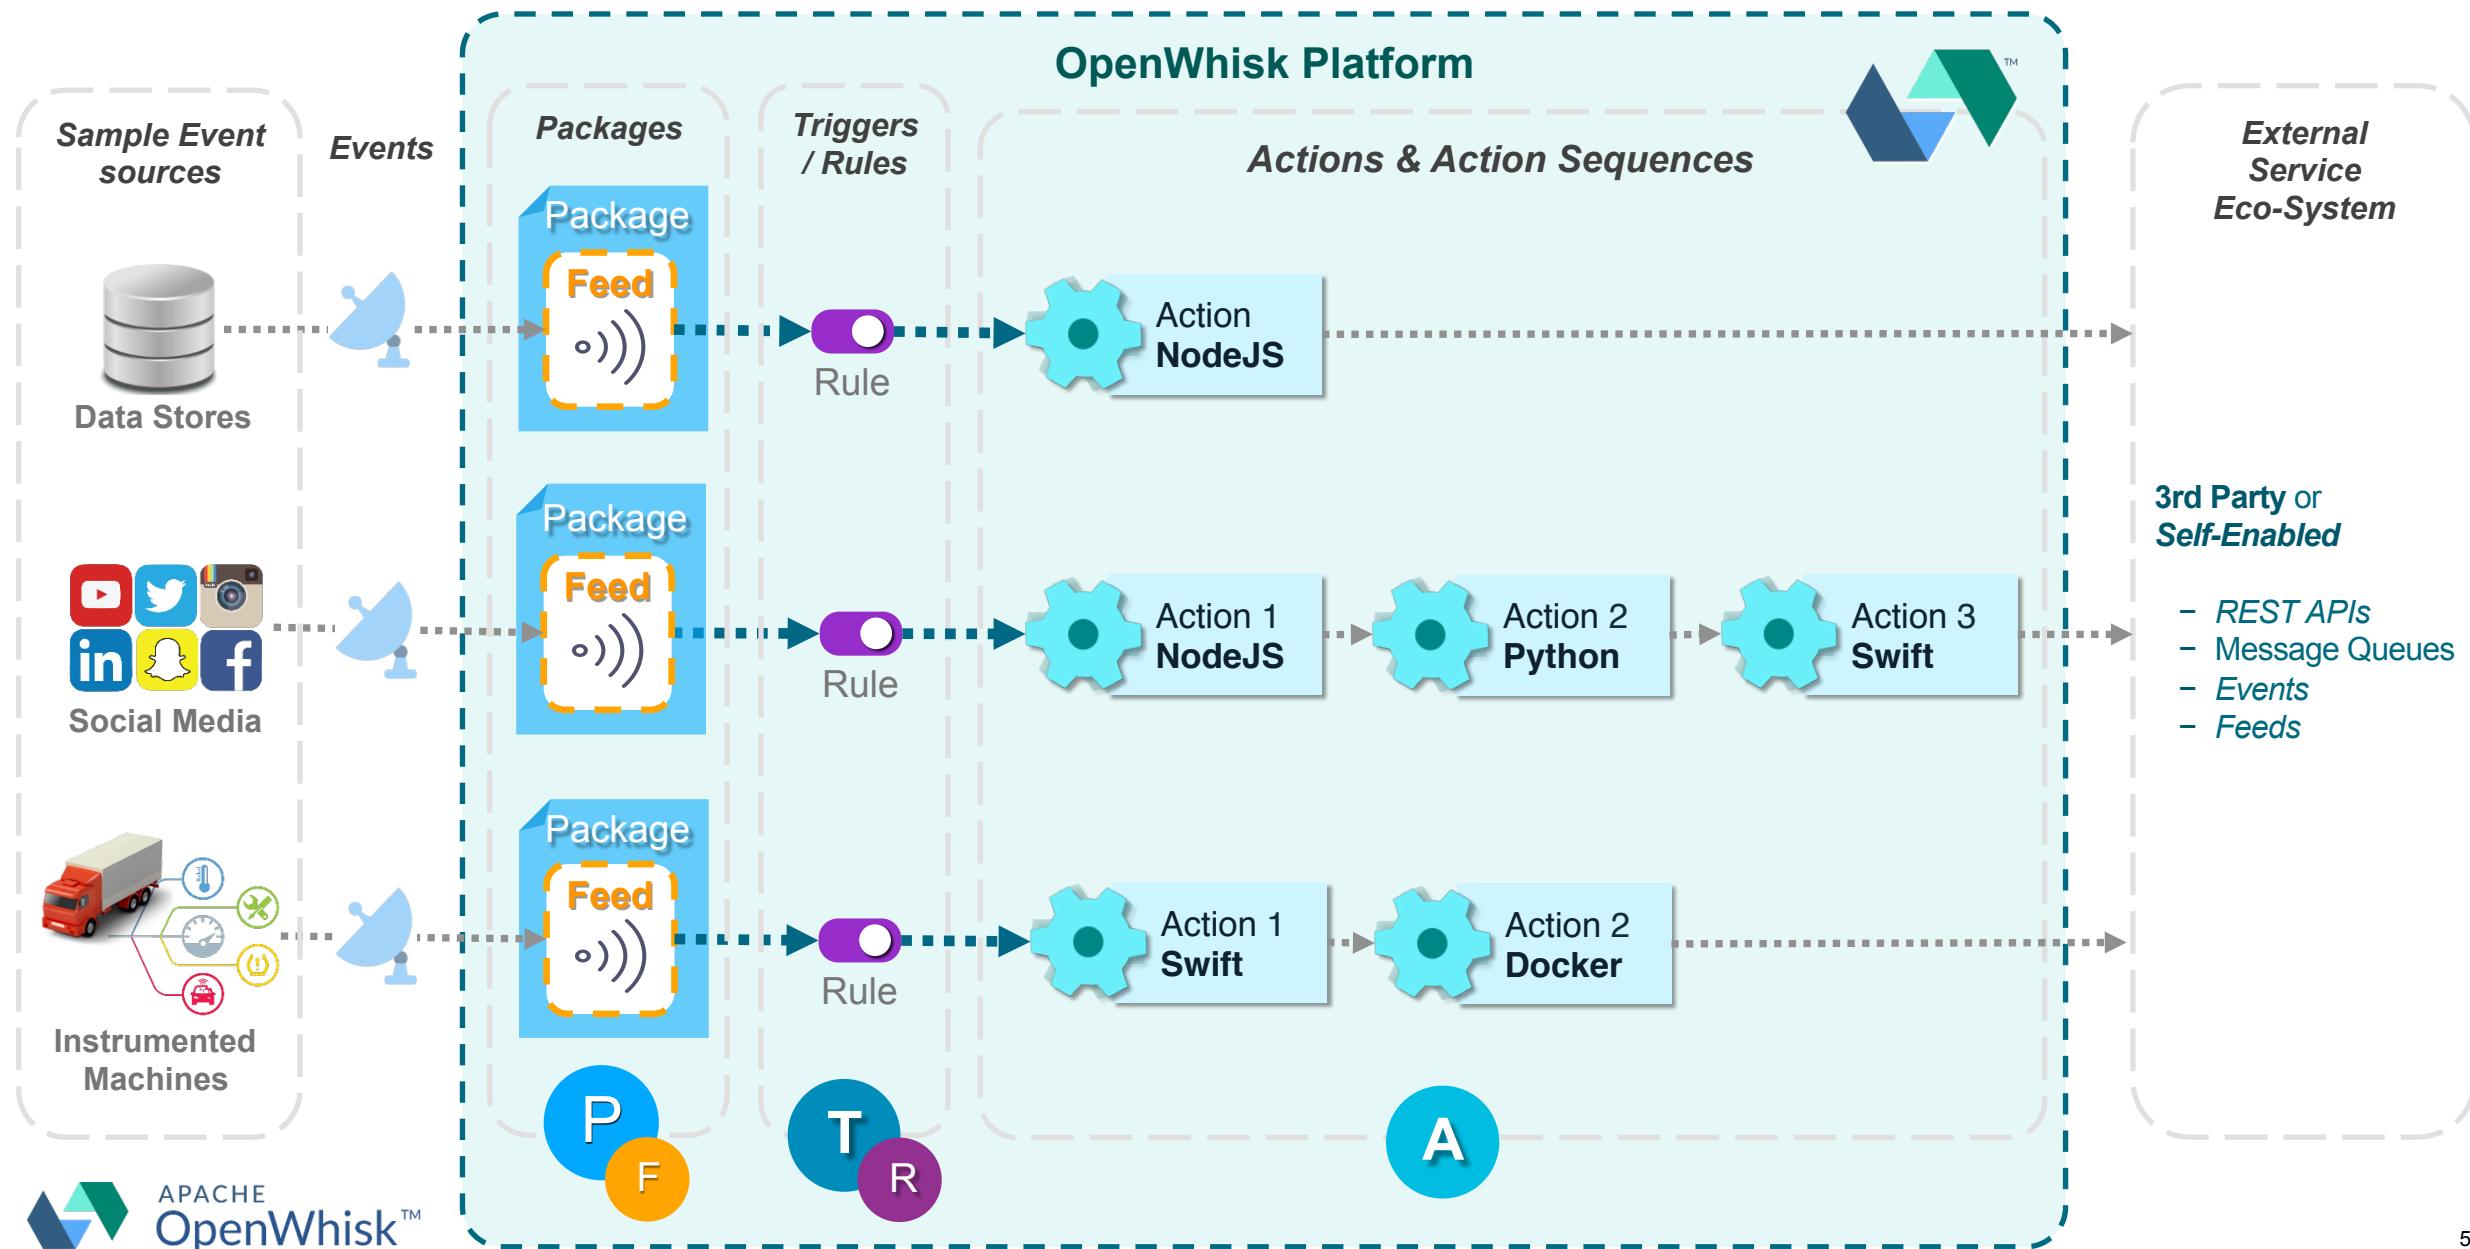
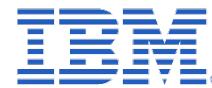


## Apache OpenWhisk offers:

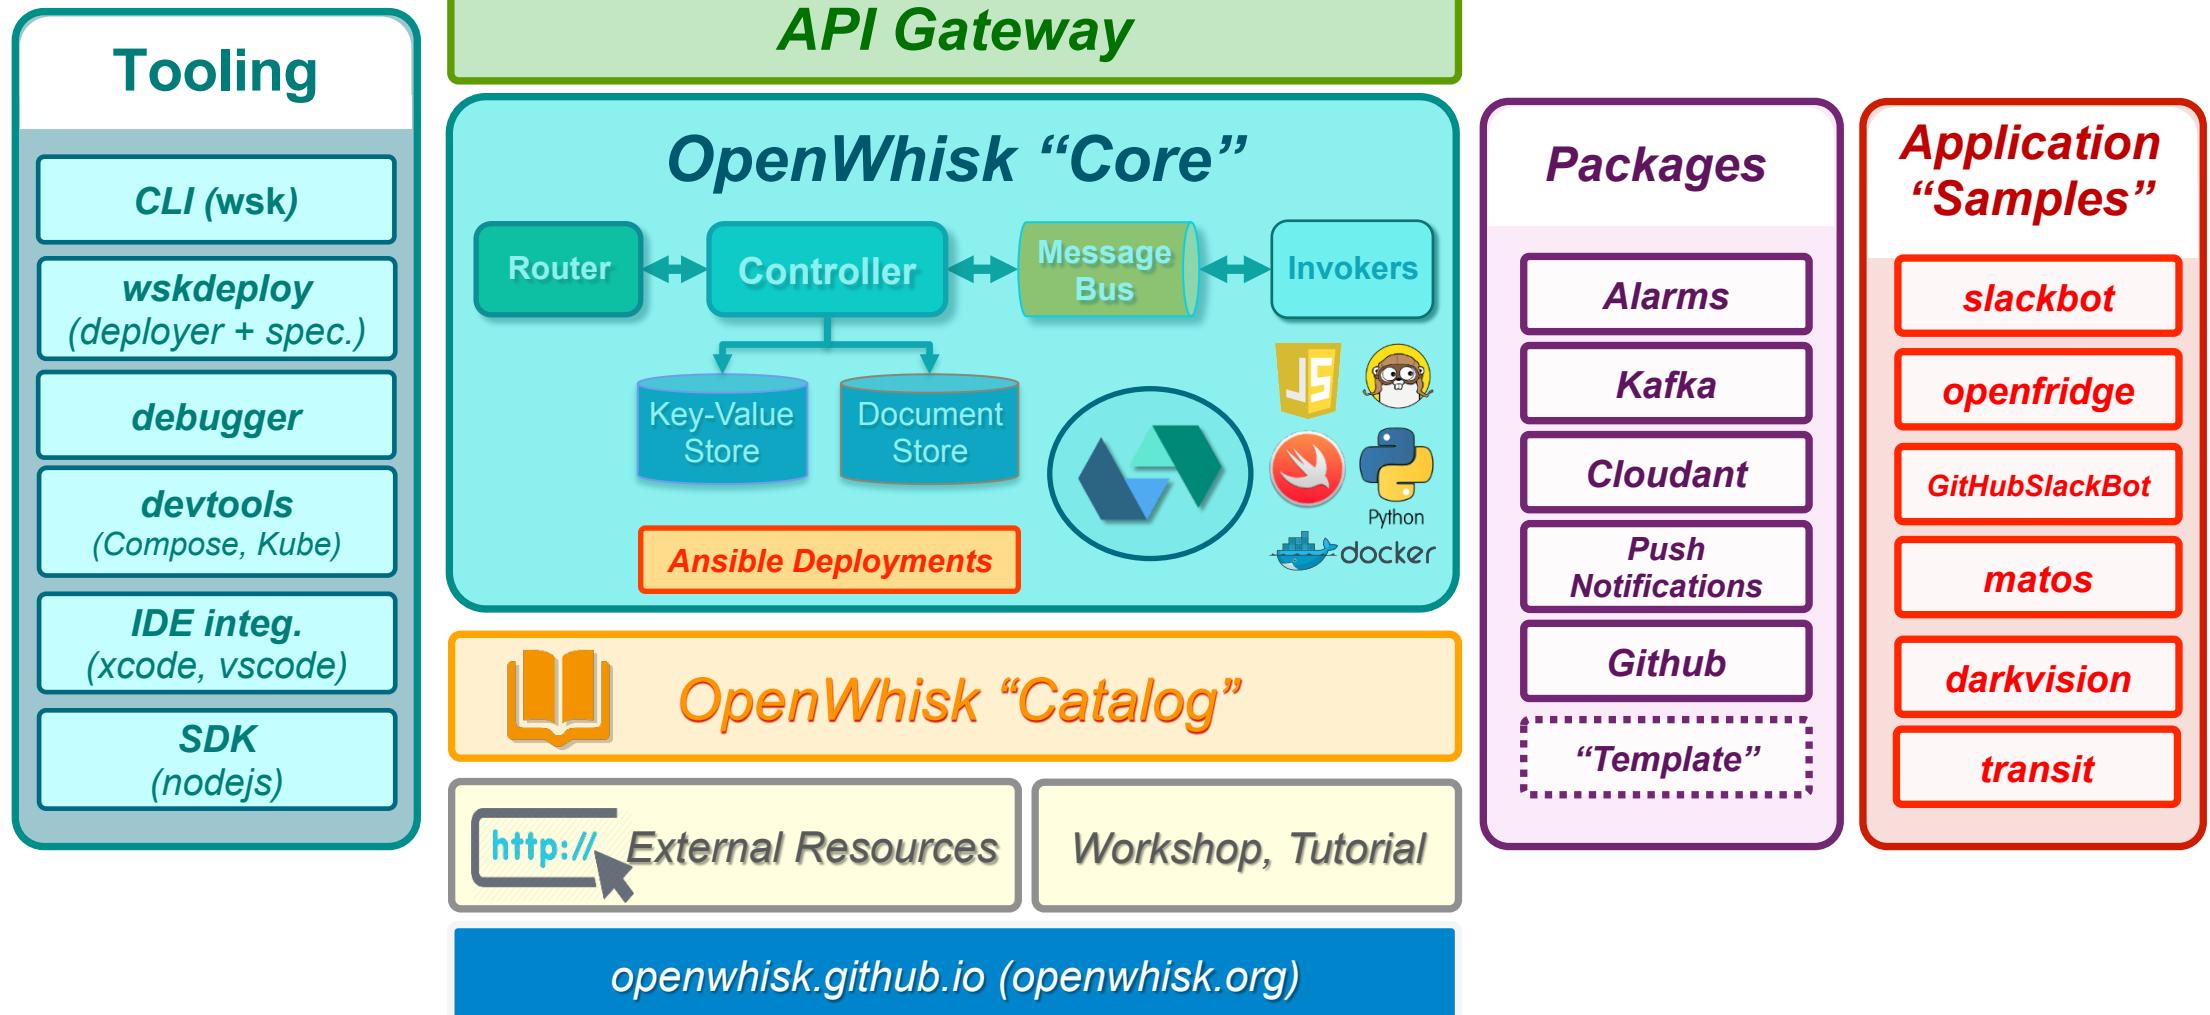
- **Apache Software Foundation (ASF)**
  - *True community-driven open source (Apache 2 License)*
- **Proven on IBM's Bluemix Cloud Platform**
  - *Exact same code in open source*



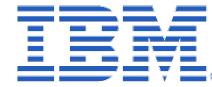
# Apache OpenWhisk – Event-Trigger-Rule-Action Processing



# Apache OpenWhisk “Eco-System” Overview

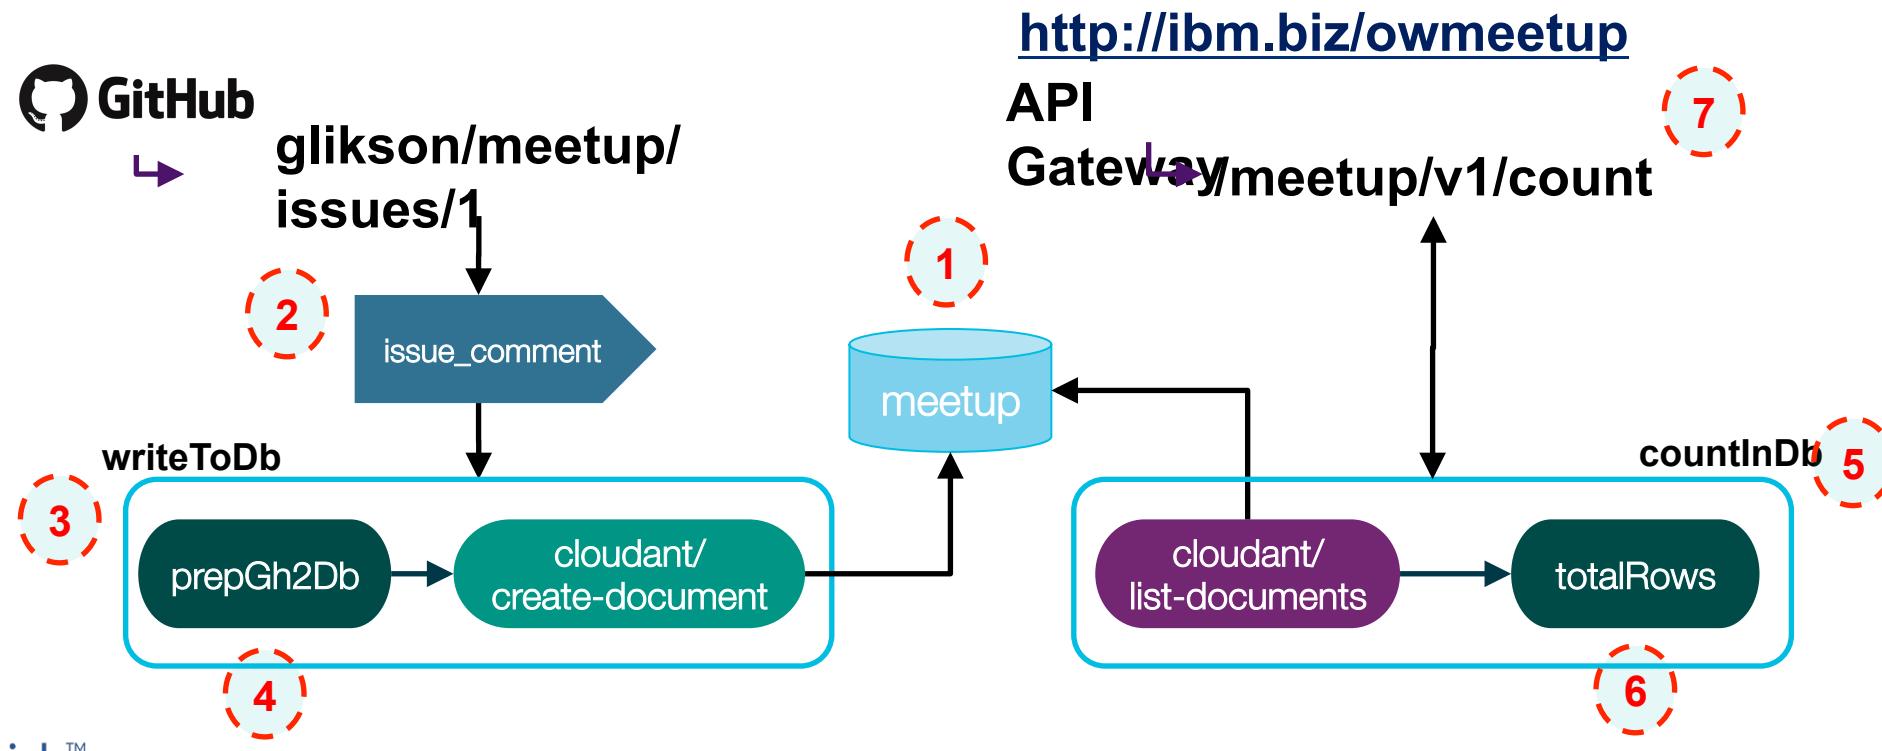


# Example: Monitoring Github comments



## ■ Goals:

1. Watch for new comments in a given github.com repository
2. Save github.com events in a DB (Cloudant)
3. Provide REST API to access data in the DB (e.g., aggregate stats)



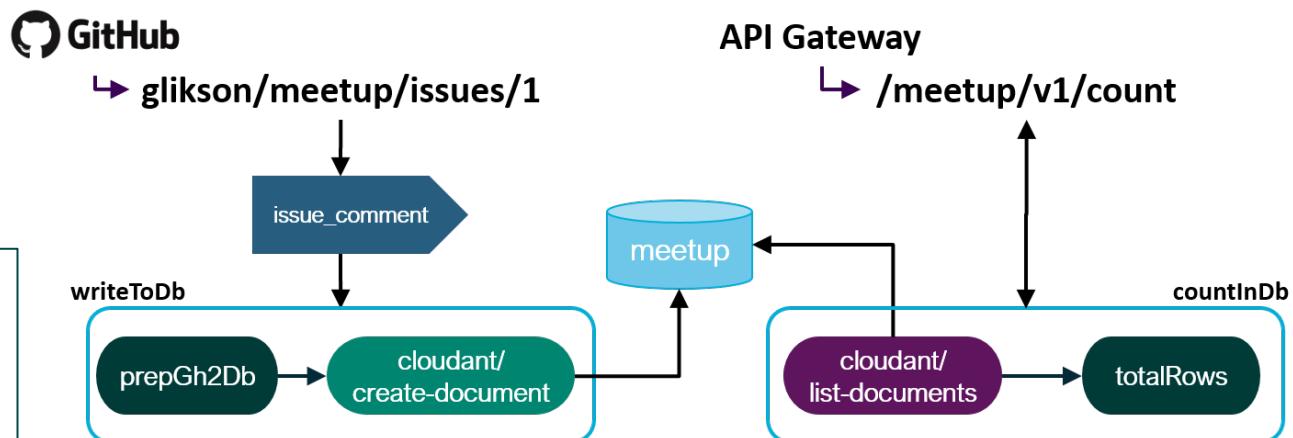
# Example: monitoring Github comments



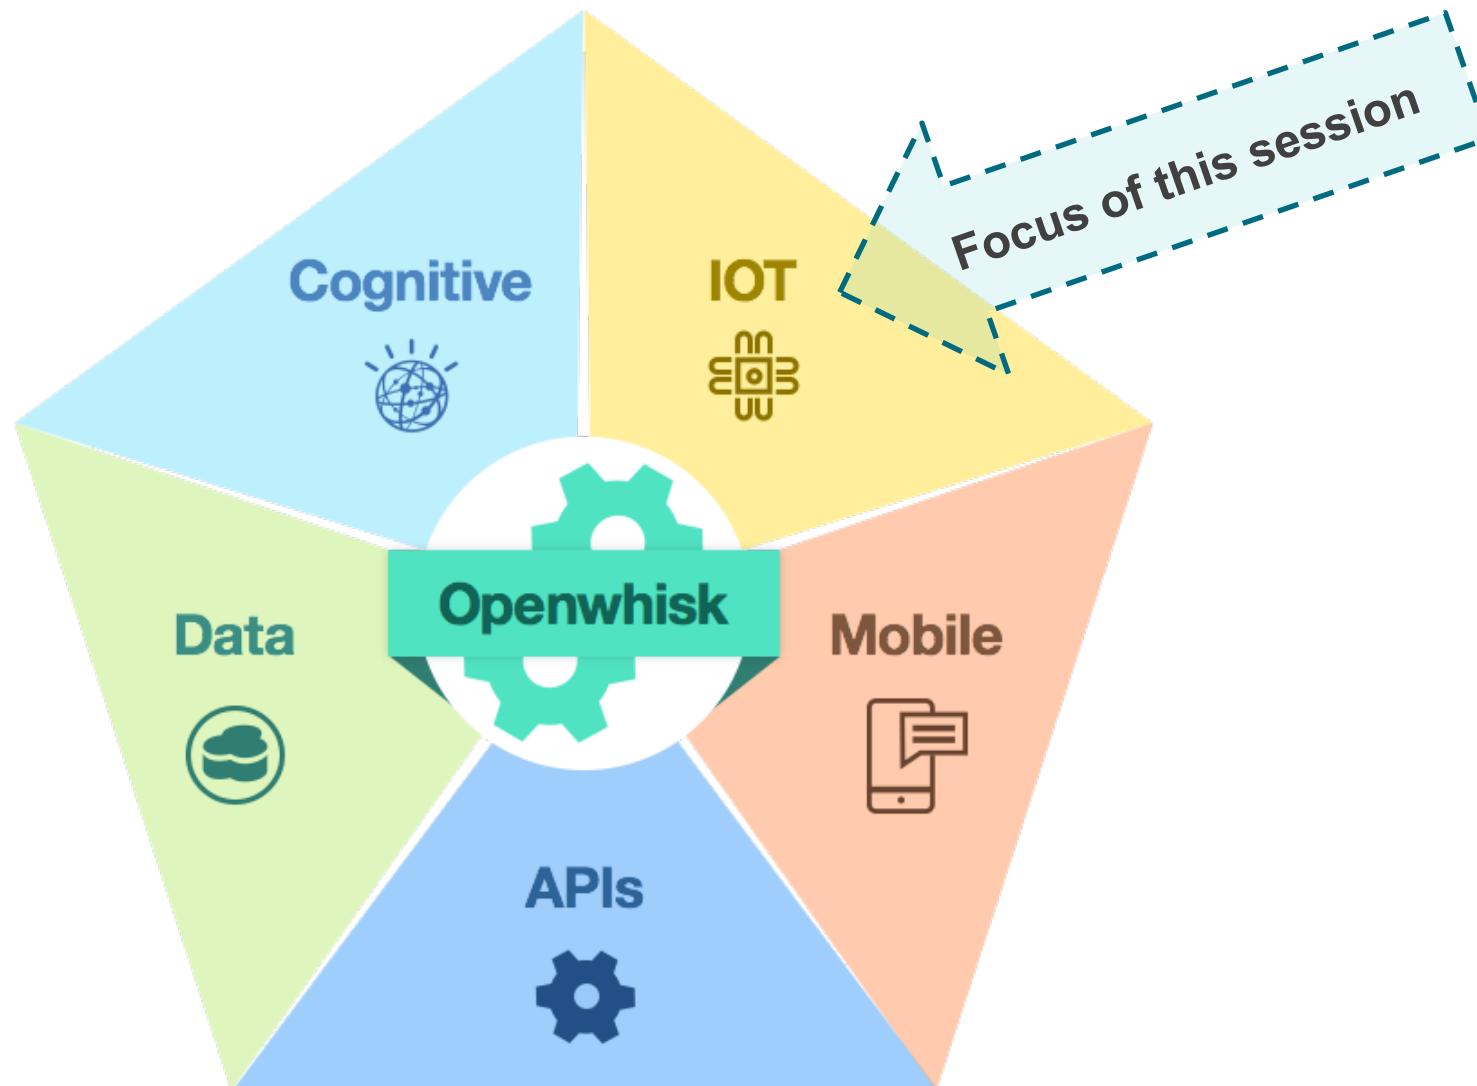
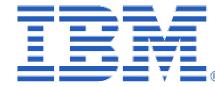
```
// prepGh2Db.js : converts Github event payload
// to be written to Cloudant DB
// using /whisk.system/cloudant/create-document
function main(hgPayload) {
    return {
        "dbname": hgPayload.repository.name,
        "doc": hgPayload
    };
}

// totalRows.js : returns total count
// of rows in a Cloudant DB.
// should be invoked following
// cloudant/list-documents
function main(docs) {
    return { "count": docs.total_rows};
}
```

```
[1] $ wsk action create --kind nodejs:6 prepGh2Db prepGh2Db.js
[2] $ wsk action create --kind nodejs:6 totalRows totalRows.js
[3] $ wsk package refresh
[4] $ wsk action create --sequence writeToDb prepGh2Db,Bluemix_db_meetup/create-document
[5] $ wsk action create --sequence countInDb Bluemix_db_meetup/list-documents,totalRows
[6] $ wsk trigger create github-meetup-issues --feed /whisk.system/github/webhook -p username \
glikson -p repository meetup -p accessToken xxx -p events issue_comment
[7] $ wsk rule create gh2db github-meetup-issues writeToDb
[8] $ wsk api-experimental create /meetup/v1 /count get countInDb
[9] $ curl https://ffac6dba-9f5c-4444-gws.api-gw.mybluemix.net/meetup/v1/count?dbname=meetup
{
    "count": 9
}
```



# Serverless (and OpenWhisk) is applicable in many domains

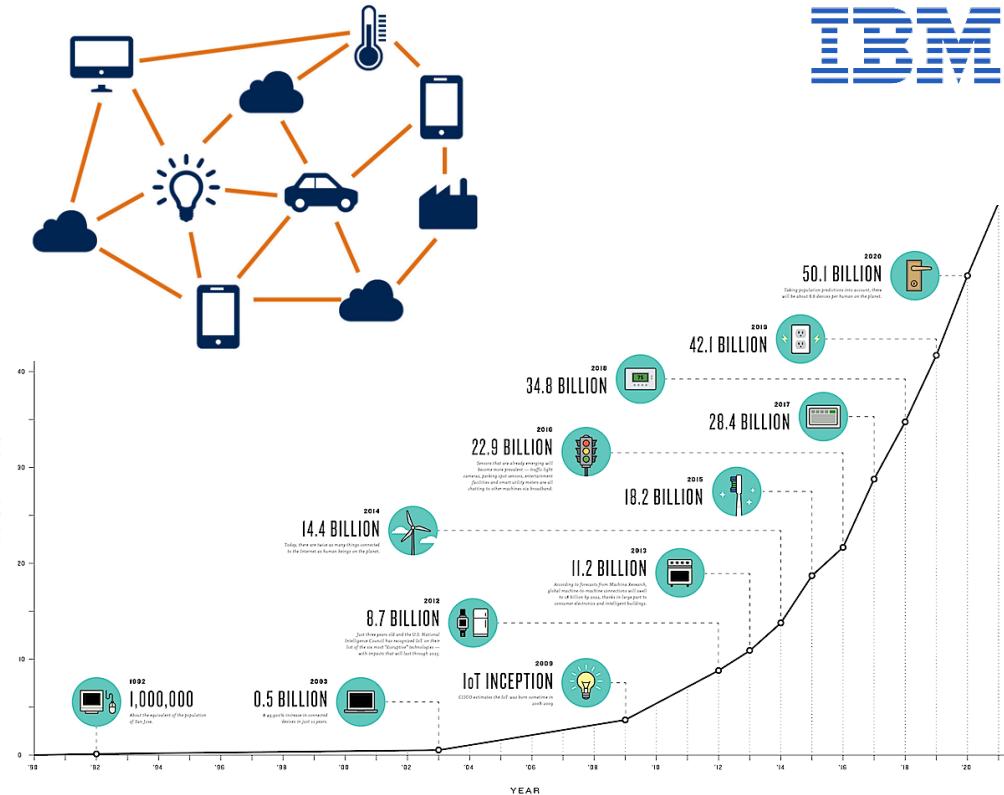


## 2. The Internet of Things (IoT) and IoT Applications

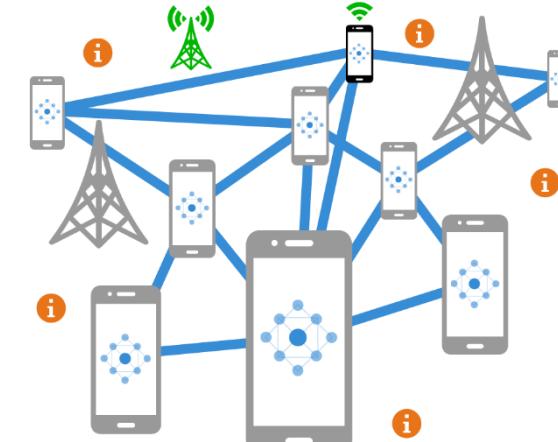
# What is Internet of Things (IoT)?

## ▪ Internet of Things (IoT) applications:

- Rapidly emerging application development paradigm, taking advantage of the ubiquitous **connectivity** and **programmability** of physical **devices**:
  1. Environment sensors (e.g., cities, private weather stations)
  2. Consumer electronics (e.g., wearables, home appliances)
  3. Industry (e.g., factories, agriculture, buildings)
  4. Mobility (e.g., personal navigation, connected cars)



- Synergetic with the phenomenal adoption of **smartphones** and the rapid growth of the enabling communication infrastructure, providing **commodity** (low-cost, ubiquitous) connectivity
  - 3G, 4G, 5G ...



# Characteristics of a Typical IoT Application



## 1. Big Data: large volume (e.g., # devices) and variety (e.g., message formats)

- Sensors (e.g., location, temperature, audio, motion)
- Actuators (e.g., valve, alarm, motor)

## 2. Devices are typically connected via Gateways

- Point of aggregation and control for devices in physical proximity
- Device-to-gateway communicating is often not TCP/IP-based (e.g., Bluetooth Low Energy)

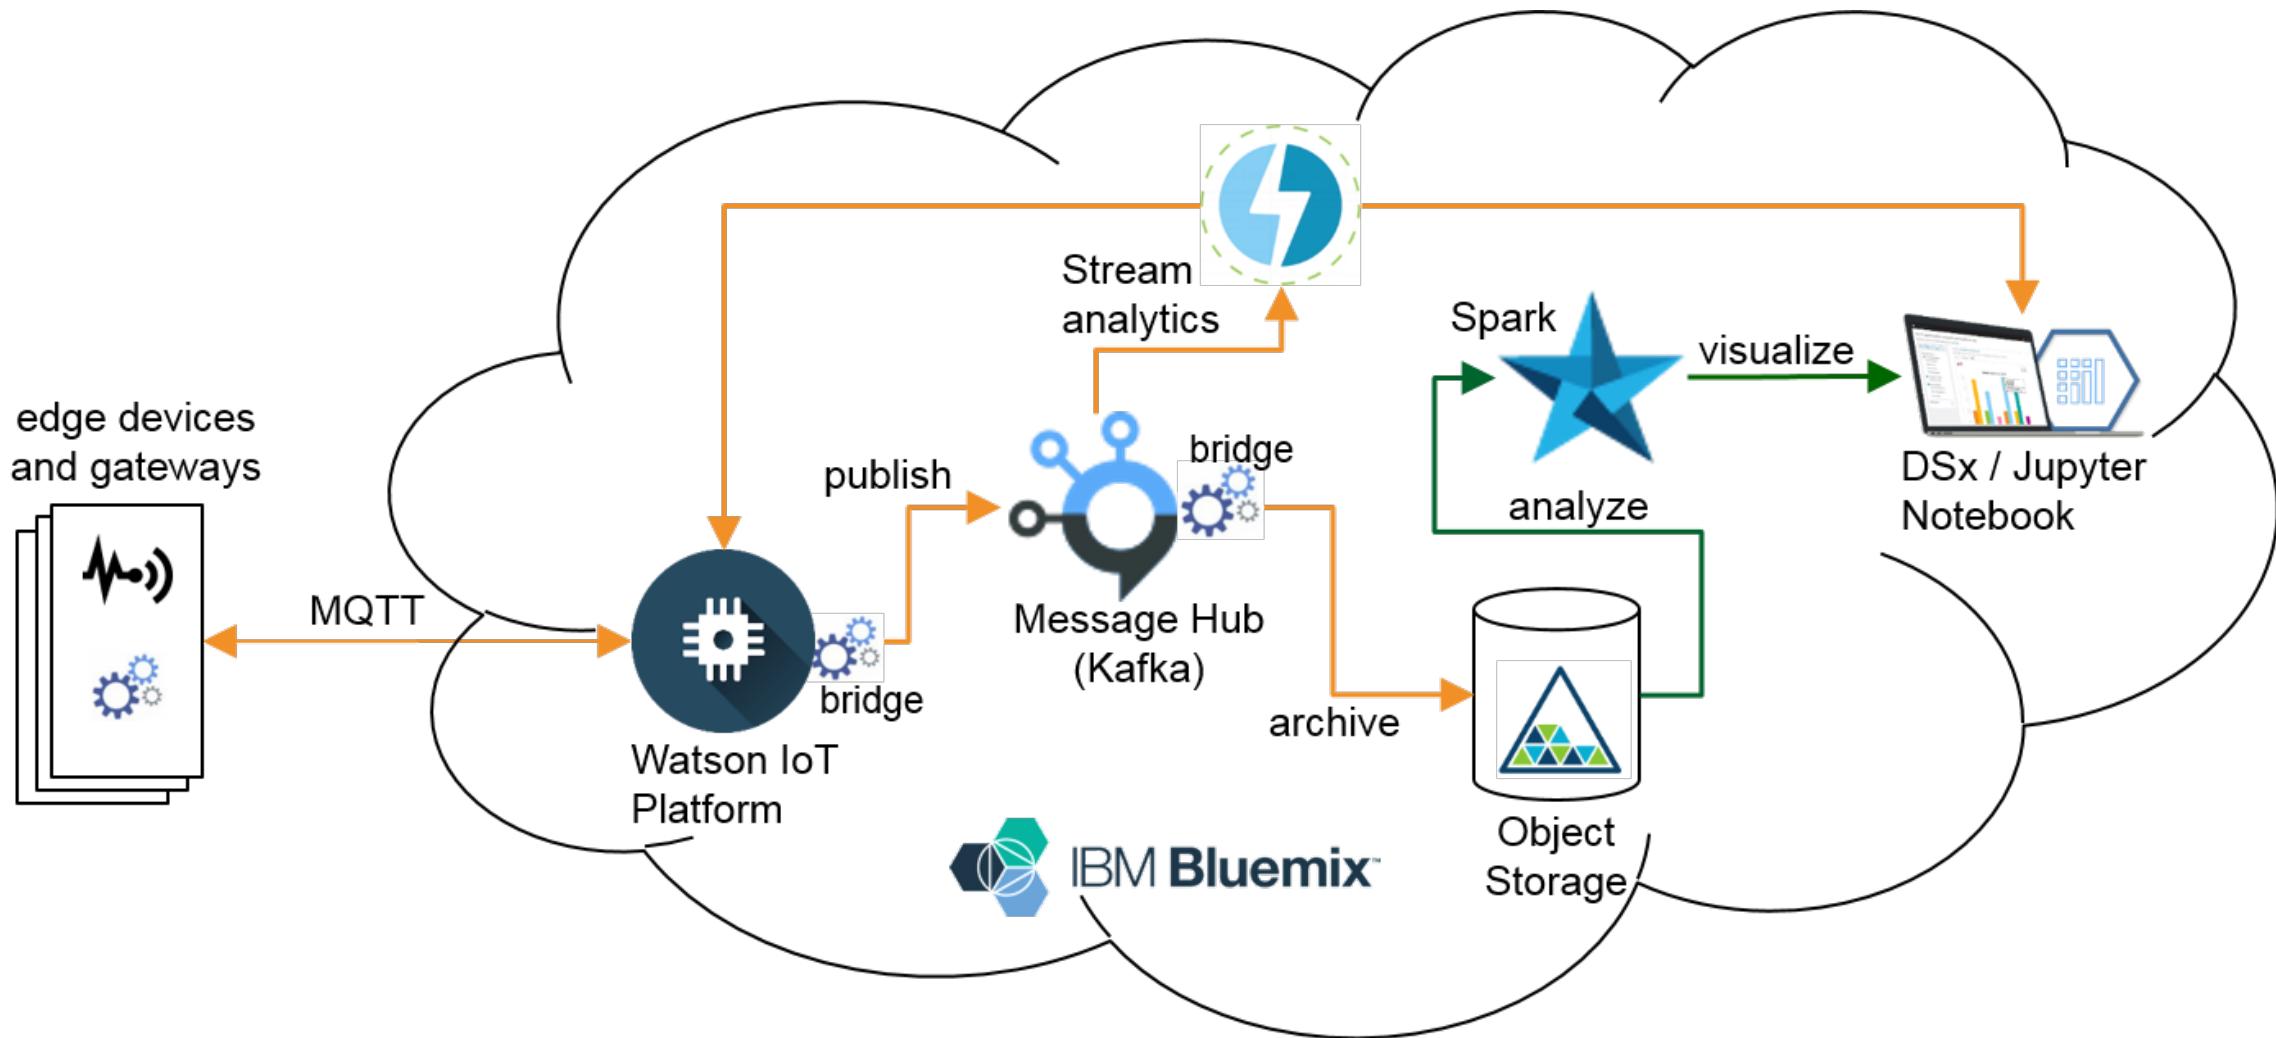
## 3. Data often undergoes initial processing at the Gateway

- Motivated by latency, bandwidth or privacy constraints
- Often comprises simple real-time processing (e.g., property-based filtering, time-based aggregation)
  - May trigger local actuation (e.g., based on pre-defined rules)

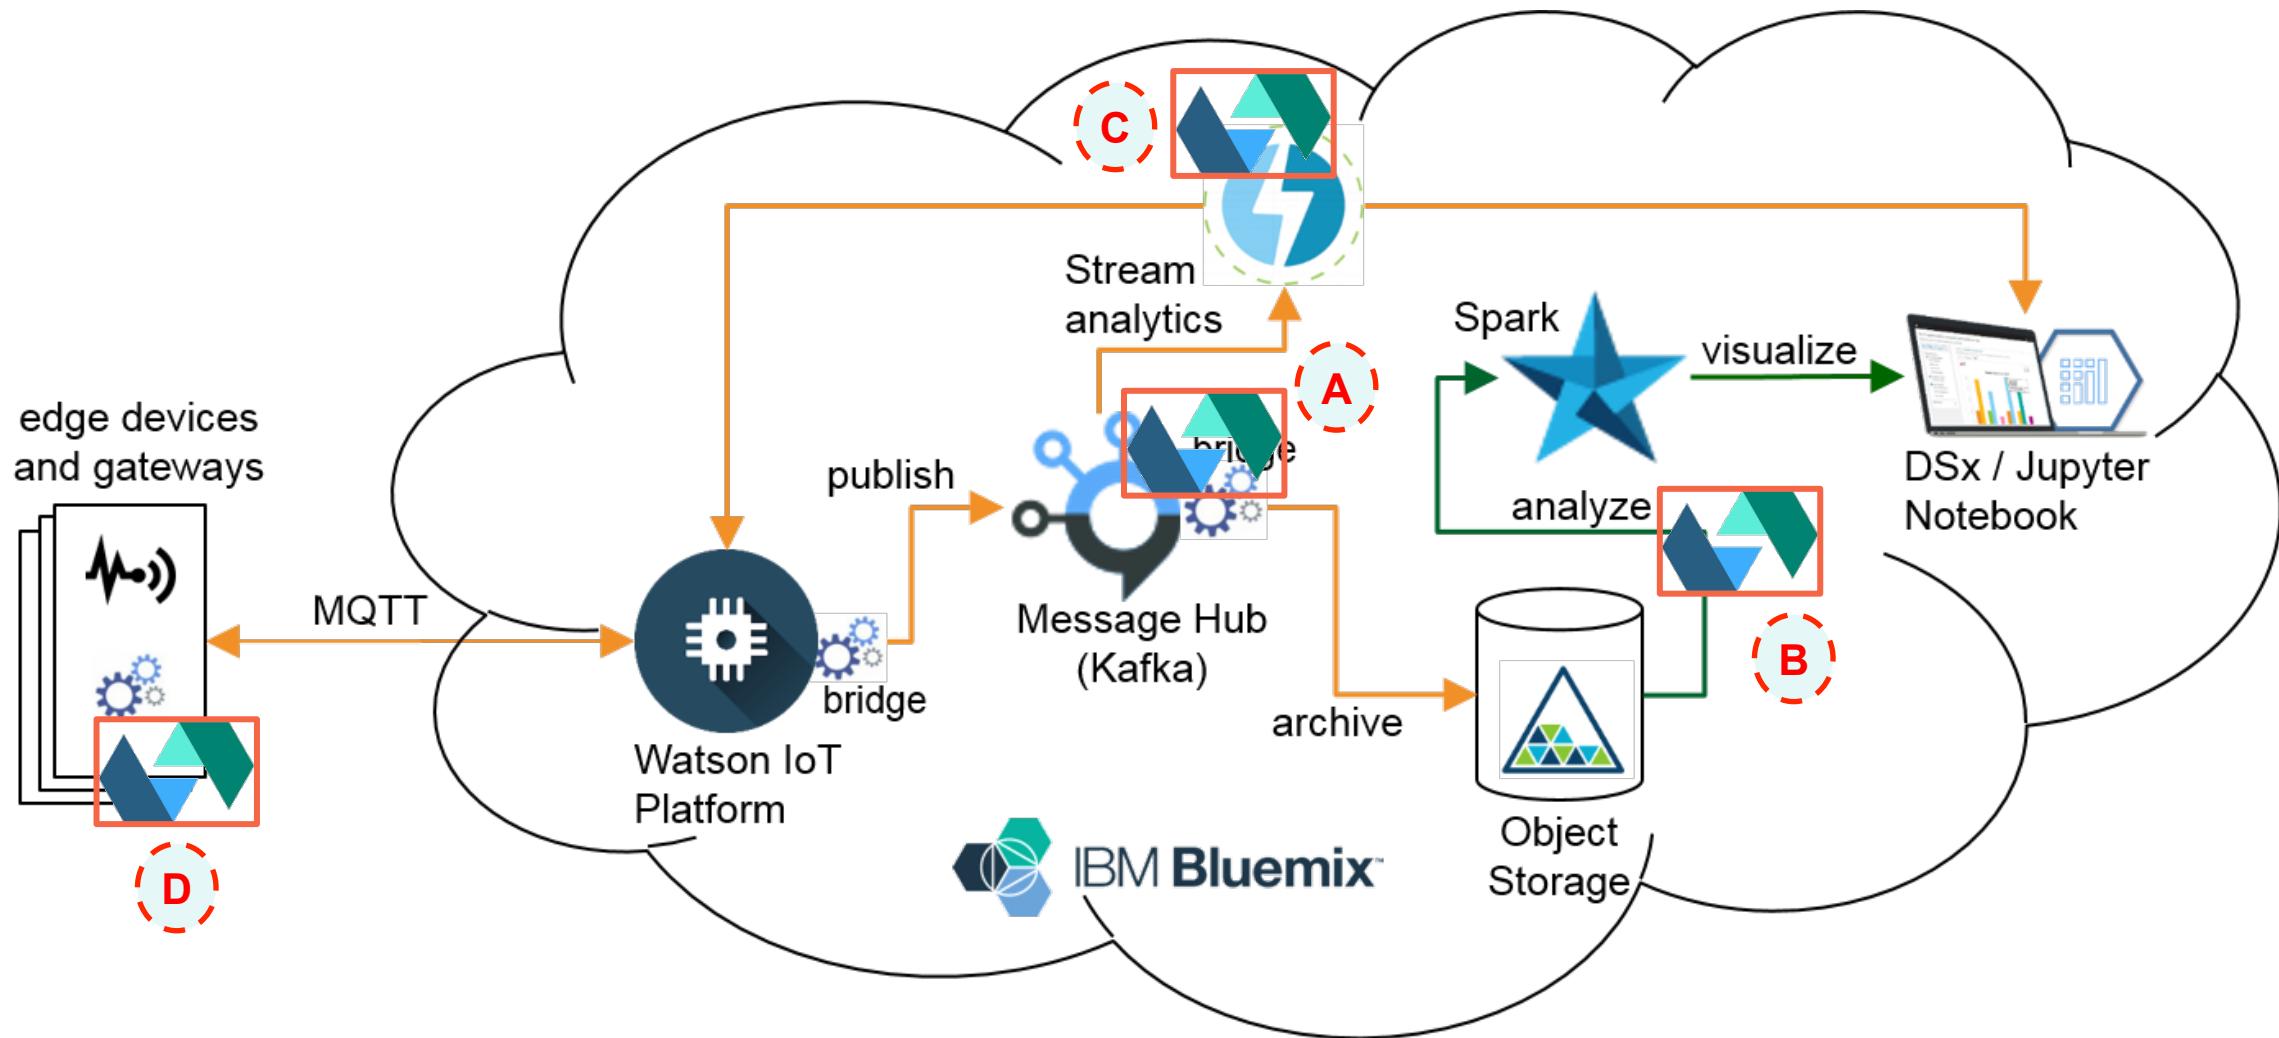
## 4. Data is shipped to a centralized cloud for storage and analytics

- Both (near) real-time and batch
- May trigger a feedback loop to the gateways
  - E.g., actions triggered by a detected anomaly, updates of filtering thresholds

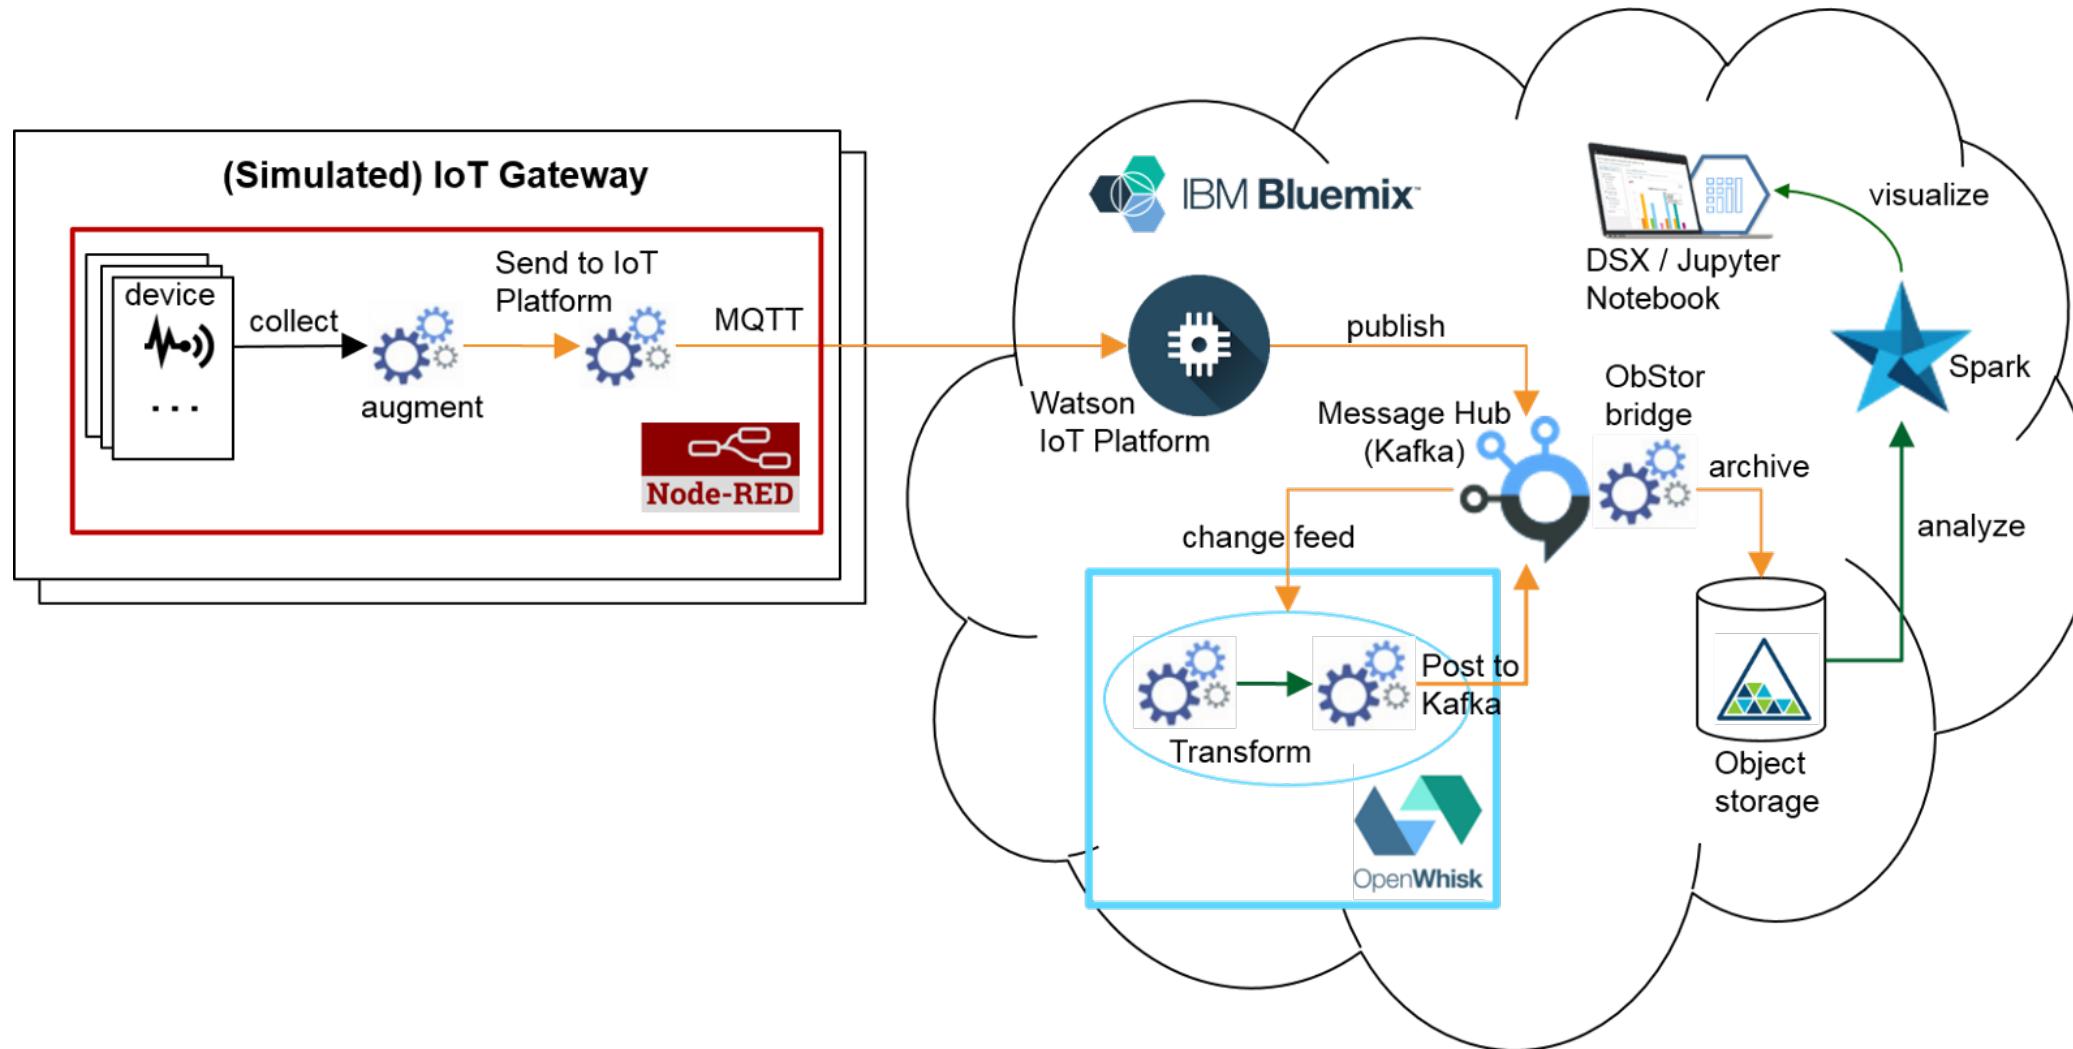
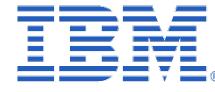
# Example of IoT Application in IBM Bluemix



# Opportunities for Serverless/FaaS in IoT

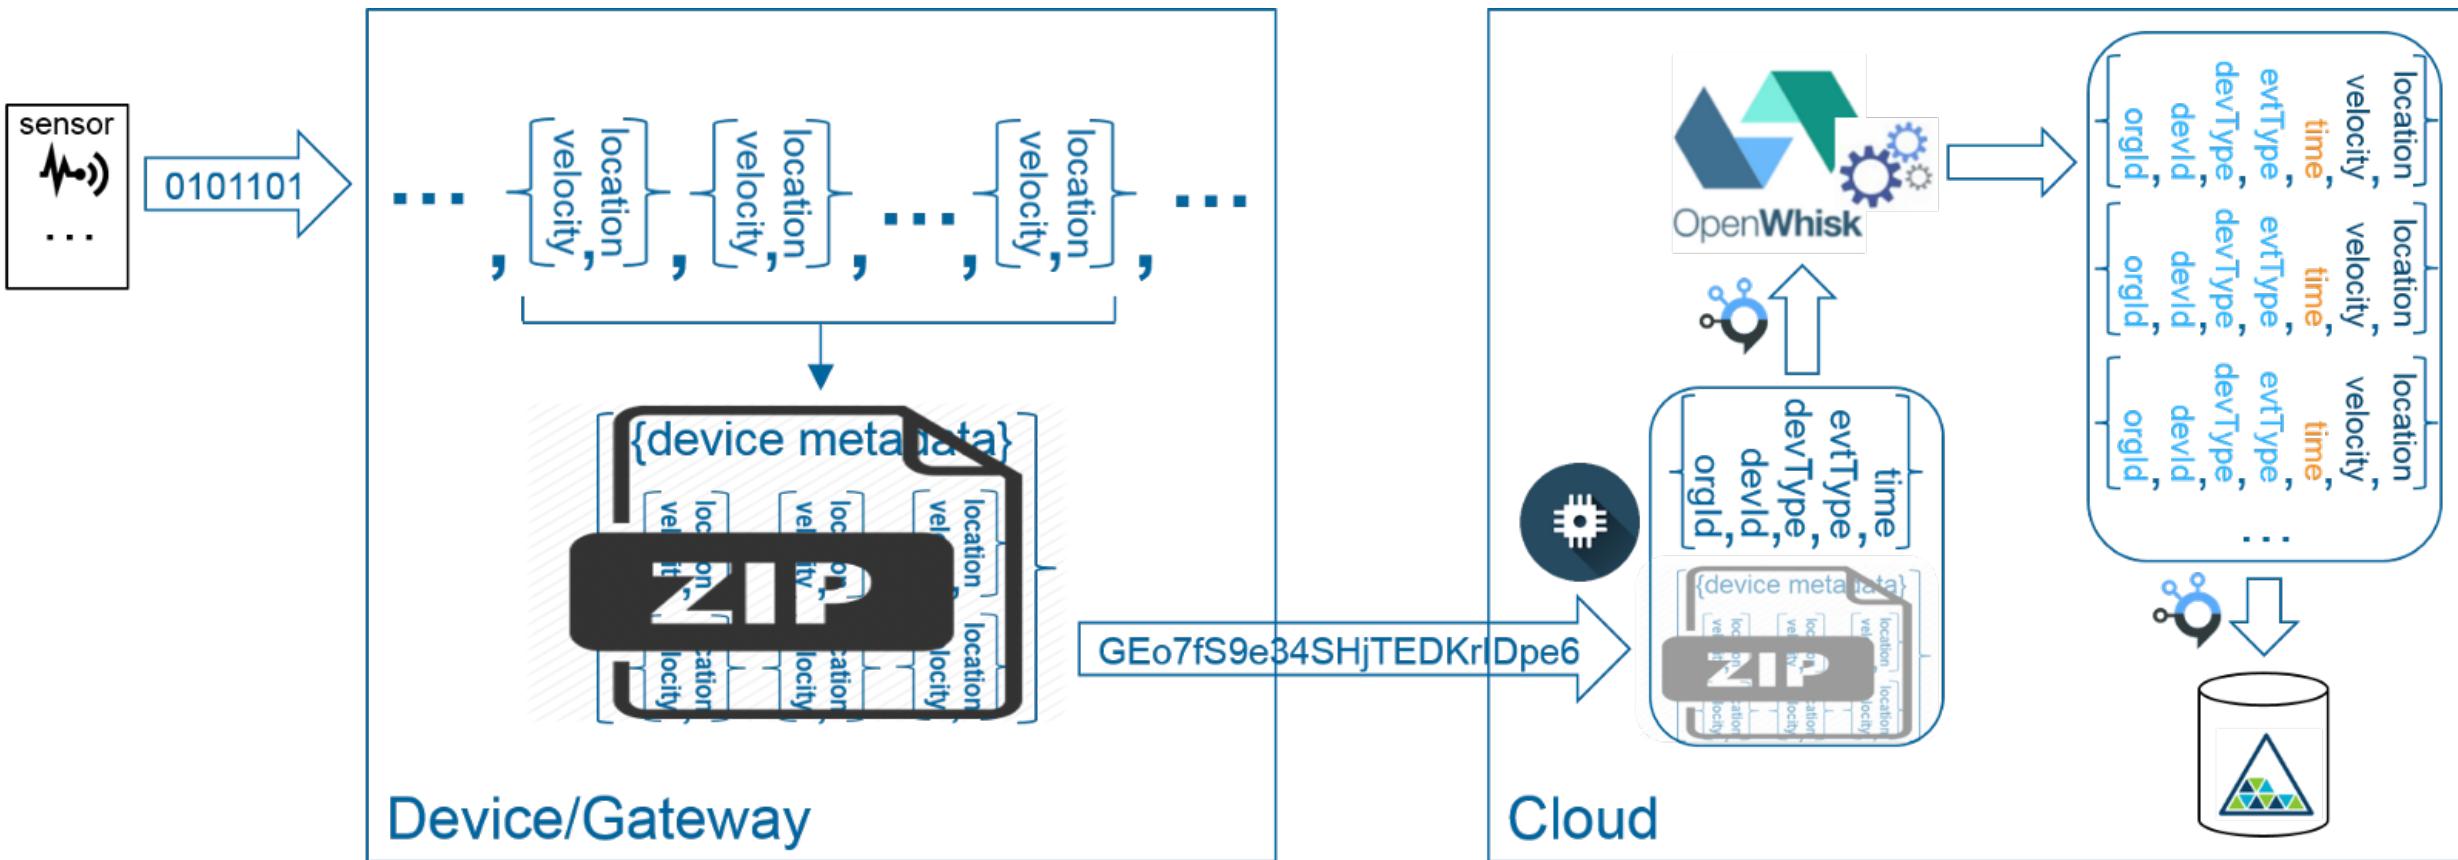


# TRANSIT: Serverless Transformation of IoT Data-in-Motion

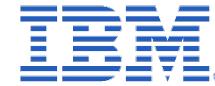


<https://medium.com/openwhisk/serverless-transformation-of-iot-data-in-motion-with-openwhisk-272e36117d6c>

# Schematic end-to-end data flow

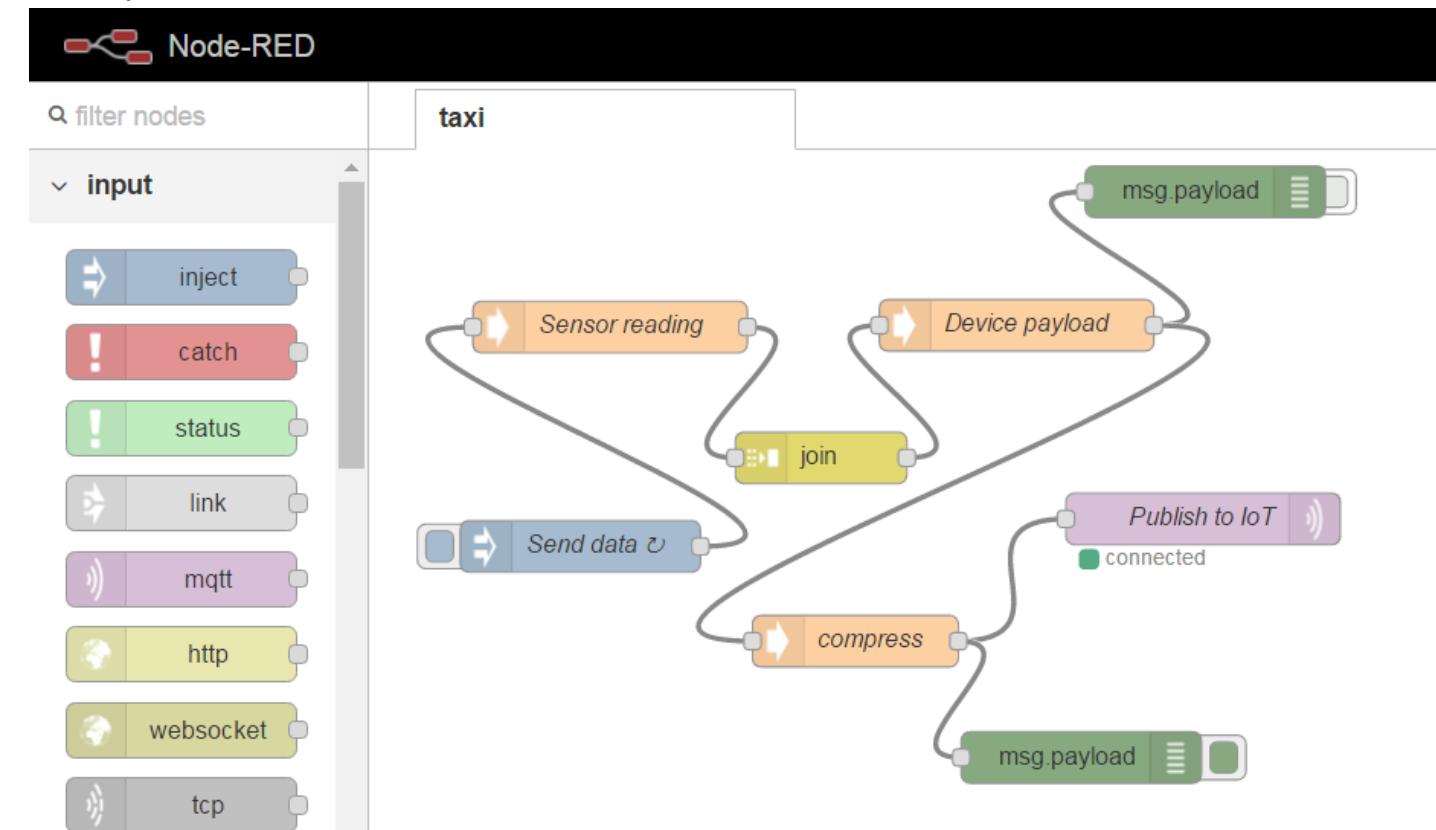


# End-to-end data flow – 1 (Node-RED)



## ■ Node-RED (simulated device):

- The ‘inject node’ triggers the flow at given intervals (e.g., every second).
- A sample event JSON is generated (by a custom ‘function node’).
- It is then aggregated using the ‘join’ node,
- compressed (using a custom function node), and
- sent to the Watson IoT Platform  
(using MQTT with properly configured orgId, device type, event type, device id and security token).



## End-to-end data flow – 2 (Watson IoT Platform)



- **Watson IoT Platform (with historical data storage extension to Message Hub):**

- Receives MQTT events from devices (after proper auth validation),
- Augments them (e.g., adding timestamp and context metadata) and
- Publishes in small batches to a Kafka topic (Message Hub).

# End-to-end data flow – 3 (OpenWhisk)



## ▪ OpenWhisk (with custom transformation action):

- The Message Hub feed provider in OpenWhisk subscribes to messages in the *iotp* Kafka topic,
- Aggregates small batches of messages and
- Sends them as a payload to the OpenWhisk trigger, which in turn
- Triggers a sequence of actions
- Performing the message transformation and
- Publishing the results to the *transformed* Kafka topic

```
function main(params) {
  zlib = require('zlib');
  return new Promise(function(resolve, reject) {
    if (!params.messages || !params.messages[0] || !params.messages[0].key || !params.messages[0].value) {
      reject("Invalid arguments. Must include 'messages' JSON array with 'key' and 'value' fields");
    }
    var msgs = params.messages;
    var out = [];
    for (var i=0; i<msgs.length; i++){
      var msg = msgs[i];
      console.log ("Processing MSG=" + JSON.stringify(msg));
      var key = JSON.parse(msg.key);
      var sensorReadings = JSON.parse(zlib.gunzipSync(msg.value))["d"].sensor_data;
      console.log ("sensor data: " + sensorReadings);
      for (var j=0; j<sensorReadings.length; j++) {
        newmsg = {
          orgId: key.orgId,
          deviceType: key.deviceType,
          deviceId: key.deviceId,
          eventType: key.eventType,
          timestamp: key.timestamp.substr(0,19),
          sensorData: JSON.parse(sensorReadings[j])
        };
        out.push(newmsg);
        console.log("result: " + JSON.stringify(newmsg));
      }
    }
    resolve({ "messages": out });
  });
}
```

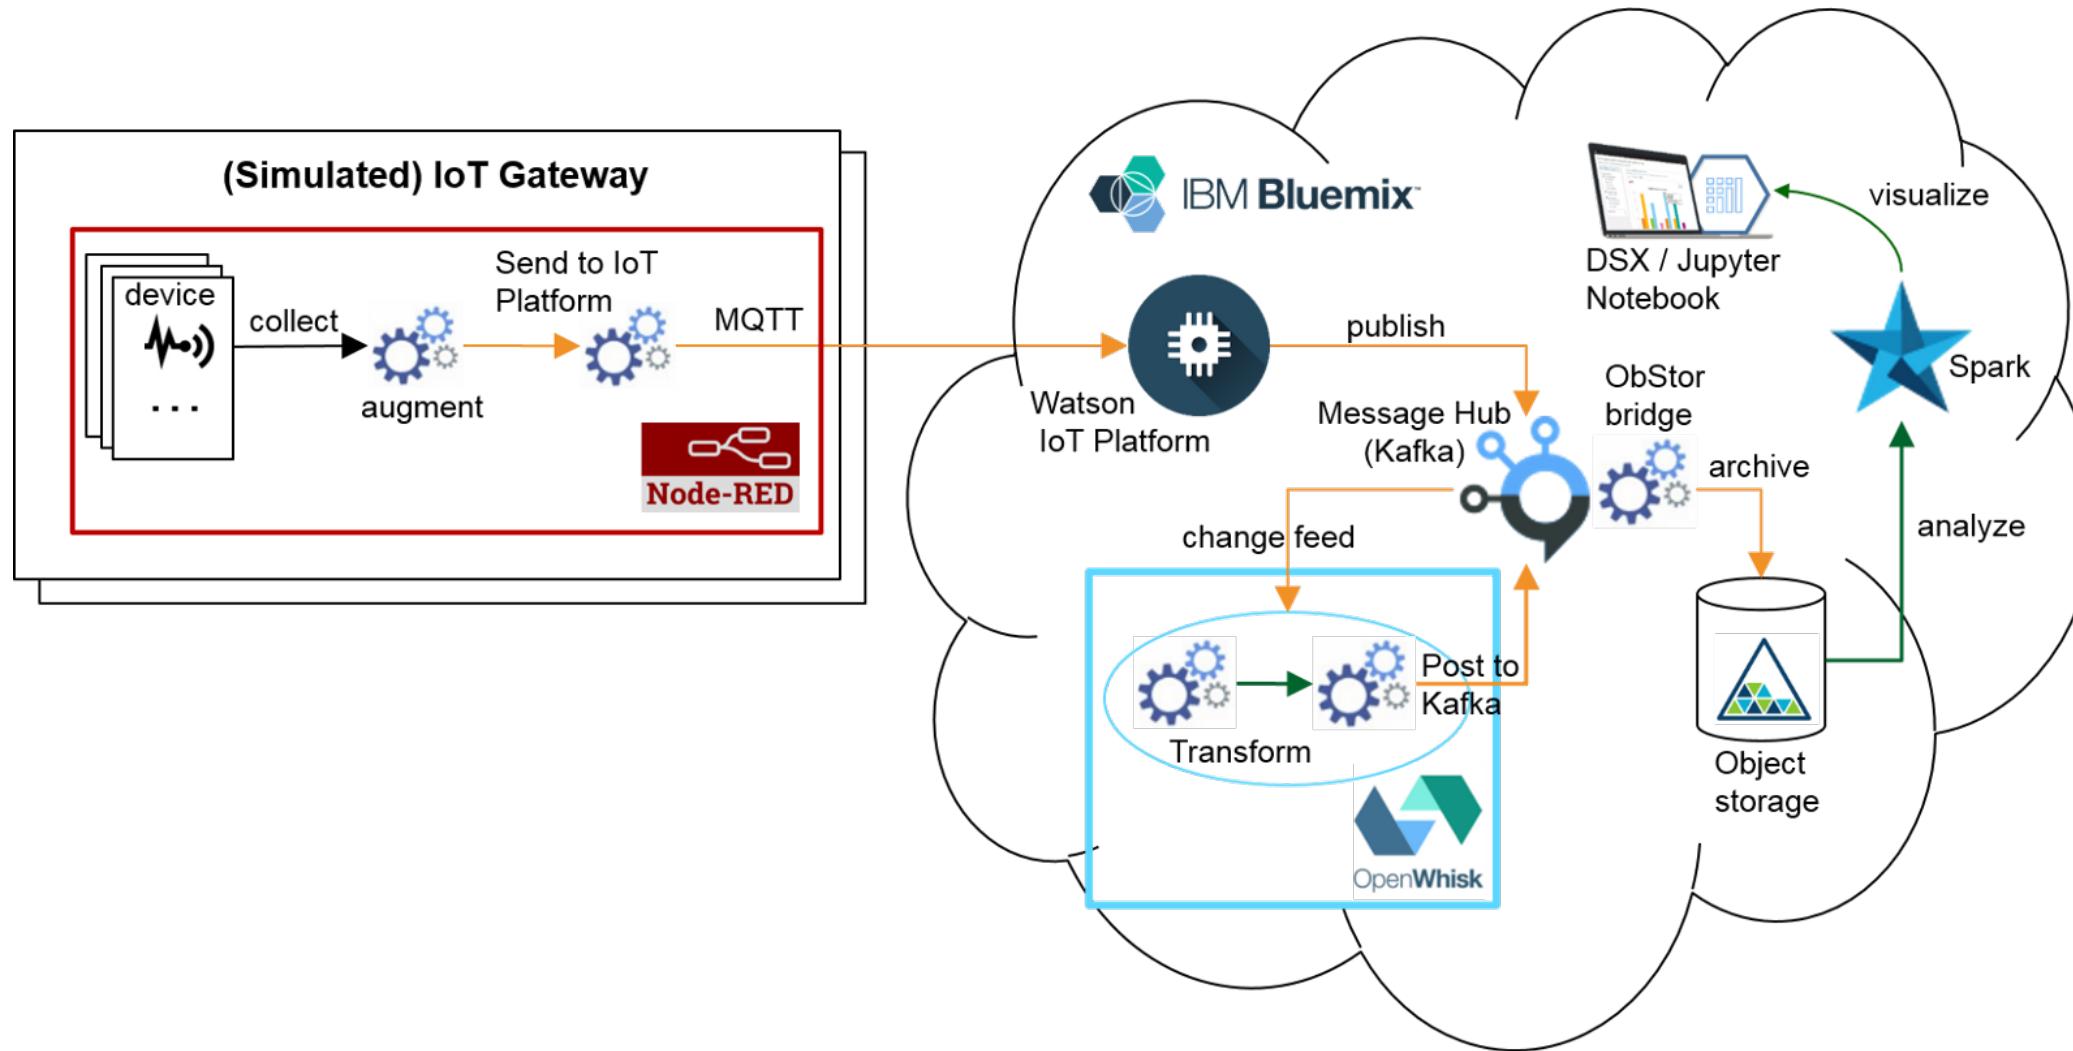
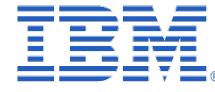
# End-to-end data flow – 4 (Message Hub)



## ▪ **Message Hub (bridge to Object Storage):**

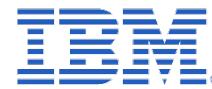
- Aggregates messages published to the transformed topic according to a specified policy (e.g., files of up to 1 hour or 1MB of data), and
- Uploads the files to Object Storage according to the specified layout (e.g., applying date-based partitioning).

# TRANSIT: Serverless Transformation of IoT Data-in-Motion

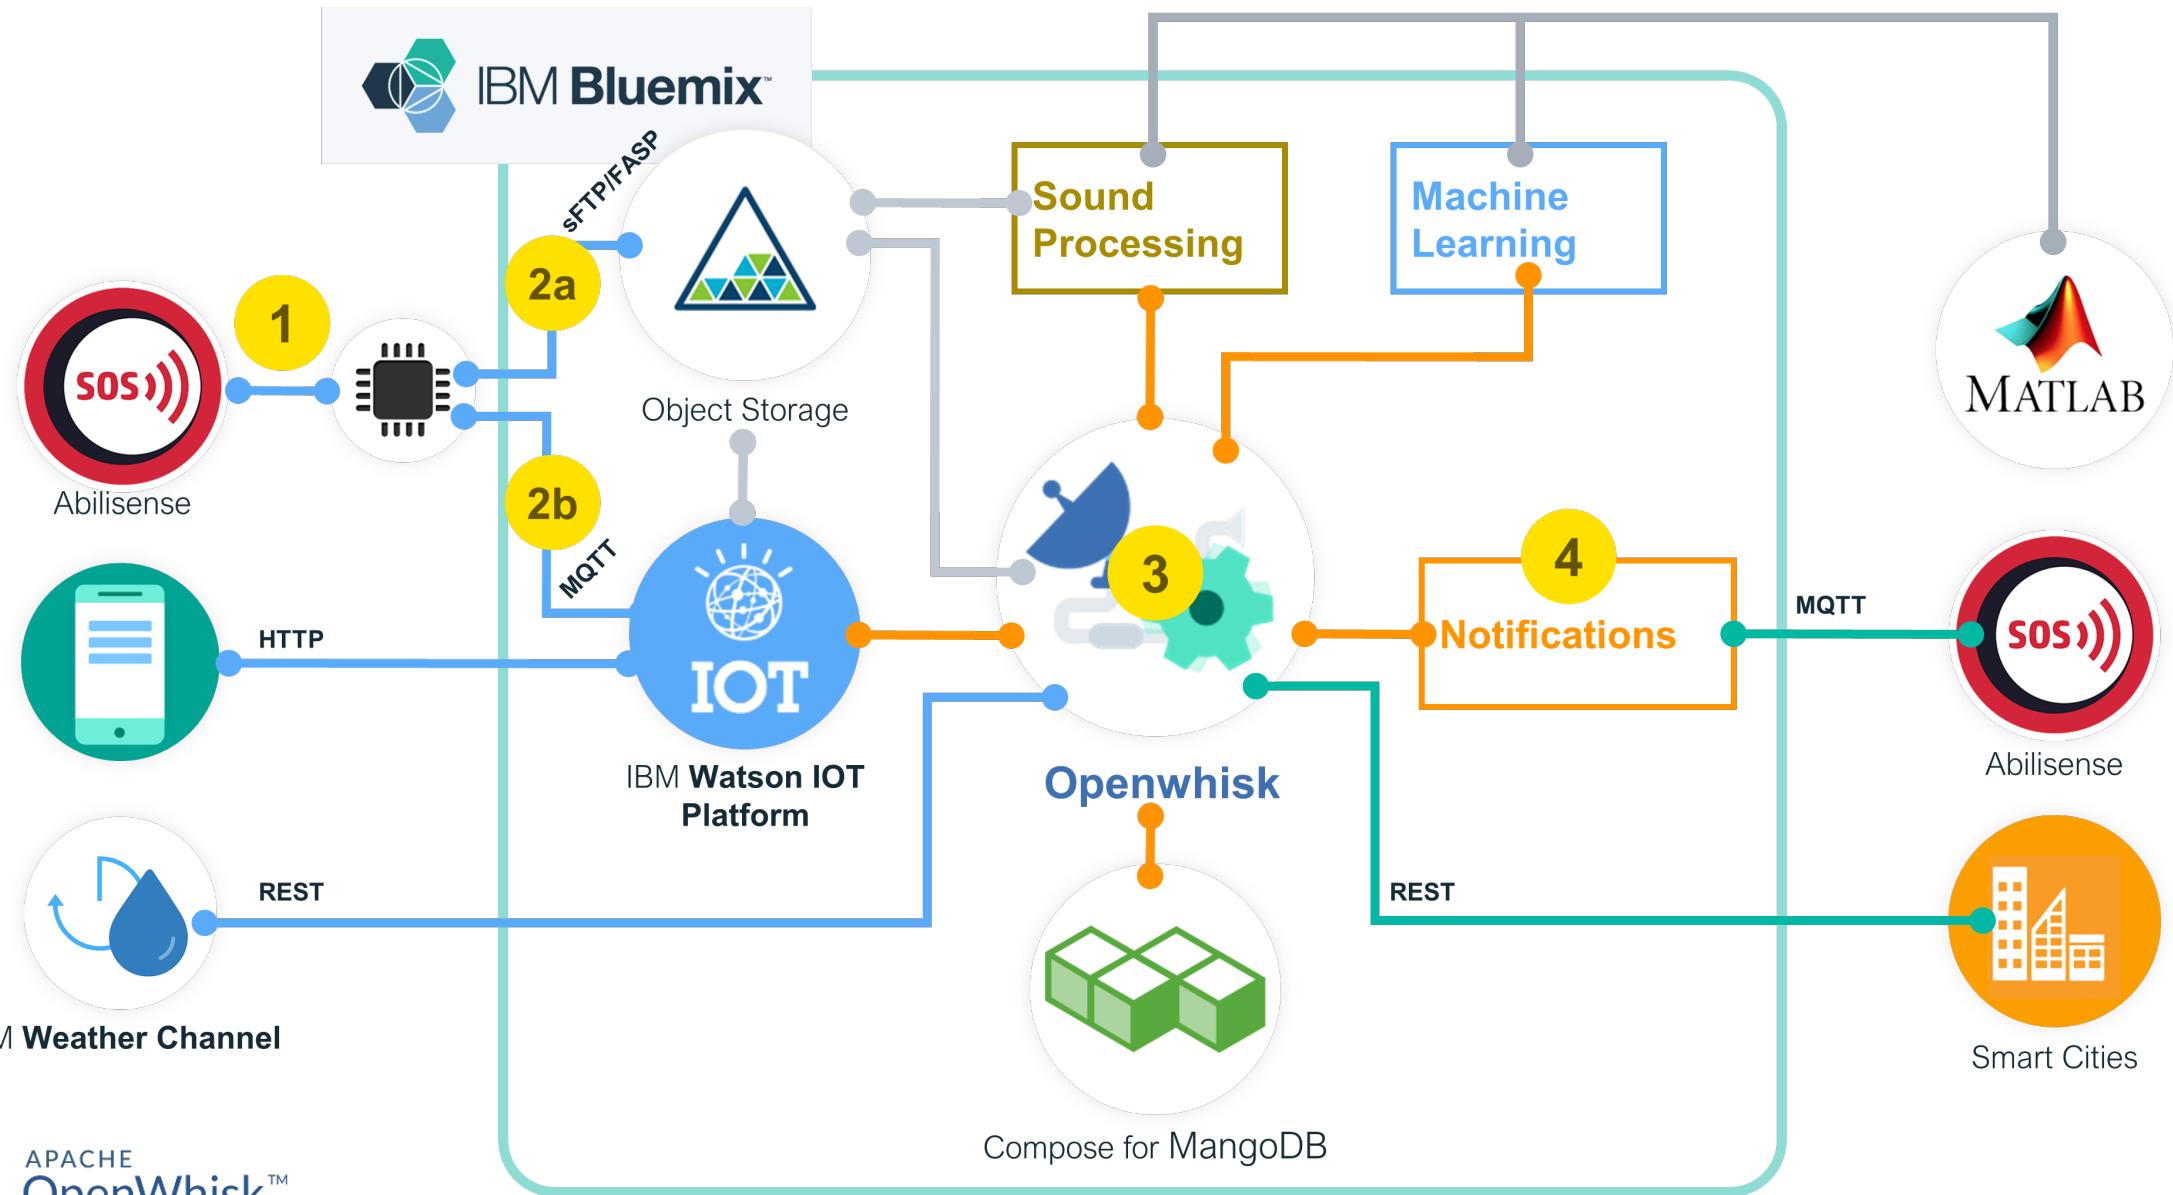


<https://medium.com/openwhisk/serverless-transformation-of-iot-data-in-motion-with-openwhisk-272e36117d6c>

# Example: Citizen Safety



<https://www.abilisense.com/>



# 3. Serverless Application Patterns for the IoT

# 4 Serverless/FaaS Patterns for IoT Applications

## Pattern #1: FaaS-based micro-services (decoupled, event-driven)

- a) Business logic is decoupled into self-contained micro-services, implemented as stateless Functions
- b) The business logic is triggered by IoT events (e.g., ingested via MQTT)
- c) Functions communicate via events or state changes in dedicated persistent services (e.g., data stores)

## Pattern #2: FaaS as a ‘glue’ integrating services in an event- or data-driven flow

- a) The application is heavily relying on existing services (e.g., cognitive, data, analytics)
- b) The ‘orchestration’ involves lightweight custom code, triggered by events or data
- c) Each handler interacting with external services is encapsulated as a stateless short-living Function

## Pattern #3: Acting upon outstanding conditions in stream processing

- a) IoT events are ingested via a high-throughput channel (e.g. Message Bus)
- b) Stateful stream processing is applied on incoming events (e.g., anomaly detection)
- c) Upon detection of pre-defined conditions, a Function is triggered to apply custom business logic

## Pattern #4: FaaS-enabled edge analytics

- a) Business logic deployed on individual devices or gateways (e.g., due to latency, bandwidth, privacy)
- b) Functions are portable across cloud and edge, enabling streamlined app development and operation

- **FaaS platforms enable rapid development of loosely-coupled event-driven applications**
  - Increased agility, manageability, elasticity
- **OpenWhisk is a leading open source FaaS project**
  - Incubating under Apache, and available commercially by several vendors (including IBM)
- **Serverless design patterns are still emerging, but there are many examples of leveraging FaaS in IoT applications**
  - Pattern #1: FaaS-based micro-services (decoupled, event-driven)
  - Pattern #2: FaaS as a ‘glue’ integrating services in an event- or data-driven flow
  - Pattern #3: Acting upon outstanding conditions in stream processing
  - Pattern #4: FaaS-enabled Edge Analytics (\*\*work in progress\*\*)

# Questions?

Alex Glikson

<https://medium.com/@glikson>

<http://slideshare.net/AlexGlikson/>

@glikson

# The Evolution of Cloud-Native Applications

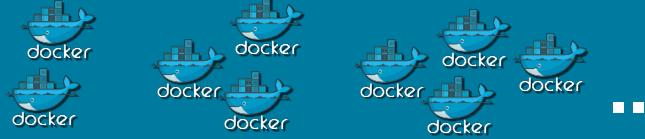


## Cloud Platforms

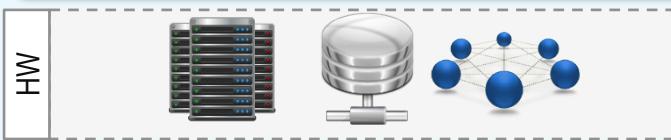
### Function-as-a-Service (FaaS)



### Micro Services (CaaS)



### Infrastructure-as-a-Service (IaaS)



Evolution ↑

### Serverless (“3.0”): Functions-as-a-Service (FaaS)

- **Architecture:** event-driven flow of single-tasked functions
- **Unit of deployment:** Functions, collection of functions
- **Connectivity:** de-coupled (e.g., via state changes/events)
- **Scaling:** per-request (with optimization of warm-up times)

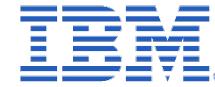
### Micro Services (“2.0”): Container-as-a-Service (CaaS)

- **Architecture:** composition of small, self-contained services, calling each other via a REST API
- **Unit of deployment:** group of containers comprising a service
- **Connectivity:** run-time service discovery and routing
- **Scaling:** container-based auto-scaling of each micro-service

### Scale-out (“1.0”): Infrastructure-as-a-Service (IaaS)

- **Architecture:** layered into few application tiers (often 3-tier)
- **Unit of deployment:** Virtual Machine (VM), collection of VMs
- **Connectivity:** configuration-time routing between tiers
- **Scaling:** VM-based (auto-) scaling within each tier

# Applying the Serverless Paradigm to IoT Applications



## ■ The Serverless/FaaS paradigm has great potential for IoT

- Event-driven programming model: a natural fit for many IoT applications
- Rapid programmability: lack of standards results in high heterogeneity and the need for customization
- Low entry barrier (pay-per-use, low ops overhead): attractive for the ‘long tail’ of DIY-style applications

## ■ Challenges

- New paradigm, lack of best practices ('patterns')
- Lack of mature developer tools
- Need more powerful programming model (e.g., for stateful stream processing)
- Need a holistic approach that spans the cloud and the edge

